

# Machine Learning Sberbank Kaggle Competition

April 24, 2024

Authors: Lidor Erez 318661444, Daniel Mizrachi 208727230, Dvir Rehavi 207206624

## 1 Libraries

Install if you need:

```
[8]: # !pip install plotly
      # !pip install hyperopt
```

```
[2]: # Data Manipulation
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.io as pio
pio.renderers.default='svg'

# Machine Learning
from sklearn.model_selection import train_test_split, cross_val_score, \
    RandomizedSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.impute import KNNImputer
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error
from hyperopt import fmin, tpe, hp, STATUS_OK, Trials
from hyperopt.pyll.base import scope
```

```
[3]: train = pd.read_csv("https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/
    ↪main/Data/train.csv")
test = pd.read_csv("https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/
    ↪main/Data/test.csv")
```

## 2 Baseline Models

```
[11]: cols = ['full_sq', 'life_sq', 'floor', 'max_floor', 'material',  
            'build_year', 'num_room', 'kitch_sq', 'state', 'product_type',  
            ↪ 'sub_area', 'price_doc']  
train_cp = train.copy(deep=True)[cols]  
test_cp = test.copy(deep=True)[cols[:-1]]
```

### 2.1 Missing Values Handling

```
[12]: train_cp.isna().sum()
```

```
[12]: full_sq          0  
life_sq          6383  
floor           167  
max_floor       9572  
material        9572  
build_year     13605  
num_room        9572  
kitch_sq        9572  
state          13559  
product_type     0  
sub_area         0  
price_doc        0  
dtype: int64
```

```
[13]: test_cp.isna().sum()
```

```
[13]: full_sq          0  
life_sq          1176  
floor            0  
max_floor        0  
material          0  
build_year       1049  
num_room          0  
kitch_sq          0  
state             694  
product_type      33  
sub_area          0  
dtype: int64
```

```
[14]: # Remove rows with more than 5 missing values.  
train_cp = train_cp.loc[train_cp.isna().sum(axis=1) < 5]  
  
# Handle missing values in build year
```

```

train_cp.loc[(train_cp['build_year'] == 0) | (train_cp['build_year'] == 1),
↳'build_year'] = None
train_cp['build_year'].fillna(train_cp['build_year'].median(numeric_only=True),
↳inplace=True)

test_cp.loc[(test_cp['build_year'] == 0) | (test_cp['build_year'] == 1),
↳'build_year'] = None
test_cp['build_year'].fillna(test_cp['build_year'].median(numeric_only=True),
↳inplace=True)

# Handle Missing values in life_sq
train_cp['life_sq'].fillna(train_cp['life_sq'].mean(numeric_only=True),
↳inplace=True)
test_cp['life_sq'].fillna(test_cp['life_sq'].mean(numeric_only=True),
↳inplace=True)

# Handle Missing values in categorical
train_cp['state'].fillna(train_cp['state'].mode()[0], inplace=True)
test_cp['state'].fillna(test_cp['state'].mode()[0], inplace=True)

train_cp['material'].fillna(train_cp['material'].mode()[0], inplace=True)
test_cp['material'].fillna(test_cp['material'].mode()[0], inplace=True)

test_cp['product_type'].fillna(test_cp['product_type'].mode()[0], inplace=True)

# Check how many NAs there are:

print(f'Test data:\n {test_cp.isna().sum()} \n\n Train Data: \n{train_cp.
↳isna().sum()}')

```

Test data:

full_sq	0
life_sq	0
floor	0
max_floor	0
material	0
build_year	0
num_room	0
kitch_sq	0
state	0
product_type	0
sub_area	0
dtype: int64	

Train Data:

full_sq	0
---------	---

```
life_sq      0
floor        0
max_floor    0
material     0
build_year   0
num_room     0
kitch_sq     0
state        0
product_type 0
sub_area     0
price_doc    0
dtype: int64
```

## 2.2 Data Types Handling

```
[15]: train_cp.dtypes
```

```
[15]: full_sq      int64
life_sq      float64
floor        float64
max_floor    float64
material     float64
build_year   float64
num_room     float64
kitch_sq     float64
state        float64
product_type  object
sub_area     object
price_doc    int64
dtype: object
```

```
[16]: test_cp.dtypes
```

```
[16]: full_sq      float64
life_sq      float64
floor        int64
max_floor    int64
material     int64
build_year   float64
num_room     int64
kitch_sq     float64
state        float64
product_type  object
sub_area     object
dtype: object
```

```
[17]: categorical_cols = ['state', 'material', 'sub_area', 'product_type']
      to_int_cols = ['num_room', 'floor', 'max_floor', 'build_year']
      to_float_cols = ['full_sq', 'price_doc']

      # Changing some of the float columns to int.
      train_cp[to_int_cols] = train_cp[to_int_cols].astype('int64')
      test_cp[to_int_cols] = test_cp[to_int_cols].astype('int64')

      # From int to float

      train_cp[to_float_cols] = train_cp[to_float_cols].astype('float64')
      test_cp[to_float_cols[:-1]] = test_cp[to_float_cols[:-1]].astype('float64')

      # Categorical columns to int.
      train_cp[categorical_cols[:2]] = train_cp[categorical_cols[:2]].astype('int64')
      test_cp[categorical_cols[:2]] = test_cp[categorical_cols[:2]].astype('int64')

      numeric_columns = list(train_cp.select_dtypes(['float64']).columns)
```

## 2.3 Outliers Handling

```
[18]: def categoricalOutlierHandler(df, cat_cols):
      vals_to_delete = {}

      for col in cat_cols:
          # Check for relative frequenecies
          freq = df[col].value_counts()
          # calculate the 10th quantile of the frequencies distribution
          q5 = freq.quantile(0.05)
          # Get rid of values that return less than the 10th quantile.
          vals_to_delete[col] = list(freq[freq < q5].index)

      for key in vals_to_delete:
          df = df[~df[key].isin(vals_to_delete[key])]

      return df

      # categorical columns without product type (product type is binary).

      train_cp = categoricalOutlierHandler(train_cp, ['state', 'material']).
      ↪reset_index(drop=True)
```

```
[19]: def numericOutlierHandler(df, numeric_cols):
    for col in numeric_cols:
        # Calculate iqr
        q25 = np.quantile(df[col], 0.25)
        q75 = np.quantile(df[col], 0.75)
        iqr = q75 - q25
        lower_val = q25 - 1.5*iqr
        upper_val = q75 + 1.5*iqr

        # Cap the outliers using the upper val and lower val.
        df.loc[df[col] > upper_val, col] = upper_val
        df.loc[df[col] < lower_val, col] = lower_val
    return df

# Get all numeric columns
train_cp = numericOutlierHandler(train_cp, numeric_columns)
```

## 2.4 Categorical Features Handling

```
[20]: # Product Type has only 2 values we can change it to binary
train_cp['product_type'] = train_cp['product_type'].apply(lambda x: (x == 'Investment')*1)
test_cp['product_type'] = test_cp['product_type'].apply(lambda x: (x == 'Investment')*1)

# Dummies for sub area
train_cp = train_cp.drop(columns = 'sub_area').join(pd.get_dummies(train_cp['sub_area'])*1).reset_index(drop=True)
test_cp = test_cp.drop(columns = 'sub_area').join(pd.get_dummies(test_cp['sub_area'])*1).reset_index(drop=True)
```

## 2.5 Feature Engineering

```
[21]: train_cp['room_sq'] = train_cp['num_room']/(train_cp['full_sq'] + 1e-10)
test_cp['room_sq'] = test_cp['num_room']/(test_cp['full_sq'] + 1e-10)
```

## 2.6 Train - Validation - Test Split

```
[22]: X = train_cp.drop(columns = 'price_doc').reset_index(drop=True)
y = train_cp['price_doc']

x_train, x_testval, y_train, y_testval = train_test_split(X, y, test_size=0.3)
x_val, x_test, y_val, y_test = train_test_split(x_testval, y_testval, test_size=0.15)
```

```
[23]: # Get all columns that are in X and not in Test Data Frame
```

```
X.columns.difference(test_cp.columns)
```

```
[23]: Index(['Poselenie Klenovskoe'], dtype='object')
```

```
[24]: test_cp['Poselenie Klenovskoe'] = 0
```

```
[25]: X.columns.difference(test_cp.columns)
```

```
[25]: Index([], dtype='object')
```

## 2.7 Random Forest Regressor

```
[26]: rf = RandomForestRegressor()  
rf.fit(x_train, y_train)  
# Get the most important features:  
importance = rf.feature_importances_
```

```
[27]: # Get the most important features  
important_features = list(x_train.columns[np.where(importance > np.  
↳median(importance))])  
x_train_subset = x_train[important_features]  
rf.fit(x_train_subset, y_train)
```

```
[27]: RandomForestRegressor()
```

```
[28]: # Cross Validation  
cv_mse =   
↳cross_val_score(rf, x_val[important_features], y_val, cv=5, scoring='neg_mean_squared_error')  
print(f'Random Forest MSE Score {np.mean(-cv_mse)}')
```

Random Forest MSE Score 5134281376847.779

```
[29]: mean_squared_error(rf.predict(x_test[important_features]), y_test)
```

```
[29]: 4673306534227.724
```

## 2.8 XGBoost Regressor

```
[30]: xgb = XGBRegressor()  
params = {  
    'learning_rate': np.arange(0.01, 0.2),  
    'max_depth' : np.arange(3, 10),  
    'gamma' : [0.0, 0.1, 0.2, 0.3, 0.4],  
    'n_estimators': [500, 600, 700, 800, 900]  
}
```

```
rs_model = RandomizedSearchCV(xgb, param_distributions=params, n_iter=5,
    ↳scoring='neg_mean_squared_error',cv=5,n_jobs=-1, verbose=3)
rs_model.fit(x_train, y_train)
rs_model.best_params_
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

```
[30]: {'n_estimators': 600, 'max_depth': 8, 'learning_rate': 0.01, 'gamma': 0.4}
```

```
[31]: new_xgb = XGBRegressor(**rs_model.best_params_)
new_xgb.fit(x_train, y_train)
cv_mse_xgb = cross_val_score(new_xgb,x_val, y_val, cv=5,
    ↳scoring='neg_mean_squared_error')
f'XGBoost Regressor MSE: {np.mean(-cv_mse_xgb)}'
```

```
[31]: 'XGBoost Regressor MSE: 4796309610025.951'
```

```
[32]: mean_squared_error(new_xgb.predict(x_test),y_test)
```

```
[32]: 4521142820948.662
```

## 2.9 Kaggle Submission Files

```
[33]: # predict test csv column and upload to kaggle.
xgb_preds = new_xgb.predict(test_cp[X.columns])
rf_preds = rf.predict(test_cp[important_features])

xgb_submit = pd.DataFrame({'id':test['id'],'price_doc':xgb_preds})
rf_submit = pd.DataFrame({'id':test['id'],'price_doc':rf_preds})
```

```
[34]: # create csv files for submissions:

xgb_submit.to_csv('xgb_baseline_submission.csv',index=False)
rf_submit.to_csv('rf_baseline_submission.csv',index=False)
```

## 3 Notebooks Summary

### 3.1 Basic Time Series Analysis & Feature Selection

#### 3.1.1 Methodology:

##### 1. Data Preprocessing:

- Few functions for data transformation were created (log transformation).
- Columns that hold more 20% missing values and above were removed from the data set.
- Duplicated features were removed.

##### 2. Exploratory Data Analysis:



- a. The distribution of log price was examined.
- b. The trend of price by months was examined.

### 3. Feature Engineering:

- a. Time frame columns: time frames columns such as year, month, and date were created.
- b. Log Price: log transformation to the response variable (price\_doc)

### 4. Feature Selection:

- a. Correlation Test between every feature in the data set
- b. Correlation Test between every feature in the data set with log price
- c. Using XGBoost for feature selection.
- d. Choosing features according both correlation test and XGBoost importance.

### 5. Data Preprocessing (Before deploying second XGBoost):

- a. Missing values imputation using most frequent value.
- b. Normalizing the data using 'l2'.
- c. Label Encoder was used to transform category features.

### 6. XGBoost Model:

- a. Training a model using K-Fold Cross Validation method with a rmse metric.
- b. Get the top 10 important features and examine their relationship with log price.

## 3.2 Results:

- 1. Some of the features are highly correlated with each other.
- 2. The highest correlation with the log price is less than 0.040.
- 3. Total area of the apartment is the most important feature to the first XGBoost model.
- 4. Time frame features created on the feature engineering phase have high importance.
- 5. Model's RMSE is 0.46%

## 3.3 Conclusions:

- 1. Reducing the number of features is a crucial step because of the high correlation.
- 2. No strong linear relation with log price was detected.
- 3. Time features seem to contribute a lot to the model even though there's no meaning to the order of the values (the values weren't sorted).
- 4. After capping outliers for each of the top 10 features they seem to have a strong relationship with log price.

### 3.4 Critics:

1. Training XGBoost using cross validation is a smart decision, however when using a time series data, it's important to consider the time which the event occurred while training the model. Thus, we suggest using a time series version of the cross-validation method.
2. Removing correlated features from the data set might overcome this problem however one can also use PCA to reduce the dimensionality in the data. Thus, we suggest also trying using other methods but removing.

### 3.5 A Very Extensive Exploratory Analysis in Python

#### 3.5.1 Methodology:

1. **Data Preprocessing:** Corrects data quality issues, ensuring the dataset's integrity and reliability.
2. **Exploratory Data Analysis:**
  - a. Missing Values: Visualize the number of missing values in each column.
  - b. Housing Internal Characteristics: Investigates features like floor area, number of rooms and building material to understand their influence on apartment prices.
  - c. Time Series Analysis: The author Investigated how does price changes in several time frames.
  - d. School Characteristics: Explores variables related to school facilities and their association with housing prices.
  - e. Cultural Characteristics: Analyzes the impact of proximity to cultural landmarks and recreational facilities on property prices.
  - f. Infrastructure Features: Examines proximity to infrastructure such as public transport, parks, and utilities in relation to housing prices.
  - g. Variable Importance: Built a Random Forest Regressor model to understand which features are most important (categorical features were encoded using label encoder).
  - h. Train – Test Comparison: Compared between the train and test data to understand if they're different or not.

#### 3.5.2 Results:

1. **Housing Internal Characteristics:**
  - a. Number of rooms in the apartment exhibits a high correlation with the total area of the apartment (makes sense).
  - b. None of the features exhibit a strong linear relation with the price though number of rooms has a correlation score of 0.48 with the price.
  - c. Housing internal features are the most important to the model.
2. **School , Cultural and Infrastructure Features:**
  - a. Some of the features exhibits high correlation with each other.

- b. There is no strong linear relation between the features and the price.

### 3.5.3 Conclusions:

1. Most important variables are housing internal features.
2. There might be redundant information in the data (high correlation between predictors).
3. The relation between price and other features is probably not linear. However, a statistical test needs to be done to support this assumption.
4. The train and test datasets are different from each other and might hold different values / number of values in each column.

### 3.5.4 Critic:

1. Categorical features were encoded using label encoder. However, in some cases it's better to use one hot encoding. Since label encoder might imply ordinality to a predictor which the ordinality is meaningless (such as sub area). Thus, we suggest using one hot encoding (dummies) instead.
2. Random Forest Regressor was used, and its performance wasn't tested. This could've been addressed by Cross Validation to provide a more robust assessment of its predictive capabilities and help validate the results of the model.
3. Instead of removing missing values, considering imputation or interpolation method could retain valuable information and prevent potential bias in the analysis.
4. While using Random Forest's feature importance is a smart decision, feature selection based on more than one method might lead to a better result. Thus, we suggest combining few feature selection methods to better understand which of the features contribute the most to the model.

## 4 Advanced Modeling

- Data Preprocessing (Integrity, Missing Values, Outliers)
- Feature Engineering
- Feature Selection (Importance, Correlation, Step-Wise regression?, LASSO)
- Hyperparameter Tuning (Bayesian Optimization, GridSearch, RandomSearch)
- Meta Learner (SVM, Linear Regression, Polynomial Regression, LASSO, Ridge)

### 4.1 Data Integrity (Quality)


We made sure the data make sense in both train and test data frames

```
[35]: full_data = pd.concat([train,test]).reset_index(drop=True)
```

```
[36]: interior_cols = ['id', 'timestamp', 'full_sq', 'floor', 'num_room',  
    ↪ 'product_type',  
    ↪ 'sub_area', 'price_doc',  
    ↪ 'life_sq', 'max_floor', 'build_year', 'kitch_sq', 'state', 'material']
```

```
[37]: interiors = full_data[interior_cols]
```

```
[38]: px.histogram(interiors['full_sq'])
```




#### 4.1.1 Full Square

```
[39]: # Life square supposed to be at least 0.3 of the full square (conservative
      ↪ value)
interiors.loc[interiors.full_sq < 10, 'full_sq'] = np.nan # full sq is too low..
      ↪ .
interiors.loc[(interiors.life_sq < 0.3*interiors.full_sq)*(interiors.full_sq >
      ↪ 210), 'full_sq'] = np.nan # full sq is probably wrong.
```

#### 4.1.2 Life Square


```
[40]: px.histogram(interiors.life_sq)
```



```
[41]: interiors.loc[(interiors.life_sq < 5), 'life_sq'] = np.nan # deal with bad life_
      ↪square.
      interiors.loc[(interiors.life_sq >= interiors.full_sq), 'life_sq'] = np.nan
```

#### 4.1.3 Kitchen Square

```
[42]: px.histogram(interiors['kitch_sq'])
```




```
[43]: # Remove too small values
interiors.loc[interiors['kitch_sq'].isin([0,1,2,3]), 'kitch_sq'] = np.nan
# Remove abnormal values

interiors.loc[(interiors['kitch_sq'] >= interiors['full_sq']) |
↳(interiors['kitch_sq'] >= interiors['life_sq']), 'kitch_sq'] = np.nan
```

#### 4.1.4 Number of rooms

```
[44]: px.histogram(interiors.num_room)
```




```
[45]: # Finding apartment with a incorrect number of rooms.
threshold_for_sq_per_room = 5

interiors.loc[(interiors['life_sq']/interiors['num_room'] <
↳threshold_for_sq_per_room) | (interiors.num_room <= 0), 'num_room'] = np.nan
```


#### 4.1.5 Floor & Max Floor

```
[46]: px.histogram(interiors.floor)
```



```
[47]: px.histogram(interiors.max_floor)
```





```
[48]: interiors.loc[interiors['floor'] == 77, 'floor'] = np.nan
interiors.loc[interiors['floor'] > interiors['max_floor'], 'max_floor'] = np.nan
interiors.loc[interiors['floor'] == 0, 'floor'] = np.nan
interiors.loc[interiors['max_floor'] == 0, 'max_floor'] = np.nan
```

#### 4.1.6 Build Year

```
[49]: interiors.build_year.value_counts().index.unique()
```

```
[49]: Index([    2014.0,    2015.0,         0.0,    2016.0,    2013.0,    2017.0,
              1.0,    1969.0,    1970.0,    1968.0,
              ...,
              1945.0,     71.0,    1904.0, 20052009.0,    1876.0,    1886.0,
              1925.0,   1691.0,     20.0,    1898.0],
            dtype='float64', name='build_year', length=127)
```

```
[50]: # Build Year
interiors.loc[(interiors['build_year'] < 1500) , 'build_year'] = np.nan
interiors.loc[interiors['build_year'] == 4965, 'build_year'] = 1965
interiors.loc[interiors['build_year'] == 20052009, 'build_year'] = np.nan
```

```
[51]: # Product type
interiors['product_type'].value_counts()
```

```
[51]: product_type
Investment      24446
OwnerOccupier   13654
Name: count, dtype: int64
```

```
[52]: interiors['product_type'].isnull().sum()
```

```
[52]: 33
```

```
[53]: interiors['product_type'].fillna(interiors['product_type'].mode()[0],
    ↪inplace=True)
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3817222879.py:1
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[54]: # State
print(interiors['state'].value_counts())
# Remove outlier.
interiors.loc[interiors['state'] == 33.0, 'state'] = 3
```

```
state
2.0      8506
3.0      7703
1.0      7121
4.0       549
33.0        1
Name: count, dtype: int64
```

```
[55]: # Material
interiors['material'].value_counts()
```

```
[55]: material
1.0      19438
2.0       3951
5.0       2048
4.0       1963
6.0       1159
3.0         2
Name: count, dtype: int64
```

```
[56]: full_data[interiors.columns] = interiors
```

## 4.2 Data Cleaning

### 4.2.1 Missing Values - Shallow Cleaning

The way we cleaned our data:

1. Got the list of the features that has more than 30% missing values.
2. Calculated each feature (from the list above) correlation with the response variable.
3. If a feature had more than 30% missing values and a correlation below 0.3 (in absolute value) it was removed.

```
[57]: missing_vals = pd.DataFrame(full_data.isnull().mean()*100)
features = missing_vals.index
vals = missing_vals.values.flatten()

# Create a dataframe with the pct of missing values for each feature
missing_vals = pd.DataFrame({'Features':features,'vals':vals})

# Get the feature
features_to_delete = missing_vals[missing_vals['vals'] > 30]['Features']

# Get the corrs with price
corss = np.round(full_data[['price_doc'] + features_to_delete.to_list()].
    ↪corr(),2).loc['price_doc']

# Remove the first correlation (price with itself).
corrs_with_price = corss.to_list()[1:]

corr_w_price = pd.DataFrame({'Features to delete':features_to_delete,'corr with_
    ↪price':corrs_with_price})

corr_w_price
```

```
[57]:
```

	Features to delete	corr with price
3	life_sq	0.56
5	max_floor	0.13
7	build_year	0.04
9	kitch_sq	0.38
10	state	0.13
24	hospital_beds_raion	0.15
160	cafe_sum_500_min_price_avg	0.04
161	cafe_sum_500_max_price_avg	0.04
162	cafe_avg_price_500	0.04

We decided to not delete max floor, state, build year and hospital beds since they give us valuable information about surroundings, building and apartment

```
[58]: corr_w_price.drop([7,24,5,10], inplace=True)
```

```
[59]: # Get the features that has low correlation with the response variable
features_to_delete = corr_w_price[abs(corr_w_price['corr with price']) < 0.
↳3]['Features to delete'].tolist()
print(features_to_delete)
full_data = full_data.drop(columns = features_to_delete)
```

```
['cafe_sum_500_min_price_avg', 'cafe_sum_500_max_price_avg',
'cafe_avg_price_500']
```

#### 4.2.2 Missing Values - Deep cleaning

By reading the data dictionary text file, we understood that some of the neighborhood features can be separated into groups. Thus, we created separated text files for each group of features and uploaded it to github so it'd be much easier to divide them into groups

```
[60]: import requests

# Paths to text files from github
paths = ['https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/areas.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/buildings.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/demographics.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/distances.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/education.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/facilities.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/interior.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/surroundings.txt']

# Groups
groups = {
    'areas': [],
    'buildings': [],
    'demographics': [],
    'distances': [],
    'education': [],
    'facilities': [],
    'interior': [],
    'surroundings': []
}
```

```

}

keys = list(groups.keys())

for i in range(len(paths)):
    # Http request to get the file data
    response = requests.get(paths[i])
    # read File lines
    lines = response.text.splitlines()
    for line in lines:
        # Split by \t
        col = line.split('\t')[0]
        # Interior file has : in it.
        if keys[i] == 'interior':
            col = col.split(':')[0]
        groups[keys[i]].append(col)

```

```

[61]: def naIndicator(data, groups):
        nas_indicators = []
        for group in groups:
            data_cols = data.columns[np.where(data.columns.isin(groups[group]))]
            na_indicator = np.any(data[data_cols].isnull().sum() > 0)
            nas_indicators.append((group, na_indicator))
        return nas_indicators

```

```

[62]: naIndicator(full_data, groups)

```

```

[62]: [('areas', False),
        ('buildings', True),
        ('demographics', False),
        ('distances', True),
        ('education', True),
        ('facilities', True),
        ('interior', True),
        ('surroundings', True)]

```

It seems that few groups are still suffering from missing data.

```

[63]: # Create a data frame for each group:
        groups_dfs = {group:full_data[full_data.columns[np.where(full_data.columns.
↪isin(groups[group]))]] for group in groups}

```

## Buildings

```

[64]: buildings = groups_dfs['buildings']
        # Get rows with missing values.
        missing_vals = buildings.loc[buildings.isna().sum(axis=1) > 0]

```

```
# Check if the missing values df has some data in it.
print(f'Number of rows in df: {missing_vals.shape[0]} \n\n Number of missing_
↪values in each columns: \n\n {missing_vals.isna().sum()}')
```

Number of rows in df: 6209

Number of missing values in each columns:

```
raion_build_count_with_material_info    6209
build_count_block                       6209
build_count_wood                        6209
build_count_frame                       6209
build_count_brick                       6209
build_count_monolith                    6209
build_count_panel                       6209
build_count_foam                        6209
build_count_slag                        6209
build_count_mix                         6209
raion_build_count_with_builddate_info    6209
build_count_before_1920                  6209
build_count_1921-1945                    6209
build_count_1946-1970                    6209
build_count_1971-1995                    6209
build_count_after_1995                   6209
dtype: int64
```

In order to fill these columns we need domain knowledge. However each of these columns represent an amount therefore we'll replace nan values with 0.

```
[65]: buildings[buildings.filter(like='count').columns] = buildings[buildings.
↪filter(like='count').columns].replace(np.nan, 0) # 0 for count data
buildings[buildings.filter(like='info').columns] = buildings[buildings.
↪filter(like='info').columns].replace(np.nan,-1)
buildings.isnull().sum()
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/222925215.py:1:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/222925215.py:2:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[65]: raion_build_count_with_material_info      0
      build_count_block                        0
      build_count_wood                        0
      build_count_frame                       0
      build_count_brick                       0
      build_count_monolith                    0
      build_count_panel                       0
      build_count_foam                       0
      build_count_slag                       0
      build_count_mix                         0
      raion_build_count_with_builddate_info    0
      build_count_before_1920                 0
      build_count_1921-1945                   0
      build_count_1946-1970                   0
      build_count_1971-1995                   0
      build_count_after_1995                  0
      dtype: int64
```

```
[66]: groups_dfs['buildings'] = buildings
```

```
[67]: groups_with_missing_vals = ['interior', 'distances', 'surroundings']
```

### Interior

```
[68]: interior = groups_dfs['interior']
      # Get interior Data Frame
      interior_na = interior.isnull().sum()
      print(interior_na[interior_na > 0])
```

```
full_sq      40
life_sq     11762
floor        177
max_floor    11711
material     9572
build_year   16116
num_room     9600
kitch_sq    18118
state       14253
price_doc    7662
dtype: int64
```

**Few Helpers:**

```

[69]: def handleFullsqLifesq(data, nas_full, nas_life, sa_med_fullsq,
    ↪sa_life_full_prop):
    data['missing_life'] = data['life_sq'].isnull().astype(int)
    data['missing_full'] = data['full_sq'].isnull().astype(int)
    for index in data.index:
        sub_area = data.loc[index, 'sub_area']

        fullsq = data.loc[index, 'full_sq']
        lifesq = data.loc[index, 'life_sq']

        mask_sa_pct = (sa_life_full_prop['sub_area'].isin([sub_area]))
        mask_sa_full = (sa_med_fullsq['sub_area'].isin([sub_area]))

        sa_prop = sa_life_full_prop.loc[mask_sa_pct]['PCT'].values[0]
        full_sq_med = sa_med_fullsq.loc[mask_sa_full]['median'].values[0]

        flag = 0

        """
        If life sq is missing - fill it with full sq * avg pct of life sq out
        ↪of full sq.
        If full sq is missing - fill it with life sq / avg pct of life sq out
        ↪of full sq.
        If both are missing - fill full sq with median of full sq and then
        ↪calculate life sq using the median value.
        """
        if not nas_full[index] and nas_life[index]:
            flag = 2
            data.loc[index, 'life_sq'] = fullsq * sa_prop
        elif not nas_life[index] and nas_full[index]:
            flag = 3
            data.loc[index, 'full_sq'] = lifesq / sa_prop
        elif nas_life[index] and nas_full[index]:
            flag = 4
            data.loc[index, 'full_sq'] = full_sq_med # get median value to fill
            ↪na
            data.loc[index, 'life_sq'] = full_sq_med * sa_prop # calculate
            ↪life sq using the median value

        # if check to see if the algorithm work correctly.
        if np.any(data['life_sq'] >= data['full_sq']):
            print(flag)
            break
    return {1: 'Done', 'Full SQ': data['full_sq'].isna().sum(), 'Life SQ':
    ↪data['life_sq'].isna().sum(), 'Life SQ >= Full SQ': (data['life_sq'] >=
    ↪data['full_sq']).any()}

```



```
[70]: def HandleFloorMaxFloor(data, sa):
    data['missing_floor'] = data['floor'].isnull().astype(int)
    data['missing_maxfloor'] = data['max_floor'].isnull().astype(int)
    # fill na in max floor
    med_max_floor = sa['max_floor'].transform('median')
    med_max_floor = np.round(med_max_floor.fillna(med_max_floor.median())) #
    ↪fill the missing value with the median of medians.
    data['max_floor'] = data['max_floor'].fillna(med_max_floor)
    # Fill floor:
    med_floor = sa['floor'].transform('median')
    med_floor = np.round(med_floor.fillna(med_floor.median())) # fill the
    ↪missing value with the median of medians.
    data.loc[:, 'floor'] = data['floor'].fillna(med_floor)

    data.loc[data['floor'] > data['max_floor'], 'floor'] = data['max_floor']

    return {1: 'Done', 'Max Floor': data['max_floor'].isna().sum(), 'Floor':
    ↪data['floor'].isna().sum(), 'Floor > Max Floor': (data['floor'] >
    ↪data['max_floor']).any()}
```

```
[71]: def HandleNumRooms(data, sa):
    data['missing_num_room'] = data['num_room'].isnull().astype(int)
    # fill na with the median number of rooms for each sub area and product
    ↪type.
    data.loc[:, 'num_room'] = data['num_room'].fillna(sa['num_room'].
    ↪transform('median'))

    # Make sure data makes sense.
    mask = data['life_sq']/data['num_room'] < 5
    # 5 square meter per room is a conservative value.
    data.loc[mask | (data['num_room'].isna()), 'num_room'] = data['life_sq'] //
    ↪5
    data.loc[data['num_room'] == 0, 'num_room'] = 1

    return {1: 'Done', 'Num Room': data['num_room'].isna().sum(), 'LifeSQ/
    ↪NumRoom < 5': (data['life_sq']/data['num_room'] < 5).any()}
```

```
[72]: def HandleBuildYear(data):
    data['missing_build_year'] = data['build_year'].isnull().astype(int)

    sa_median_build_year = data.groupby(['sub_area'])['build_year'].
    ↪transform('median')

    data.loc[:, 'build_year'] = data['build_year'].fillna(sa_median_build_year)

    return {1: "Done", "Build_Year": data['build_year'].isna().sum()}
```

```
[73]: def handleStateMaterial(data):
    data['missing_state'] = data['state'].isnull().astype(int)
    data['missing_material'] = data['material'].isnull().astype(int)

    data['year'] = pd.to_datetime(data['timestamp']).dt.year
    data['age'] = data['year'] - data['build_year']

    state_modes = data['state'].fillna(data.
↳groupby(['age', 'sub_area'])['state'].transform(lambda val: val.mode()[0] if
↳len(val.mode())>0 else None))
    state_modes = state_modes.fillna(state_modes.mode()[0])

    data.loc[:, 'state'] = data['state'].fillna(state_modes)

    material_modes = data['material'].fillna(data.
↳groupby(['build_year', 'sub_area'])['material'].transform(lambda val: val.
↳mode()[0] if len(val.mode())>0 else None))
    material_modes = material_modes.fillna(material_modes.mode()[0])
    data.loc[:, 'material'] = data['material'].fillna(material_modes)

    return {1: 'Done', 'State': data['state'].isna().sum(), 'Material':
↳data['material'].isna().sum()}
```

```
[74]: def HandleKitchSQ(data, sa_kitch_life_prop):
    data['missing_kitch_sq'] = data['kitch_sq'].isnull().astype(int)
    for idx in data.index:
        life_sq = data.loc[idx, 'life_sq']
        kitch_sq = data.loc[idx, 'kitch_sq']
        product_type = data.loc[idx, 'product_type']
        sub_area = data.loc[idx, 'sub_area']

        if not np.isnan(kitch_sq) and kitch_sq < life_sq:
            continue
        mask = (sa_kitch_life_prop['sub_area'].isin([sub_area]))
        avg_ratio = sa_kitch_life_prop.loc[mask]['PCT'].values[0]

        data.loc[idx, 'kitch_sq'] = life_sq * avg_ratio

    return {1: 'Done', 'Kitch SQ': data['kitch_sq'].isna().sum(), 'Kitch SQ >=
↳Full SQ': (data['kitch_sq'] >= data['full_sq']).any(), 'Kitch SQ >= Life SQ':
↳(data['kitch_sq'] >= data['life_sq']).any()}
```

```
[75]: # Calculate life sq / full sq
interior.loc[:, 'life_full_ratio'] = interior['life_sq']/interior['full_sq']
(interior['life_full_ratio'] >= 1).any()
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/4061903555.py:2
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

[75]: False

```
[76]: # product type and sub area group by
sa_grps = interior.groupby(['sub_area'])

# Calculate the avg pct of life sq out of full sq for each sub area:
avg_life_full_prop = (sa_grps['life_full_ratio'].mean()).reset_index().
    ↪rename(columns={'life_full_ratio': 'PCT'})

[77]: # median full sq
sa_med_fullsq = sa_grps['full_sq'].median().reset_index().rename(columns = ↪
    ↪{'full_sq': 'median'}) # cause median is robust to outliers.

# Get the NAS of full sq
nas_full = interior['full_sq'].isnull()
# Get the NAS of life sq
nas_life = interior['life_sq'].isnull()

[78]: handleFullsqLifesq(interior, nas_full, nas_life, sa_med_fullsq, ↪
    ↪avg_life_full_prop)
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/2830605670.py:2
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/2830605670.py:3
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

[docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://docs.stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[78]: {1: 'Done', 'Full SQ': 0, 'Life SQ': 0, 'Life SQ >= Full SQ': False}
```

```
[79]: interior = interior.drop(columns = 'life_full_ratio')
```

```
[80]: HandleFloorMaxFloor(interior, sa_grps)
```

```
[80]: {1: 'Done', 'Max Floor': 0, 'Floor': 0, 'Floor > Max Floor': False}
```

```
[81]: HandleNumRooms(interior, sa_grps)
```

```
[81]: {1: 'Done', 'Num Room': 0, 'LifeSQ/NumRoom < 5': False}
```

```
[82]: # Sanity check:
      interior.isna().sum()
```

```
[82]: id                0
      timestamp         0
      full_sq           0
      life_sq           0
      floor             0
      max_floor         0
      material          9572
      build_year        16116
      num_room          0
      kitch_sq          18118
      state             14253
      product_type       0
      sub_area           0
      price_doc          7662
      missing_life       0
      missing_full       0
      missing_floor      0
      missing_maxfloor   0
      missing_num_room   0
      dtype: int64
```

```
[83]: HandleBuildYear(interior)
```

```
[83]: {1: 'Done', 'Build_Year': 0}
```

```
[84]: handleStateMaterial(interior)
```

```
[84]: {1: 'Done', 'State': 0, 'Material': 0}
```

```
[85]: interior.drop(columns = ['year', 'age'], inplace=True) # Drop cols.
```

```
[86]: interior.isna().sum()
```

```
[86]: id                0
      timestamp        0
      full_sq          0
      life_sq           0
      floor            0
      max_floor        0
      material         0
      build_year       0
      num_room         0
      kitch_sq        18118
      state           0
      product_type     0
      sub_area         0
      price_doc        7662
      missing_life     0
      missing_full     0
      missing_floor    0
      missing_maxfloor 0
      missing_num_room 0
      missing_build_year 0
      missing_state    0
      missing_material 0
      dtype: int64
```

```
[87]: # Fill missing values in kitch_sq column

      # Get the kitch_sq/full sq
      interior['kitch_life_ratio'] = interior['kitch_sq']/interior['life_sq']
      interior['kitch_life_ratio'] = interior['kitch_life_ratio'].
      ↪ fillna(interior['kitch_life_ratio'].mean()) # Fill missing values with the
      ↪ mean.

      # Get the avg kitch_sq/full sq for each sub area and product type.
      sa_avg_kitch_life = interior.groupby(['sub_area'])['kitch_life_ratio'].mean().
      ↪ reset_index().rename(columns={'kitch_life_ratio': 'PCT'})
```

```
[88]: HandleKitchSQ(interior, sa_avg_kitch_life)
```

```
[88]: {1: 'Done',
      'Kitch SQ': 0,
      'Kitch SQ >= Full SQ': False,
      'Kitch SQ >= Life SQ': False}
```

```
[89]: groups_dfs['interior'] = interior.drop(columns = ['kitch_life_ratio'])
```

## Distances

```
[90]: # Get distances data frame
distances = groups_dfs['distances']
distances_na = distances.isnull().sum()
# Get the features that have missing values.
distances_na[distances_na > 0]
```

```
[90]: metro_min_walk      59
metro_km_walk          59
railroad_station_walk_km  59
railroad_station_walk_min 59
ID_railroad_station_walk 59
dtype: int64
```

```
[91]: distances['sub_area'] = full_data['sub_area']
for c in distances_na[distances_na > 0].index:
    if 'ID' not in c:
        distances[c].fillna(distances.groupby('sub_area')[c].
↳ transform('median'), inplace=True)
    else:
        distances[c].fillna(distances.groupby('sub_area')[c].transform(lambda x:
↳ x.mode()[0]), inplace=True)
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:1
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:4
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:4
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:4  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:4  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/3585862785.py:6  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[92]: distances.drop(columns='sub_area', inplace=True)
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_66217/2186917241.py:1  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[93]: groups_dfs['distances'] = distances
```

### Surroundings

```
[94]: # Get surroundings data frame  
surroundings = groups_dfs['surroundings']  
surr_na = surroundings.isnull().sum()
```

```
# Get the features that have missing values.
surr_na[surr_na > 0]
```

```
[94]: cafe_sum_1000_min_price_avg    7746
      cafe_sum_1000_max_price_avg    7746
      cafe_avg_price_1000           7746
      cafe_sum_1500_min_price_avg    5020
      cafe_sum_1500_max_price_avg    5020
      cafe_avg_price_1500           5020
      green_part_2000                19
      cafe_sum_2000_min_price_avg    2149
      cafe_sum_2000_max_price_avg    2149
      cafe_avg_price_2000           2149
      cafe_sum_3000_min_price_avg    1173
      cafe_sum_3000_max_price_avg    1173
      cafe_avg_price_3000           1173
      prom_part_5000                 270
      cafe_sum_5000_min_price_avg     425
      cafe_sum_5000_max_price_avg     425
      cafe_avg_price_5000            425
      dtype: int64
```

To address these missing values, we decided to use KNNImputer since each feature here represent an avg price of a shop within a certain distance. However, before doing that we need to check the variability of each feature which is the STD/Mean.

```
[95]: # Get the columns
      surr_na_cols = surr_na[surr_na > 0].index
      # Check the std/mean ratio
      surroundings[surr_na_cols].std()/surroundings[surr_na_cols].mean()
```

```
[95]: cafe_sum_1000_min_price_avg    0.318793
      cafe_sum_1000_max_price_avg    0.288062
      cafe_avg_price_1000           0.298280
      cafe_sum_1500_min_price_avg    0.271707
      cafe_sum_1500_max_price_avg    0.245080
      cafe_avg_price_1500           0.253864
      green_part_2000                0.676771
      cafe_sum_2000_min_price_avg    0.275928
      cafe_sum_2000_max_price_avg    0.250875
      cafe_avg_price_2000           0.259165
      cafe_sum_3000_min_price_avg    0.286584
      cafe_sum_3000_max_price_avg    0.269686
      cafe_avg_price_3000           0.275521
      prom_part_5000                 0.549153
      cafe_sum_5000_min_price_avg    0.196017
      cafe_sum_5000_max_price_avg    0.182153
      cafe_avg_price_5000            0.187165
```



dtype: float64

```
[96]: # Create an imputer
imputer = KNNImputer(n_neighbors=10)

# Use imputer to fill missing values.
surroundings.loc[:,surr_na_cols] = imputer.
↳fit_transform(surroundings[surr_na_cols])
```

```
[97]: # Check if knn imputer worked
np.any(surroundings.isna().sum() > 0)
```

[97]: False

```
[98]: # Update surroundings data frame
groups_dfs['surroundings'] = surroundings
```

### Education

```
[99]: education = groups_dfs['education']
education.isna().sum()
```

```
[99]: preschool_quota                8284
preschool_education_centers_raion    0
school_quota                        8280
school_education_centers_raion        0
school_education_centers_top_20_raion 0
university_top_20_raion              0
sport_objects_raion                  0
additional_education_raion           0
dtype: int64
```

each of the columns with missing values give an information about the number of seats.. so we decided to fill them with -1 to indicate there's no data.

```
[100]: education.loc[:, 'preschool_quota'] = education['preschool_quota'].fillna(-1)
education.loc[:, 'school_quota'] = education['school_quota'].fillna(-1)
```

```
[101]: groups_dfs['education'] = education
```

### Facilities

```
[102]: faci = groups_dfs['facilities']
faci.isna().sum()
```

```
[102]: hospital_beds_raion            17859
healthcare_centers_raion           0
shopping_centers_raion              0
office_raion                       0
thermal_power_plant_raion           0
```

```

incineration_raion      0
oil_chemistry_raion     0
radiation_raion         0
railroad_terminal_raion 0
big_market_raion        0
nuclear_reactor_raion   0
detention_facility_raion 0
dtype: int64

```

```
[103]: # Again a feature that indicates a count... can be fillied with -1 - no data.
```

```

faci.loc[:, 'hospital_beds_raion'] = faci['hospital_beds_raion'].fillna(-1)

```

```
[104]: groups_dfs['facilities'] = faci
```

### Cleaning the data

```
[105]: # Check if we missed any missing value:
```

```

group_na_indicators = []

```

```

for group in groups_dfs:
    indicator = np.any(groups_dfs.get(group).isna().sum() > 0)
    group_na_indicators.append((group, indicator))

```

```

group_na_indicators

```

```
[105]: [('areas', False),
        ('buildings', False),
        ('demographics', False),
        ('distances', False),
        ('education', False),
        ('facilities', False),
        ('interior', True),
        ('surroundings', False)]

```

```
[106]: nas = full_data.isna().sum()
```

```

# Change np.nan values into -1 so they could be replaced with the real value_
↳from groups dfs

```

```

full_data.loc[:, nas[nas > 0].index] = train.loc[:, nas[nas > 0].index].
↳replace(np.nan, -999)

```

```

for group in groups_dfs:
    df = groups_dfs.get(group)
    # Insert clean columns to df.
    full_data.loc[df.index, df.columns] = df

```

```
[107]: full_data = full_data.replace(-999,np.nan)
```

```
[108]: # get columns with missing values - supposed to be the distances columns only
x = full_data.isna().sum()
print(x[x>0])
```

```
price_doc      7662
dtype: int64
```

```
[109]: # full_data.to_csv("~/desktop/cleand_data.csv", index=False)
```

### 4.3 Loading Cleaned Data Frame

```
[4]: import requests

# Paths to text files from github
paths = ['https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/areas.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/buildings.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/demographics.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/distances.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/education.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/facilities.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/interior.txt',
         'https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/main/
↳text_files/surroundings.txt']

# Groups
groups = {
    'areas': [],
    'buildings': [],
    'demographics': [],
    'distances': [],
    'education': [],
    'facilities': [],
    'interior': [],
    'surroundings': []
}

keys = list(groups.keys())

for i in range(len(paths)):
    # Http request to get the file data
```

```

response = requests.get(paths[i])
# read File lines
lines = response.text.splitlines()
for line in lines:
    # Split by \t
    col = line.split('\t')[0]
    # Interior file has : in it.
    if keys[i] == 'interior':
        col = col.split(':')[0]
    groups[keys[i]].append(col)

```

```

[5]: full_data = pd.read_csv("https://raw.githubusercontent.com/LidorErez98/
↳Sberbank_ML/main/cleand_data.csv")

```

```

[6]: sanity = full_data.isna().sum()
sanity[sanity > 0]

```

```

[6]: price_doc    7662
dtype: int64

```

```

[7]: groups_dfs = {group:full_data[full_data.columns[np.where(full_data.columns.
↳isin(groups[group]))]] for group in groups}

```

## 4.4 Feature Engineering

```

[8]: def create_time_features(data):
    data['year'] = pd.to_datetime(data['timestamp']).dt.year
    data['month'] = pd.to_datetime(data['timestamp']).dt.month
    data['day'] = pd.to_datetime(data['timestamp']).dt.day
    data['day_of_week'] = pd.to_datetime(data['timestamp']).dt.dayofweek
    data['week_of_year'] = pd.to_datetime(data['timestamp']).dt.days_in_month
    data['quarter'] = pd.to_datetime(data['timestamp']).dt.quarter

    # Part 5
    data['monthyear'] = data['year']*100 + data['month']
    data['weekyear'] = data['year']*100 + data['week_of_year']

    month_year_counts = data['monthyear'].value_counts().to_dict()
    data['monthyear_count'] = data['monthyear'].map(month_year_counts)

    week_year_counts = data['weekyear'].value_counts().to_dict()
    data['weekyear_count'] = data['weekyear'].map(week_year_counts)

[9]: full_data['timestamp'] = pd.to_datetime(full_data['timestamp'])
create_time_features(full_data)
full_data['log_price'] = np.log(full_data['price_doc'] + 1)
full_data['building_age'] = full_data['year'] - full_data['build_year']

```

```
full_data['price_sq'] = full_data['price_doc']/full_data['full_sq']
```

```
[10]: dens_columns = [col + '_dens' for col in groups_dfs['demographics'].columns]
```

```
[11]: full_data[dens_columns] = full_data[groups_dfs['demographics'].columns].  
      ↪apply(lambda x:x/full_data['area_m'])
```

```
[12]: groups_dfs['demographics'].loc[:,dens_columns] = full_data[dens_columns]
```

```
[13]: full_data["ratio_life_sq_full_sq"] = full_data["life_sq"] / full_data["full_sq"]
```

```
[14]: # adding ratio_life_sq_full_sq to interior group  
      groups_dfs['interior'].loc[:, 'ratio_life_sq_full_sq'] =  
      ↪full_data['ratio_life_sq_full_sq']
```

/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/198719290.py:2:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[15]: #kitchen ratio from the life sq  
      full_data["ratio_kitch_sq_full_sq"] = full_data["kitch_sq"] /  
      ↪full_data["full_sq"]
```

```
[16]: # adding ratio_kitch_sq_full_sq to interior group  
      groups_dfs['interior'].loc[:, 'ratio_kitch_sq_full_sq'] =  
      ↪full_data['ratio_kitch_sq_full_sq']
```

/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/975391471.py:2:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[17]: # calculate the room size  
      full_data["room_size"] = (full_data["life_sq"] - full_data['kitch_sq']) /  
      ↪full_data["num_room"]
```

```
[18]: # adding room size to interior group
groups_dfs['interior'].loc[:, 'room_size'] = full_data['room_size']
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/4252726323.py:2
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[19]: full_data['extra_sq'] = full_data['full_sq'] - full_data['life_sq']
```

```
[20]: full_data['area_diff'] = full_data['full_sq'] - full_data['kitch_sq'] # Part 5
```

```
[21]: # adding extra sq to interior group
groups_dfs['interior'].loc[:, 'extra_sq'] = full_data['extra_sq']
groups_dfs['interior'].loc[:, 'area_diff'] = full_data['area_diff']
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/2361000648.py:2
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/2361000648.py:3
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[22]: # calculate the floor ratio
full_data['floor_ratio'] = full_data['floor'] / full_data['max_floor']
```

```
[23]: # adding floor ratio to interior group
groups_dfs['interior'].loc[:, 'floor_ratio'] = full_data['floor_ratio']
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3316551521.py:2
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[24]: # num of floor from top
full_data['floor_from_top'] = full_data['max_floor'] - full_data['floor']
```

```
[25]: # adding floor from top to interior group
groups_dfs['interior'].loc[:, 'floor_from_top'] = full_data['floor_from_top']
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/757007344.py:2:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
[26]: percapita = [c + '_percapita' for c in groups_dfs['education'].columns]
full_data[percapita] = full_data[groups_dfs['education'].columns].apply(lambda edu:
    edu/full_data['raion_popul'])
```

```
[27]: groups_dfs['education'].loc[:,percapita] = full_data[percapita]
```

```
[28]: minmaxscaler = MinMaxScaler()
cafe_count = full_data.filter(like="cafe_count")
minmaxscaler.fit_transform(cafe_count)
# replace with cafe count columns in full_data
full_data.loc[:,cafe_count.columns] = minmaxscaler.fit_transform(cafe_count)
```

```
[29]: groups_dfs['surroundings'].loc[:,groups_dfs['surroundings'].
    filter(like="cafe_count").columns] = full_data.filter(like="cafe_count")
```

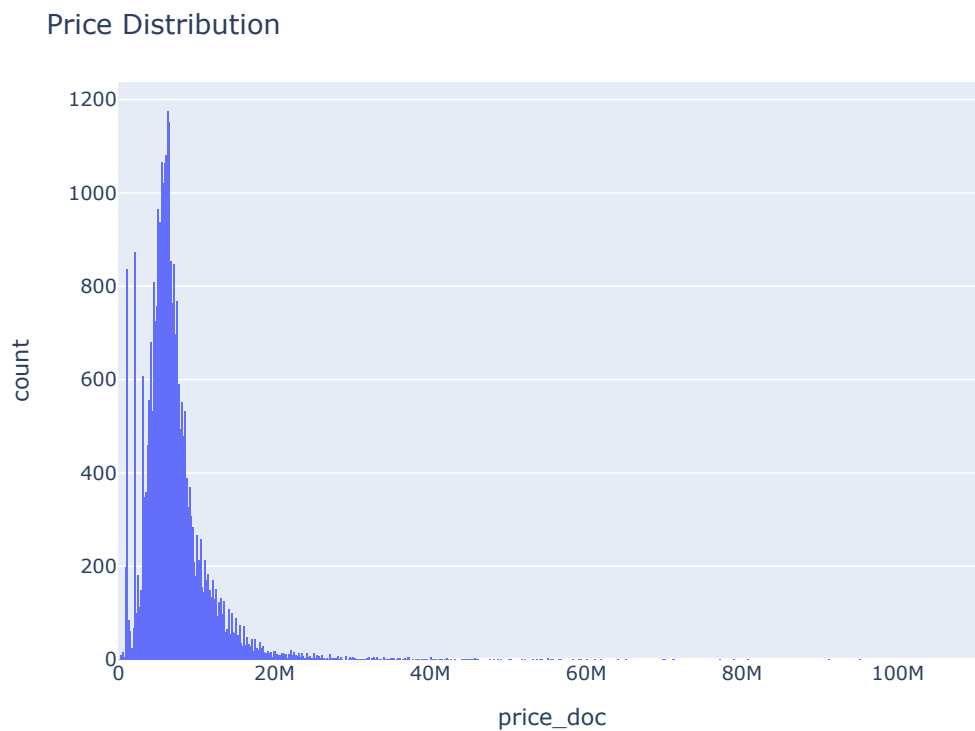
```
[30]: full_data.drop(columns = [c.replace('_dens','') for c in dens_columns],
    inplace=True)
full_data.drop(columns = [c.replace('_percapita','') for c in percapita],
    inplace=True)
```

```
groups_dfs['demographics'] = groups_dfs['demographics'].drop(columns = [c.
    ↳replace('_dens','') for c in dens_columns])
groups_dfs['education'] = groups_dfs['education'].drop(columns = [c.
    ↳replace('_percapita','') for c in percapita])
```

## 4.5 Explanatory Data Analysis

```
[31]: import plotly.express as px
```


```
[32]: fig = px.histogram(full_data, x='price_doc')
fig.update_layout(title = 'Price Distribution')
fig.show()
```



```
[33]: fig = px.histogram(full_data, x='log_price')
fig.update_layout(title = 'Log Price Distribution')
fig.show()
```



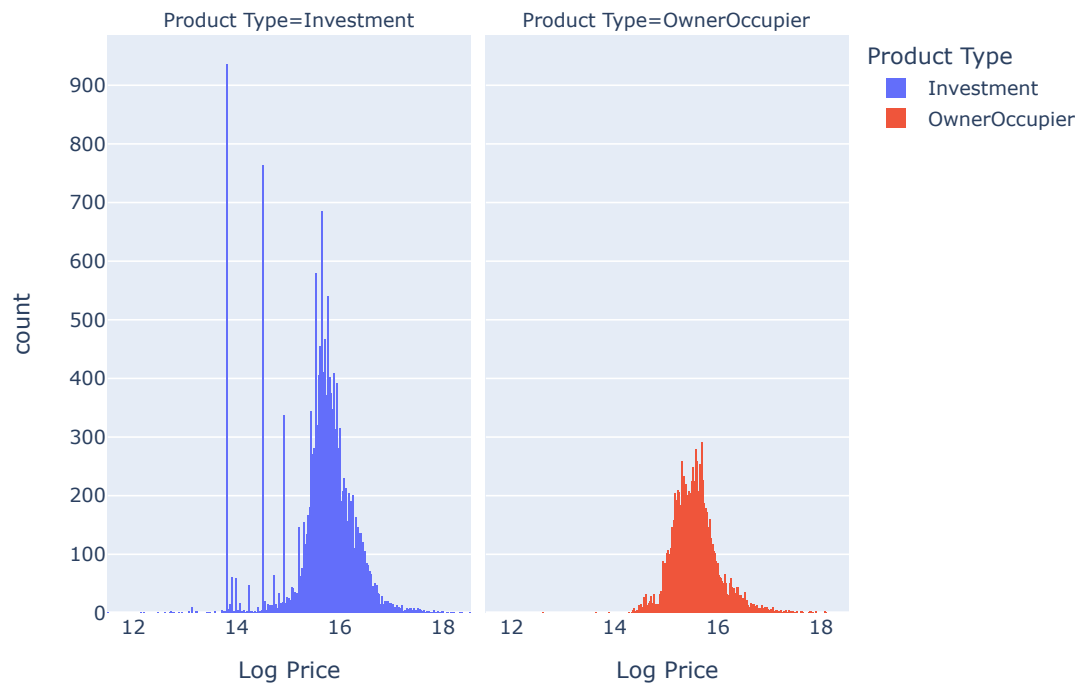
Log Price Distribution



Log price might be a better target variable than the actual price.


```
[34]: fig = px.histogram(full_data, x='log_price',  
    ↪ facet_col='product_type', labels={'log_price': 'Log Price', 'product_type':  
    ↪ 'Product Type'}, color='product_type')  
fig.update_layout(title = 'Density Plot of Log Price by Product Type')  
fig.show()
```

Density Plot of Log Price by Product Type




```
[35]: fig = px.box(full_data, x='log_price', y='product_type', labels={'log_price':
    ↳ 'Log Price', 'product_type': 'Product Type'}, color='product_type')
fig.update_layout(title = 'Density Plot of Log Price by Product Type')
fig.show()
```

Density Plot of Log Price by Product Type




There are many outliers in the log price column this might have an effect on our model predictions.

```
[36]: cp = full_data.copy()
cp['year'] = cp['timestamp'].dt.year
temp = (cp.groupby(['year', 'product_type'])['log_price'].mean()).reset_index()
temp['year'] = temp['year'].apply(lambda x: str(x))
year_order = sorted(temp['year'].unique())
fig = px.line(temp, x='year', y='log_price', facet_col = 'product_type',
              category_orders={'year': year_order}, color='product_type')
fig.update_layout(title = 'Log Price over the year by Product Type')
fig.show()
```




Prices different behavior among product types indicate that we might need to train separated models for each product type.


```
[37]: temp = (full_data.groupby("build_year")['price_doc'].mean()).reset_index()
fig = px.line(temp, x='build_year',y='price_doc')
fig.show()
```



```
[38]: temp = full_data.loc[full_data.build_year <= 1960]
      fig = px.histogram(temp['price_doc'])
      fig.show()
```



```
[39]: temp = (full_data['build_year'].value_counts()).reset_index()
fig = px.histogram(temp,x='build_year',y='count')
fig.show()
```



We found that for old building prices are higher when the build year is less than 1960 so we will add a binary feature to describe if a house is “vintage”

```
[40]: full_data.loc[:, 'is_vintage'] = full_data['build_year'].apply(lambda year:
    ↪(year <= 1960)).astype(int)
```

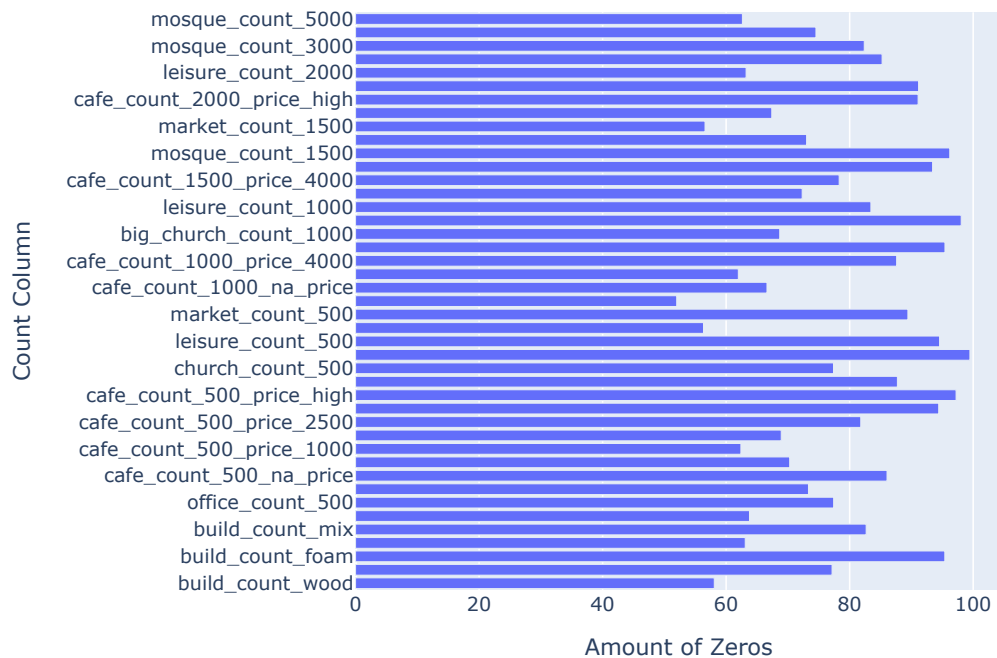
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/2529656510.py:1  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
[41]: # Checking for sparsity in count data:

zeroes_pct = np.round((full_data.filter(like='count') == 0).mean()*100,2).
    ↪reset_index().rename(columns = {0:'Amount of Zeros', 'index':'Count Column'})

px.bar(zeroes_pct[zeroes_pct['Amount of Zeros'] >= 50], x='Amount of
    ↪Zeros', y='Count Column')
```



We checked the amount of zeros in our count columns to figure out how sparse they are. We then created a bar plot of the count columns that have at least 50% zeros. To cope with this issue we decided to turn each of the count column that is shown on the plot to a binary column which will indicate the absence of data

```
[42]: sparse_count_columns = zeroes_pct[zeroes_pct['Amount of Zeros'] >= 50]['Count_
↳Column'].values
sparse_count_columns
```

```
[42]: array(['build_count_wood', 'build_count_frame', 'build_count_foam',
'build_count_slag', 'build_count_mix', 'build_count_before_1920',
'office_count_500', 'trc_count_500', 'cafe_count_500_na_price',
'cafe_count_500_price_500', 'cafe_count_500_price_1000',
'cafe_count_500_price_1500', 'cafe_count_500_price_2500',
'cafe_count_500_price_4000', 'cafe_count_500_price_high',
'big_church_count_500', 'church_count_500', 'mosque_count_500',
'leisure_count_500', 'sport_count_500', 'market_count_500',
'office_count_1000', 'cafe_count_1000_na_price',
'cafe_count_1000_price_2500', 'cafe_count_1000_price_4000',
'cafe_count_1000_price_high', 'big_church_count_1000',
'mosque_count_1000', 'leisure_count_1000', 'market_count_1000',
'cafe_count_1500_price_4000', 'cafe_count_1500_price_high',
```



```
'mosque_count_1500', 'leisure_count_1500', 'market_count_1500',
'cafe_count_2000_price_4000', 'cafe_count_2000_price_high',
'mosque_count_2000', 'leisure_count_2000',
'cafe_count_3000_price_high', 'mosque_count_3000',
'cafe_count_5000_price_high', 'mosque_count_5000'], dtype=object)
```

```
[43]: for col in sparse_count_columns:
        full_data.loc[:, col + '_binary'] = ((full_data[col] > 0)).astype(int)

full_data.filter(like='_binary')
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/3714908506.py:2  
: PerformanceWarning:

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all

columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
/var/folders/2y/5v1st1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/3714908506.py:2
: PerformanceWarning:
```

DataFrame is highly fragmented. This is usually the result of calling ``frame.insert`` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use ``newframe = frame.copy()``

```
[43]:
```

	build_count_wood_binary	build_count_frame_binary \
0	0	0
1	1	0
2	0	0
3	1	1
4	0	0
...	...	...
38128	0	0
38129	0	0
38130	1	1
38131	0	0
38132	0	0

	build_count_foam_binary	build_count_slag_binary \
0	0	0
1	0	0
2	0	1
3	0	1
4	0	1
...	...	...
38128	0	0
38129	0	0
38130	0	1
38131	0	0
38132	0	0

	build_count_mix_binary	build_count_before_1920_binary \
--	------------------------	----------------------------------



0	0	0
1	0	1
2	0	1
3	1	1
4	1	1
...	...	...
38128	0	0
38129	0	0
38130	1	1
38131	0	0
38132	0	0

	office_count_500_binary	trc_count_500_binary \
0	0	0
1	0	0
2	0	0
3	0	0
4	1	1
...	...	...
38128	1	1
38129	0	0
38130	1	1
38131	0	1
38132	1	1

	cafe_count_500_na_price_binary	cafe_count_500_price_500_binary ... \
0	0	0 ...
1	0	1 ...
2	0	0 ...
3	0	0 ...
4	1	1 ...
...	...	...
38128	0	1 ...
38129	0	0 ...
38130	1	1 ...
38131	0	1 ...
38132	0	0 ...

	leisure_count_1500_binary	market_count_1500_binary \
0	0	1
1	1	0
2	0	1
3	0	1
4	1	1
...	...	...
38128	1	1
38129	0	0

38130	1	1
38131	0	1
38132	0	0

	cafe_count_2000_price_4000_binary	cafe_count_2000_price_high_binary \
0	1	0
1	1	0
2	0	0
3	0	0
4	1	1
...	...	...
38128	0	0
38129	0	0
38130	1	1
38131	1	0
38132	0	0

	mosque_count_2000_binary	leisure_count_2000_binary \
0	0	0
1	0	1
2	0	0
3	0	0
4	0	1
...	...	...
38128	0	1
38129	0	0
38130	1	1
38131	0	0
38132	1	0

	cafe_count_3000_price_high_binary	mosque_count_3000_binary \
0	0	0
1	0	0
2	0	0
3	0	0
4	1	1
...	...	...
38128	0	1
38129	0	0
38130	1	1
38131	0	0
38132	0	1

	cafe_count_5000_price_high_binary	mosque_count_5000_binary
0	0	1
1	0	1
2	0	0

3	1	0
4	1	1
...	...	...
38128	1	1
38129	0	0
38130	1	1
38131	0	0
38132	0	1

[38133 rows x 43 columns]

```
[44]: full_data = full_data.drop(columns = sparse_count_columns)
```

```
[45]: for group in groups_dfs:
        if np.any(groups_dfs[group].columns.isin(sparse_count_columns)):
            idx = np.where(groups_dfs[group].columns.isin(sparse_count_columns))
            cols_to_drop = groups_dfs[group].columns[idx]
            groups_dfs[group] = groups_dfs[group].drop(columns = cols_to_drop)
```

## 4.6 Macro CSV

```
[46]: macro = pd.read_csv("https://raw.githubusercontent.com/LidorErez98/Sberbank_ML/
    ↪main/Data/macro.csv")
min_date, max_date = full_data['timestamp'].min(), full_data['timestamp'].max()
macro['date'] = pd.to_datetime(macro['timestamp'])
macro = macro[(macro['date'] >= min_date) & (macro['date'] <= max_date)]
macro['year'] = macro['date'].dt.year
macro['month'] = macro['date'].dt.month
macro['quarter'] = macro['date'].dt.quarter
```

```
[47]: # split to sub dataframes
inflation = macro[['date', 'quarter', 'year', 'month', 'cpi', 'ppi',
    ↪'gdp_deflator']]
gdp = macro[['date', 'quarter', 'year', 'month', 'gdp_quart',
    ↪'gdp_quart_growth', 'gdp_deflator', 'gdp_annual', 'gdp_annual_growth']]
# salary = macro[['date', 'quarter', 'year', 'month', 'salary',
    ↪'salary_growth', 'real_dispos_income_per_cap_growth', 'salary_growth']]
# mortgage = macro[['date', 'quarter', 'year', 'month', 'mortgage_rate',
    ↪'mortgage_growth', 'deposits_rate', 'deposits_growth']]
# investmeent = macro[['date', 'quarter', 'year', 'month',
    ↪'invest_fixed_assets', 'invest_fixed_assets_phys',
    ↪'profitable_enterpr_share', 'unprofitable_enterpr_share',
    ↪'share_own_revenues', 'overdue_wages_per_cap', 'fin_res_per_cap',
    ↪'invest_fixed_assets']]
```


```
# consumption = macro[['date', 'quarter', 'year', 'month', 'income_per_cap',
↳ 'real_dispos_income_per_cap_growth', 'salary', 'salary_growth',
↳ 'retail_trade_turnover', 'retail_trade_turnover_per_cap',
↳ 'retail_trade_turnover_growth', 'labor_force']]
# interest = macro[['date', 'quarter', 'year', 'month', 'deposits_rate',
↳ 'deposits_growth', 'mortgage_rate', 'mortgage_growth']]
# governmenr = macro[['date', 'quarter', 'year', 'month', 'balance_trade',
↳ 'balance_trade_growth', 'usdrub', 'eurrrub', 'micex_rgbi_tr', 'micex']]
```

#### 4.6.1 INFLATION

```
[48]: inflation_df = inflation.copy()
      # using plotly express to plot the data over time
      def deflator_to_inflation_rate(deflator1, defltor0):
          return (deflator1 - defltor0)/defltor0

[49]: fig = px.line(inflation_df, x='date', y='gdp_deflator', title='GDP deflator_
↳ over time')
      fig.show()
      print(inflation_df['gdp_deflator'].describe())
```

GDP deflator over time



count 1746.000000

```

mean      111.344730
std       12.220722
min       86.721000
25%      100.000000
50%      113.465000
75%      123.661000
max       133.160000
Name: gdp_deflator, dtype: float64

```

### Inflation Integrity

```

[50]: # fix the gdp deflator of 2011 to 100 and the rest of them accordingly
fix = inflation_df[inflation_df['year'] == 2011]['gdp_deflator'].values[0] - 100
inflation_df['gdp_deflator'] = inflation_df['gdp_deflator'] - fix
inflation_df['inflation_from_8_2011'] = inflation_df['gdp_deflator'].
    ↪ apply(lambda x: deflator_to_inflation_rate(x, 100))

```

```

[51]: # for annual inflation growth rate we will use the inflation from 8/2011
annual_inflation_df = pd.DataFrame(inflation_df[['year', 'gdp_deflator',
    ↪ 'inflation_from_8_2011']].drop_duplicates(subset=['year']))
annual_inflation_df.set_index('year', inplace=True)
for i in range(2012, 2017):
    annual_inflation_df.loc[i, 'inflation_growth'] = annual_inflation_df.loc[i,
    ↪ 'inflation_from_8_2011'] - annual_inflation_df.loc[i-1,
    ↪ 'inflation_from_8_2011']
annual_inflation_df.loc[2011, 'inflation_growth'] = 0
annual_inflation_df

```

```

[51]:      gdp_deflator  inflation_from_8_2011  inflation_growth
year
2011      100.000          0.00000          0.00000
2012      113.279          0.13279          0.13279
2013      121.578          0.21578          0.08299
2014      126.744          0.26744          0.05166
2015      136.940          0.36940          0.10196
2016      146.439          0.46439          0.09499


```

```

[52]: # annual inflation growth rate vs year ployline
fig = px.line(annual_inflation_df, x=annual_inflation_df.index,
    ↪ y='inflation_growth', title='Annual inflation growth rate vs year')
fig.show()

```

Annual inflation growth rate vs year



```
[53]: # add the inflation growth rate to the inflation dataframe
inflation_df.set_index('year', inplace=True)
inflation_df['annual_inflation_growth'] =
    ↪annual_inflation_df['inflation_growth']
inflation_df.reset_index(inplace=True)
inflation_df
```

```
[53]:
```

	year	date	quarter	month	cpi	ppi	gdp_deflator	\
0	2011	2011-08-20	3	8	354.0	420.7	100.000	
1	2011	2011-08-21	3	8	354.0	420.7	100.000	
2	2011	2011-08-22	3	8	354.0	420.7	100.000	
3	2011	2011-08-23	3	8	354.0	420.7	100.000	
4	2011	2011-08-24	3	8	354.0	420.7	100.000	
...	...	...	...	...	...	...	...	
1741	2016	2016-05-26	2	5	523.2	584.0	146.439	
1742	2016	2016-05-27	2	5	523.2	584.0	146.439	
1743	2016	2016-05-28	2	5	523.2	584.0	146.439	
1744	2016	2016-05-29	2	5	523.2	584.0	146.439	
1745	2016	2016-05-30	2	5	523.2	584.0	146.439	

```
inflation_from_8_2011  annual_inflation_growth
```


0	0.00000	0.00000
1	0.00000	0.00000
2	0.00000	0.00000
3	0.00000	0.00000
4	0.00000	0.00000
...	...	...
1741	0.46439	0.09499
1742	0.46439	0.09499
1743	0.46439	0.09499
1744	0.46439	0.09499
1745	0.46439	0.09499

[1746 rows x 9 columns]

### cpi and ppi correlation with inflation rate

```
[54]: # plot 3 lines together
fig = px.line(inflation_df, x='date', y=['cpi', 'ppi', 'inflation_from_8_2011'],
              title='Inflation rates over time',
              color='variable')
fig.show()
```

Inflation rates over time




```
[55]: inflation_df['cpi'].values[0]
```

```
[55]: 354.0
```

```
[56]: # turn cpi and ppi to growths from 8 2011
cpifix = 100 - inflation_df['cpi'].values[0]
ppifix = 100 - inflation_df['ppi'].values[0]
inflation_df['cpi'] = inflation_df['cpi'] + cpifix
inflation_df['ppi'] = inflation_df['ppi'] + ppifix
inflation_df['cpi_growth'] = inflation_df['cpi'].apply(lambda x:
    ↪deflator_to_inflation_rate(x, 100))
inflation_df['ppi_growth'] = inflation_df['ppi'].apply(lambda x:
    ↪deflator_to_inflation_rate(x, 100))
```


```
[57]: gdp_dfd = gdp.copy(deep=True)
# check the gdp growth over time( annually vs quarterly) plotly express
fig = px.line(gdp_dfd, x='date', y='gdp_annual_growth', title='GDP annual_
    ↪growth over time')
fig.show()
fig = px.line(gdp_dfd, x='date', y='gdp_quart_growth', title='GDP quarterly_
    ↪growth over time')
fig.show()
```

GDP annual growth over time





GDP quarterly growth over time



```
[58]: # create a df with the quart gdp growth for each quarter in each year no
      ↳ duplicates
gdp_dfd_quart = pd.DataFrame(gdp_dfd[['year', 'quarter', 'gdp_quart_growth']].
      ↳ drop_duplicates())
# setting year, quarter as index, and divide by 100 to get the percentage
gdp_dfd_quart['gdp_quart_growth'] = gdp_dfd_quart['gdp_quart_growth'].
      ↳ apply(lambda x: x/100)
gdp_dfd_quart.set_index(['year', 'quarter'], inplace=True)
gdp_dfd_quart['gdp_quart_growth_since_2011'] =
      ↳ gdp_dfd_quart['gdp_quart_growth'].cumsum()
gdp_dfd_quart
```

```
[58]:
```

		gdp_quart_growth	gdp_quart_growth_since_2011
year	quarter		
2011	3	0.033	0.033
	4	0.050	0.083
2012	1	0.052	0.135
	2	0.047	0.182
	3	0.042	0.224
	4	0.031	0.255
2013	1	0.020	0.275

	2	0.007	0.282
	3	0.012	0.294
	4	0.013	0.307
2014	1	0.021	0.328
	2	0.006	0.334
	3	0.007	0.341
	4	0.009	0.350
2015	1	0.004	0.354
	2	-0.028	0.326
	3	-0.045	0.281
	4	-0.037	0.244
2016	1	-0.038	0.206
	2	-0.012	0.194

```
[59]: # adding the gdp growth sum to relevant rows in the gdp df
gdp_dfd.set_index(['year', 'quarter'], inplace=True)
gdp_dfd['gdp_quart_growth_since_2011'] =
↳ gdp_dfd_quart['gdp_quart_growth_since_2011']
gdp_dfd.reset_index(inplace=True)
gdp_dfd.head(17)
```

```
[59]:
```

	year	quarter	date	month	gdp_quart	gdp_quart_growth \
0	2011	3	2011-08-20	8	14313.7	3.3
1	2011	3	2011-08-21	8	14313.7	3.3
2	2011	3	2011-08-22	8	14313.7	3.3
3	2011	3	2011-08-23	8	14313.7	3.3
4	2011	3	2011-08-24	8	14313.7	3.3
5	2011	3	2011-08-25	8	14313.7	3.3
6	2011	3	2011-08-26	8	14313.7	3.3
7	2011	3	2011-08-27	8	14313.7	3.3
8	2011	3	2011-08-28	8	14313.7	3.3
9	2011	3	2011-08-29	8	14313.7	3.3
10	2011	3	2011-08-30	8	14313.7	3.3
11	2011	3	2011-08-31	8	14313.7	3.3
12	2011	3	2011-09-01	9	14313.7	3.3
13	2011	3	2011-09-02	9	14313.7	3.3
14	2011	3	2011-09-03	9	14313.7	3.3
15	2011	3	2011-09-04	9	14313.7	3.3
16	2011	3	2011-09-05	9	14313.7	3.3

	gdp_deflator	gdp_annual	gdp_annual_growth	gdp_quart_growth_since_2011
0	86.721	46308.5	0.045037	0.033
1	86.721	46308.5	0.045037	0.033
2	86.721	46308.5	0.045037	0.033
3	86.721	46308.5	0.045037	0.033
4	86.721	46308.5	0.045037	0.033
5	86.721	46308.5	0.045037	0.033

6	86.721	46308.5	0.045037	0.033
7	86.721	46308.5	0.045037	0.033
8	86.721	46308.5	0.045037	0.033
9	86.721	46308.5	0.045037	0.033
10	86.721	46308.5	0.045037	0.033
11	86.721	46308.5	0.045037	0.033
12	86.721	46308.5	0.045037	0.033
13	86.721	46308.5	0.045037	0.033
14	86.721	46308.5	0.045037	0.033
15	86.721	46308.5	0.045037	0.033
16	86.721	46308.5	0.045037	0.033

```
[60]: gdp_dfd_year = pd.DataFrame(gdp_dfd[['year', 'gdp_annual_growth']].
↳ drop_duplicates())
gdp_dfd_year.set_index('year', inplace=True)
gdp_dfd_year['gdp_annual_growth_since_2011'] =
↳ gdp_dfd_year['gdp_annual_growth'].cumsum()
gdp_dfd_year
```

```
[60]:      gdp_annual_growth  gdp_annual_growth_since_2011
year
2011          0.045037          0.045037
2012          0.042644          0.087681
2013          0.035179          0.122859
2014          0.012795          0.135654
2015          0.007065          0.142719
2016         -0.037267          0.105452
```

```
[61]: # insert the new feature to gdp_dfd
gdp_dfd.set_index('year', inplace=True)
gdp_dfd['gdp_annual_growth_since_2011'] =
↳ gdp_dfd_year['gdp_annual_growth_since_2011']
gdp_dfd.reset_index(inplace=True)
gdp_dfd['gdp_quart_growth'] = gdp_dfd['gdp_quart_growth'].apply(lambda x: x/100)
```

```
[62]: #adjusting the gdp growth rates so they will start from 0 in 2011
gdp_dfd['gdp_annual_growth'] = gdp_dfd['gdp_annual_growth'] -
↳ gdp_dfd[gdp_dfd['year'] == 2011]['gdp_annual_growth'].values[0]
gdp_dfd['gdp_annual_growth_since_2011'] =
↳ gdp_dfd['gdp_annual_growth_since_2011'] - gdp_dfd[gdp_dfd['year'] ==
↳ 2011]['gdp_annual_growth_since_2011'].values[0]
gdp_dfd['gdp_quart_growth'] = gdp_dfd['gdp_quart_growth'] -
↳ gdp_dfd[gdp_dfd['year'] == 2011]['gdp_quart_growth'].values[0]
gdp_dfd['gdp_quart_growth_since_2011'] = gdp_dfd['gdp_quart_growth_since_2011'] -
↳ gdp_dfd[gdp_dfd['year'] == 2011]['gdp_quart_growth_since_2011'].values[0]
```

```
[63]: gdp_dfd.head(17)
```

```

[63]:
    year  quarter      date  month  gdp_quart  gdp_quart_growth  \
0  2011         3  2011-08-20     8    14313.7          0.0
1  2011         3  2011-08-21     8    14313.7          0.0
2  2011         3  2011-08-22     8    14313.7          0.0
3  2011         3  2011-08-23     8    14313.7          0.0
4  2011         3  2011-08-24     8    14313.7          0.0
5  2011         3  2011-08-25     8    14313.7          0.0
6  2011         3  2011-08-26     8    14313.7          0.0
7  2011         3  2011-08-27     8    14313.7          0.0
8  2011         3  2011-08-28     8    14313.7          0.0
9  2011         3  2011-08-29     8    14313.7          0.0
10 2011         3  2011-08-30     8    14313.7          0.0
11 2011         3  2011-08-31     8    14313.7          0.0
12 2011         3  2011-09-01     9    14313.7          0.0
13 2011         3  2011-09-02     9    14313.7          0.0
14 2011         3  2011-09-03     9    14313.7          0.0
15 2011         3  2011-09-04     9    14313.7          0.0
16 2011         3  2011-09-05     9    14313.7          0.0

    gdp_deflator  gdp_annual  gdp_annual_growth  gdp_quart_growth_since_2011  \
0          86.721    46308.5             0.0              0.0
1          86.721    46308.5             0.0              0.0
2          86.721    46308.5             0.0              0.0
3          86.721    46308.5             0.0              0.0
4          86.721    46308.5             0.0              0.0
5          86.721    46308.5             0.0              0.0
6          86.721    46308.5             0.0              0.0
7          86.721    46308.5             0.0              0.0
8          86.721    46308.5             0.0              0.0
9          86.721    46308.5             0.0              0.0
10         86.721    46308.5             0.0              0.0
11         86.721    46308.5             0.0              0.0
12         86.721    46308.5             0.0              0.0
13         86.721    46308.5             0.0              0.0
14         86.721    46308.5             0.0              0.0
15         86.721    46308.5             0.0              0.0
16         86.721    46308.5             0.0              0.0

    gdp_annual_growth_since_2011
0              0.0
1              0.0
2              0.0
3              0.0
4              0.0
5              0.0
6              0.0
7              0.0

```

8	0.0
9	0.0
10	0.0
11	0.0
12	0.0
13	0.0
14	0.0
15	0.0
16	0.0

### merge inflation gdp

```
[64]: # merge new features together to one df
inf_gdp_merged = inflation_df.merge(gdp_dfd, on=['date', 'quarter', 'year', 'month'])
# date feature and growth features only
inf_gdp_merged_total_growth = inf_gdp_merged[['date', 'quarter', 'year', 'month', 'inflation_from_8_2011', 'gdp_quart_growth_since_2011', 'gdp_annual_growth_since_2011']]
inf_gdp_merged_period_growth = inf_gdp_merged[['date', 'quarter', 'year', 'month', 'annual_inflation_growth', 'gdp_quart_growth', 'gdp_annual_growth']]
inf_gdp_merged_total_growth
```

```
[64]:
```

	date	quarter	year	month	inflation_from_8_2011 \
0	2011-08-20	3	2011	8	0.00000
1	2011-08-21	3	2011	8	0.00000
2	2011-08-22	3	2011	8	0.00000
3	2011-08-23	3	2011	8	0.00000
4	2011-08-24	3	2011	8	0.00000
...	...	...	...	...	...
1741	2016-05-26	2	2016	5	0.46439
1742	2016-05-27	2	2016	5	0.46439
1743	2016-05-28	2	2016	5	0.46439
1744	2016-05-29	2	2016	5	0.46439
1745	2016-05-30	2	2016	5	0.46439


  

	gdp_quart_growth_since_2011	gdp_annual_growth_since_2011
0	0.000	0.000000
1	0.000	0.000000
2	0.000	0.000000
3	0.000	0.000000
4	0.000	0.000000
...	...	...
1741	0.161	0.060415
1742	0.161	0.060415
1743	0.161	0.060415
1744	0.161	0.060415
1745	0.161	0.060415

[1746 rows x 7 columns]

```
[65]: # plot all total_growth features together over time using plotly express
fig = px.line(inf_gdp_merged_total_growth, x='date',
             y=['inflation_from_8_2011', 'gdp_quart_growth_since_2011',
               'gdp_annual_growth_since_2011'], title='Total growth rates over time',
             color='variable')
fig.show()
```

Total growth rates over time



```
[66]: inf_gdp_merged_period_growth
```

```
[66]:
```

	date	quarter	year	month	annual_inflation_growth	\
0	2011-08-20	3	2011	8	0.00000	
1	2011-08-21	3	2011	8	0.00000	
2	2011-08-22	3	2011	8	0.00000	
3	2011-08-23	3	2011	8	0.00000	
4	2011-08-24	3	2011	8	0.00000	
...	...	...	...	...	...	
1741	2016-05-26	2	2016	5	0.09499	
1742	2016-05-27	2	2016	5	0.09499	


1743	2016-05-28	2	2016	5	0.09499
1744	2016-05-29	2	2016	5	0.09499
1745	2016-05-30	2	2016	5	0.09499

	gdp_quart_growth	gdp_annual_growth
0	0.000	0.000000
1	0.000	0.000000
2	0.000	0.000000
3	0.000	0.000000
4	0.000	0.000000
...	...	...
1741	-0.045	-0.082304
1742	-0.045	-0.082304
1743	-0.045	-0.082304
1744	-0.045	-0.082304
1745	-0.045	-0.082304

[1746 rows x 7 columns]

```
[67]: # plot period growth features together over time using plotly express
fig = px.line(inf_gdp_merged_period_growth, x='date',
             y=['annual_inflation_growth', 'gdp_quart_growth', 'gdp_annual_growth'],
             title='Period growth rates over time', color='variable')
fig.show()
```

Period growth rates over time



```
[68]: # group the period growth features by year and month
inf_gdp_merged_monthly_growth = inf_gdp_merged_period_growth.groupby(['year',
↪ 'month']).mean()
inf_gdp_merged_monthly_growth.reset_index(inplace=True)
inf_gdp_merged_monthly_growth
## the same per quarter
inf_gdp_merged_quarterly_growth = inf_gdp_merged_period_growth.groupby(['year',
↪ 'quarter']).mean()
inf_gdp_merged_quarterly_growth.reset_index(inplace=True)
inf_gdp_merged_quarterly_growth
```

```
[68]:
```

	year	quarter	date	month	annual_inflation_growth \
0	2011	3	2011-09-09 12:00:00	8.714286	0.00000
1	2011	4	2011-11-15 12:00:00	11.000000	0.00000
2	2012	1	2012-02-15 00:00:00	2.000000	0.13279
3	2012	2	2012-05-16 00:00:00	5.000000	0.13279
4	2012	3	2012-08-15 12:00:00	7.989130	0.13279
5	2012	4	2012-11-15 12:00:00	11.000000	0.13279
6	2013	1	2013-02-14 12:00:00	2.000000	0.08299
7	2013	2	2013-05-16 00:00:00	5.000000	0.08299
8	2013	3	2013-08-15 12:00:00	7.989130	0.08299
9	2013	4	2013-11-15 12:00:00	11.000000	0.08299
10	2014	1	2014-02-14 12:00:00	2.000000	0.05166
11	2014	2	2014-05-16 00:00:00	5.000000	0.05166
12	2014	3	2014-08-15 12:00:00	7.989130	0.05166
13	2014	4	2014-11-15 12:00:00	11.000000	0.05166
14	2015	1	2015-02-14 12:00:00	2.000000	0.10196
15	2015	2	2015-05-16 00:00:00	5.000000	0.10196
16	2015	3	2015-08-15 12:00:00	7.989130	0.10196
17	2015	4	2015-11-15 12:00:00	11.000000	0.10196
18	2016	1	2016-02-15 00:00:00	2.000000	0.09499
19	2016	2	2016-04-30 12:00:00	4.500000	0.09499

	gdp_quart_growth	gdp_annual_growth
0	0.000	0.000000
1	0.017	0.000000
2	0.019	-0.002394
3	0.014	-0.002394
4	0.009	-0.002394
5	-0.002	-0.002394
6	-0.013	-0.009858
7	-0.026	-0.009858
8	-0.021	-0.009858
9	-0.020	-0.009858



10	-0.012	-0.032242
11	-0.027	-0.032242
12	-0.026	-0.032242
13	-0.024	-0.032242
14	-0.029	-0.037972
15	-0.061	-0.037972
16	-0.078	-0.037972
17	-0.070	-0.037972
18	-0.071	-0.082304
19	-0.045	-0.082304

```
[69]: # for all ratings we found creating a wining function
def wine_rate(input_rating, input_period_in_year, output_period):
    return (1 + input_rating)**(output_period/input_period_in_year) - 1
```

```
[70]: inf_gdp_merged_quarterly_growth['wined_annual_growth_tquart'] =
    ↪inf_gdp_merged_quarterly_growth['annual_inflation_growth'].apply(lambda x:
    ↪wine_rate(x, 1, 4))
inf_gdp_merged_quarterly_growth['wined_gdp_quart_growth_tquart'] =
    ↪inf_gdp_merged_quarterly_growth['gdp_quart_growth'].apply(lambda x:
    ↪wine_rate(x, 1, 4))
```

```
[71]: # predict quarterly inflation growth rate using the annual inflation growth
    ↪rate, and the gdp growth rates
X = inf_gdp_merged_quarterly_growth[['annual_inflation_growth',
    ↪'gdp_quart_growth', 'gdp_annual_growth']]
# use a wined annual inflation growth rate as the constant for the model to
    ↪predict the quarterly inflation growth rate
```

```
[72]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Define the independent variables (features)
X = inf_gdp_merged_quarterly_growth[['annual_inflation_growth',
    ↪'gdp_quart_growth', 'gdp_annual_growth']]

# Use the annual inflation growth rate as a constant for prediction
constant = inf_gdp_merged_quarterly_growth['annual_inflation_growth'].mean()

# Predict the quarterly inflation growth rate using a simple aggregation (e.g.,
    ↪mean or median)
predicted_quarterly_inflation_growth = X.mean(axis=1) # You can also use other
    ↪aggregation functions like median

# Print the predicted quarterly inflation growth rate
```

```

print("Predicted Quarterly Inflation Growth Rate:")
print(predicted_quarterly_inflation_growth)
# MAKING sure to add a zero in the first row
predicted_quarterly_inflation_growth = pd.concat([pd.Series([0]),
↪predicted_quarterly_inflation_growth], ignore_index=True)
# Print the predicted quarterly inflation growth rate
print("Predicted Quarterly Inflation Growth Rate:")
print(predicted_quarterly_inflation_growth)

```

Predicted Quarterly Inflation Growth Rate:

```

0      0.000000
1      0.005667
2      0.049799
3      0.048132
4      0.046465
5      0.042799
6      0.020044
7      0.015711
8      0.017377
9      0.017711
10     0.002473
11    -0.002527
12    -0.002194
13    -0.001527
14     0.011663
15     0.000996
16    -0.004671
17    -0.002004
18    -0.019438
19    -0.010771

```

dtype: float64

Predicted Quarterly Inflation Growth Rate:

```

0      0.000000
1      0.000000
2      0.005667
3      0.049799
4      0.048132
5      0.046465
6      0.042799
7      0.020044
8      0.015711
9      0.017377
10     0.017711
11     0.002473
12    -0.002527
13    -0.002194
14    -0.001527
15     0.011663

```

```

16    0.000996
17   -0.004671
18   -0.002004
19   -0.019438
20   -0.010771
dtype: float64

```

```

[73]: # add the predicted quarterly inflation growth rate to the dataframe
inf_gdp_merged_quarterly_growth['predicted_quarterly_inflation_growth'] =_
    ↪ predicted_quarterly_inflation_growth
inf_gdp_merged_quarterly_growth

```

```

[73]:      year  quarter      date      month  annual_inflation_growth  \
0    2011         3  2011-09-09 12:00:00    8.714286              0.00000
1    2011         4  2011-11-15 12:00:00   11.000000              0.00000
2    2012         1  2012-02-15 00:00:00    2.000000              0.13279
3    2012         2  2012-05-16 00:00:00    5.000000              0.13279
4    2012         3  2012-08-15 12:00:00    7.989130              0.13279
5    2012         4  2012-11-15 12:00:00   11.000000              0.13279
6    2013         1  2013-02-14 12:00:00    2.000000              0.08299
7    2013         2  2013-05-16 00:00:00    5.000000              0.08299
8    2013         3  2013-08-15 12:00:00    7.989130              0.08299
9    2013         4  2013-11-15 12:00:00   11.000000              0.08299
10   2014         1  2014-02-14 12:00:00    2.000000              0.05166
11   2014         2  2014-05-16 00:00:00    5.000000              0.05166
12   2014         3  2014-08-15 12:00:00    7.989130              0.05166
13   2014         4  2014-11-15 12:00:00   11.000000              0.05166
14   2015         1  2015-02-14 12:00:00    2.000000              0.10196
15   2015         2  2015-05-16 00:00:00    5.000000              0.10196
16   2015         3  2015-08-15 12:00:00    7.989130              0.10196
17   2015         4  2015-11-15 12:00:00   11.000000              0.10196
18   2016         1  2016-02-15 00:00:00    2.000000              0.09499
19   2016         2  2016-04-30 12:00:00    4.500000              0.09499

```

```

      gdp_quart_growth  gdp_annual_growth  wined_annual_growth_tquart  \
0              0.000      0.000000      0.000000
1              0.017      0.000000      0.000000
2              0.019     -0.002394      0.646636
3              0.014     -0.002394      0.646636
4              0.009     -0.002394      0.646636
5             -0.002     -0.002394      0.646636
6             -0.013     -0.009858      0.375618
7             -0.026     -0.009858      0.375618
8             -0.021     -0.009858      0.375618
9             -0.020     -0.009858      0.375618
10            -0.012     -0.032242      0.223211
11            -0.027     -0.032242      0.223211


```

12	-0.026	-0.032242	0.223211
13	-0.024	-0.032242	0.223211
14	-0.029	-0.037972	0.474563
15	-0.061	-0.037972	0.474563
16	-0.078	-0.037972	0.474563
17	-0.070	-0.037972	0.474563
18	-0.071	-0.082304	0.437608
19	-0.045	-0.082304	0.437608

	wined_gdp_quart_growth_tquart	predicted_quarterly_inflation_growth
0	0.000000	0.000000
1	0.069754	0.000000
2	0.078194	0.005667
3	0.057187	0.049799
4	0.036489	0.048132
5	-0.007976	0.046465
6	-0.050995	0.042799
7	-0.100014	0.020044
8	-0.081391	0.015711
9	-0.077632	0.017377
10	-0.047143	0.017711
11	-0.103704	0.002473
12	-0.100014	-0.002527
13	-0.092599	-0.002194
14	-0.111051	-0.001527
15	-0.222568	0.011663
16	-0.277357	0.000996
17	-0.251948	-0.004671
18	-0.255160	-0.002004
19	-0.168210	-0.019438

```
[74]: ## plot the predicted quarterly inflation growth rate vs the actual quarterly
      ↪gdp growth as well as the annual inflation growth rate and annual gdp growth
      ↪rate
      ### over time
      fig = px.line(inf_gdp_merged_quarterly_growth, x=
      ↪inf_gdp_merged_quarterly_growth.index, y=[ 'gdp_quart_growth',
      ↪'predicted_quarterly_inflation_growth'], title='Annual inflation growth rate
      ↪vs predicted quarterly inflation growth rate vs gdp growth rate')
      fig.show()
      ##plot annual igdp growth rate vs predicted quarterly inflation growth rate
      fig = px.line(inf_gdp_merged_quarterly_growth, x=
      ↪inf_gdp_merged_quarterly_growth.index, y=['gdp_annual_growth',
      ↪'annual_inflation_growth'], title='Annual gdp growth rate vs predicted
      ↪quarterly inflation growth rate')
      fig.show()
```

Annual inflation growth rate vs predicted quarterly inflation growth rate vs gd



Annual gdp growth rate vs predicted quarterly inflation growth rate



## 4.7 Categorical Variables

```
[75]: def CategoricalHelper(data, target_encoding = False):
      data = data.select_dtypes('object')
      classes = {}
      for col in data.columns:
          if data[col].nunique() == 2:
              classes[col] = data[col].unique()
          elif col == 'ecology':
              classes[col] = 'Encoder'
          else:
              if target_encoding and col == 'sub_area':
                  classes[col] = 'TargetEncoding'
              else:
                  classes[col] = 'Dummies'
      return classes
```

```
[76]: ecology = ['excellent', 'good', 'satisfactory', 'poor', 'no data']
      ecology_encoding = [4, 3, 2, 1, -1]
      encoder = dict(zip(ecology, ecology_encoding))
      encoder
```

```
[76]: {'excellent': 4, 'good': 3, 'satisfactory': 2, 'poor': 1, 'no data': -1}
```

```
[77]: def categoricalToNumbers(data, classes, encoder, m=300):
      # Data Contains categorical variables only.
      for c in classes:
          if type(classes[c]) == str:
              if classes[c] == 'Dummies':
                  dummies = pd.get_dummies(data[c])*1
                  data = data.drop(columns = c).join(dummies)
              elif classes[c] == 'TargetEncoding':
                  sa_grp = data.groupby(c)['price_sq'].transform('mean')
                  sa_count = data[c].value_counts().to_dict()
                  sa_counts = data[c].map(sa_count)
                  sa_weights = (sa_counts)/(sa_counts + m)
                  data[c] = sa_grp*(sa_weights) + data['price_sq'].
                  ↪mean()*(1-sa_weights)
              else:
                  data.loc[:,c] = data[c].apply(lambda val: encoder[val])
          else:
              data.loc[:,c] = data[c].apply(lambda val: (val == 'yes')*1)
      return data
```

```
[78]: classes = CategoricalHelper(full_data, True)
      classes.pop('product_type')
      classes
```

```
[78]: {'sub_area': 'TargetEncoding',
      'culture_objects_top_25': array(['no', 'yes'], dtype=object),
      'thermal_power_plant_raion': array(['no', 'yes'], dtype=object),
      'incineration_raion': array(['no', 'yes'], dtype=object),
      'oil_chemistry_raion': array(['no', 'yes'], dtype=object),
      'radiation_raion': array(['no', 'yes'], dtype=object),
      'railroad_terminal_raion': array(['no', 'yes'], dtype=object),
      'big_market_raion': array(['no', 'yes'], dtype=object),
      'nuclear_reactor_raion': array(['no', 'yes'], dtype=object),
      'detention_facility_raion': array(['no', 'yes'], dtype=object),
      'water_1line': array(['no', 'yes'], dtype=object),
      'big_road1_1line': array(['no', 'yes'], dtype=object),
      'railroad_1line': array(['no', 'yes'], dtype=object),
      'ecology': 'Encoder'}
```

```
[79]: full_data = categoricalToNumbers(full_data, classes, encoder)
```

## 5 Correlation Analysis

On this part we wanted to understand if our independent features correlate with each other to check if there is some redundancy in our data. So for each product type we checked the correlation within each group.

We have separated the correlation analysis for each product type

```
[80]: # Get the dtypes for each group
      groups_dtypes = {group: list(groups_dfs[group].dtypes.unique()) for group in
      ↪ groups_dfs} # we don't need buildings anymore
      groups_dtypes
```

```
[80]: {'areas': [dtype('float64'), dtype('int64')],
      'buildings': [dtype('float64')],
      'demographics': [dtype('float64')],
      'distances': [dtype('float64'), dtype('int64'), dtype('O')],
      'education': [dtype('float64')],
      'facilities': [dtype('float64'), dtype('int64'), dtype('O')],
      'interior': [dtype('int64'), dtype('O'), dtype('float64')],
      'surroundings': [dtype('O'), dtype('int64'), dtype('float64')]}
```

```
[81]: from matplotlib.colors import LinearSegmentedColormap

      def createHeatMap(corr, group, annot=True):
          plt.figure(figsize=(10,10))
```

```

cmap: LinearSegmentedColormap = sns.diverging_palette(220, 20,
↪as_cmap=True)
sns.heatmap(corr, cmap=cmap, annot=annot, fmt=".2f")
plt.title(f'Correlation Matrix of {group}')
plt.tight_layout()
plt.show()

```

```

[82]: groups_corrs = ['areas', 'demographics', 'surroundings', 'interior', 'distances'] #
↪Groups for the correlation Analysis.

```

```

[83]: # Change ID's to INT
groups_dfs['distances'].loc[:, groups_dfs['distances'].filter(like='ID').
↪columns] = groups_dfs['distances'].filter(like='ID').astype('int64')

```

```

[84]: # Create a dictionary for each group with the correlation matrix.
groups_floats = {group: groups_dfs[group].select_dtypes(['float64']) for group
↪in groups_dfs}

```

### 5.0.1 Investment

```

[85]: investment = full_data[full_data['product_type'] == 'Investment']
investment.drop(columns = 'product_type', inplace=True)

```

/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel\_10945/514358765.py:2:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```


[86]: createHeatMap(investment[groups_floats['areas'].columns].corr(), 'Areas')

```

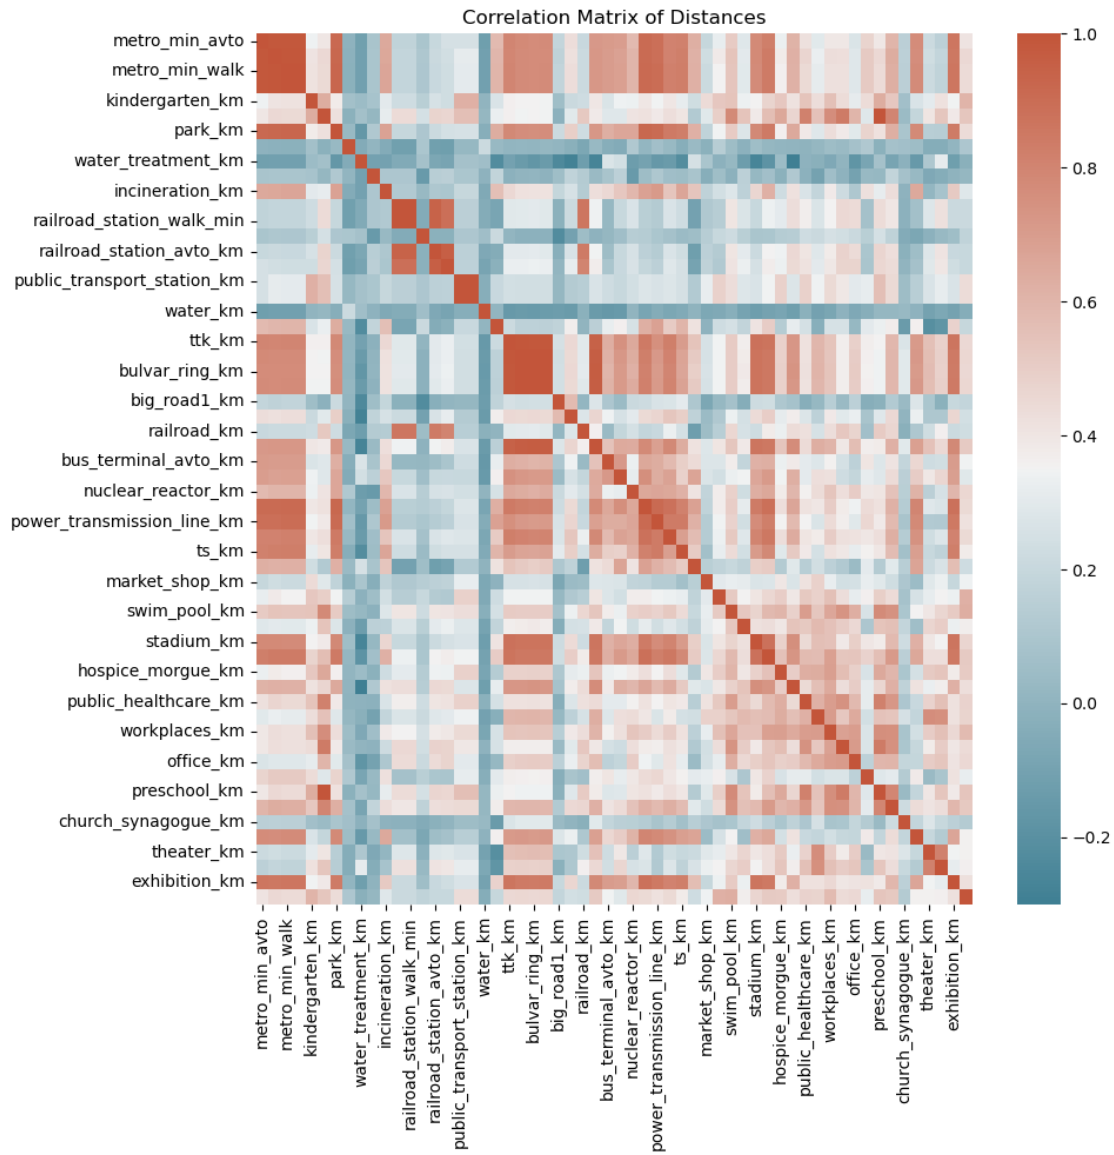





```
[87]: createHeatMap(investment[groups_floats['demographics'].columns].corr(),  
                  ↪ 'Demographics', annot=False)
```



```
[88]: createHeatMap(investment[groups_floats['distances'].columns].corr(),
↳ 'Distances', annot=False)
```




```
[89]: createHeatMap(investment[groups_floats['interior'].columns].corr(), 'Interior')
```




```
[90]: cafe_invst = [col for col in investment.filter(like='cafe').columns if
    ↳ '_binary' not in col]
    green_invst = [col for col in groups_floats['surroundings'].columns if 'green'
    ↳ in col.lower() or 'prom' in col.lower()]
```

```
[91]: createHeatMap(investment[cafe_invst].corr(), 'cafe Surroundings', annot=False)
```



```
[92]: createHeatMap(investment[green_invst].corr(), 'Green & Prom')
```



## 5.0.2 OwnerOccupier


```
[93]: occupier = full_data[full_data['product_type'] == 'OwnerOccupier']
      occupier.drop(columns = 'product_type', inplace=True)
```

```
/var/folders/2y/5vlst1hd6jz9tggyvm776y3m0000gn/T/ipykernel_10945/1144162341.py:2
: SettingWithCopyWarning:
```


A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
[94]: createHeatMap(ocuppier[groups_floats['areas'].columns].corr(), 'Areas')
```




```
[95]: createHeatMap(ocuppier[groups_dfs['demographics'].columns].corr(),  
                    ↪ 'Demographic', annot=False)
```




```
[96]: cafe_ocu = occupier[[col for col in occupier.filter(like='cafe') if '_binary'
    ↪not in col]]
green_prom_ocu = occupier[occupier.columns[np.where(occupier.columns.str.
    ↪contains('green','prom'))]]
```

```
[97]: createHeatMap(corr=cafe_ocu.corr(), group='Cafe', annot=False)
```







```
[98]: createHeatMap(green_prom_ocu.corr(), 'Green & Prom')
```



```
[99]: createHeatMap(ocuppier[groups_floats['interior'].columns].corr(), 'Interior')
```



```
[100]: createHeatMap(ocuppier[groups_floats['distances'].columns].corr(), 'Distances',
↪annot=False)
```



We have some high correlation between independent features which might indicate that we have some redundancy in our data this we might need to use PCA / carefully remove them.

## 6 Principal Component Analysis

To handle the redundancy in our data, we'll use PCA. However, when using PCA choosing the right number of components is crucial. To achieve that we'll use a scree plot and we will choose the number of components that hold between 85%-90% of the variance.

We created two functions that will help us understand what is the number of components that will preserve at least 85% of the variance.

```
[101]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline

# Plotting explained variance ratio vs. number of components
def plotPCA(pca_data, ratio):
    scaler = StandardScaler()
    pca = PCA()

    # Scaler & PCA Pipeline.
    pipe = Pipeline([('scaler',scaler),('pca',pca)])
    pipe.fit(pca_data)

    # Visualize explained variance ratio
    plt.figure(figsize = (10,8))
    plt.plot(np.cumsum(pca.explained_variance_ratio_))
    plt.xticks(np.arange(0, len(np.cumsum(pca.explained_variance_ratio_))) )
    plt.xlabel('Number of Components')
    plt.ylabel('Cumulative Explained Variance Ratio')
    plt.title('Explained Variance Ratio vs. Number of Components')
    plt.grid(True)
    plt.show()
    return np.where(np.cumsum(pca.explained_variance_ratio_) >= ratio)[0].min()
↪+ 1

def pca_df(num_components, data, name):
    # Pca & Scaler pipeline.
    pca = PCA(num_components)
    scaler = StandardScaler()
    pipe_2 = Pipeline([('scaler',scaler),('pca',pca)])
    # Getting the new features.
    x = pipe_2.fit_transform(data)
    names = [f'{name}_pca{i}' for i in range(1,num_components+1)]
    return pd.DataFrame(x, columns=names)
```

## 6.1 Investment

### 6.1.1 Surroundings PCA


We've separated the surroundings into 2 sub groups:

1. Green + Prom
2. Cafe

Then we have applied PCA to reduce the correlation within each group.

```
[102]: cafe_surroundings = groups_floats['surroundings'][cafe_invst]
green_surroundings = groups_floats['surroundings'][green_invst]
```


```
[103]: plotPCA(green_surroundings, 0.95)
```



```
[103]: 6
```

```
[104]: green_pca_invst = pca_df(6, green_surroundings, 'green')
```

```
[105]: plotPCA(cafe_surroundings, 0.95)
```




```
[105]: 5
```

```
[106]: cafe_pca_invst = pca_df(5, cafe_surroundings, 'cafe')
```

### 6.1.2 Distances PCA

```
[107]: distances = groups_floats['distances']

plotPCA(distances, 0.95)
```




[107]: 18

[108]: `distances_pca_invst = pca_df(18,distances, 'distances')`

### 6.1.3 Demographics PCA

[109]: `demographics = groups_floats['demographics']  
plotPCA(demographics, 0.95)`





[109]: 2

```
[110]: demographics_pca_invst = pca_df(2,demographics,'demographic')
```

```
[111]: # Combine the pca data frames:
```

```
# Demographics, Surroundings, Distances
pca_df_inv = pd.concat([cafe_pca_invst, green_pca_invst,
    ↳ demographics_pca_invst, distances_pca_invst], axis=1)
surr_floats = list(groups_dfs['surroundings'].select_dtypes("float64").columns)
demographics = list(groups_dfs['demographics'].filter(like='dens').columns)
distances = list(groups_dfs['distances'].select_dtypes('float64').columns)
new_inv = investment.drop(columns=surr_floats+demographics+distances).
    ↳ reset_index(drop=True)
new_inv[pca_df_inv.columns] = pca_df_inv
```

## 6.2 OwnerOccupier


### 6.2.1 Surroundings PCA

We've separated the surroundings into 2 sub groups:

1. Green + Prom
2. Cafe

Then we have applied PCA to reduce the correlation within each group.


```
[112]: plotPCA(green_prom_ocu, 0.95)
```



```
[112]: 4
```

```
[113]: green_prom_pca_ocu = pca_df(4, green_prom_ocu, 'green_prom')
```

```
[114]: plotPCA(cafe_ocu, 0.95)
```




[114]: 4

```
[115]: cafe_ocu_pca = pca_df(4, cafe_ocu, 'cafe')
```

### 6.2.2 Distances PCA

```
[116]: distances_ocu = occupier[groups_dfs['distances'].columns].
        ↪select_dtypes('float64')

        plotPCA(distances_ocu, 0.95)
```




[116]: 12

```
[117]: distances_pca_ocu = pca_df(12,distances_ocu, 'distances')
```

### 6.2.3 Demographics PCA

```
[118]: demographics_ocu = ocuppier[groups_dfs['demographics'].filter(like='dens').
↪columns]
plotPCA(demographics_ocu, 0.95)
```



[118]: 2

[119]: `demographics_pca_ocu = pca_df(2,demographics_ocu,'demographic')`

[120]: *# Combine the pca data frames:*

*# Demographics, Surroundings, Distances*

```
pca_df_ocu = pd.concat([cafe_ocu_pca, green_prom_pca_ocu,
    ↪demographics_pca_ocu,distances_pca_ocu],axis=1)
surr_floats = list(groups_dfs['surroundings'].select_dtypes("float64").columns)
demographics = list(groups_dfs['demographics'].filter(like='dens').columns)
distances = list(groups_dfs['distances'].select_dtypes('float64').columns)
new_ocu = occupier.drop(columns=surr_floats+demographics+distances).
    ↪reset_index(drop=True)
new_ocu[pca_df_ocu.columns] = pca_df_ocu
```

## 7 Preparing Data For Model Building


```
[121]: # Part 5
new_inv = new_inv[(new_inv.price_sq <= 600000) & (new_inv.price_sq >= 10000) |
↳new_inv.price_sq.isnull()] # reduces errors...
new_ocu = new_ocu[(new_ocu.price_sq <= 600000) & (new_ocu.price_sq >= 10000) |
↳new_ocu.price_sq.isnull()] # reduces errors...
```

### 7.1 Investment

Before building and tuning the models we first need to handle some bad prices values in our investment data.

RECALL:


```
[122]: px.histogram(new_inv['log_price'])
```



Bad Values:

1. 14.9 - 14.95
2. 14.5 - 14.55
3. 13.8 - 13.85

```
[123]: fig = px.histogram(new_inv[~(((new_inv['log_price']>=12) &
    ↪(new_inv['log_price'] <= 14.95)) | (new_inv['log_price'] > 17.
    ↪6))]['log_price'])
fig.update_layout(xaxis_range = [14,18])
```



Seems a little bit better...

```
[124]: inv = new_inv[~(((new_inv['log_price']>=12) & (new_inv['log_price'] <= 14.95))
    ↪| (new_inv['log_price'] > 17.6))].set_index('timestamp').drop(columns = 'id')
```

```
[125]: cat = inv.select_dtypes('object').columns
inv[cat] = inv[cat].astype('int')
response_vars = inv.columns[inv.columns.
    ↪isin(['price_doc','price_sq','log_price'])] # get response variables
inv_response_values = inv[inv.columns[inv.columns.isin(response_vars)]]
inv_features = inv[inv.columns[~inv.columns.isin(inv_response_values.columns)]]
    ↪# get features
```

## 7.2 OwnerOccupier

```
[126]: cat = new_ocu.select_dtypes('object').columns
new_ocu[cat] = new_ocu[cat].astype('int')
ocu = new_ocu.set_index('timestamp').drop(columns = 'id')
response_vars = ocu.columns[ocu.columns.
    ↳isin(['price_doc', 'price_sq', 'log_price'])] # get response variables
ocu_response_values = ocu[ocu.columns[ocu.columns.isin(response_vars)]]
ocu_features = ocu[ocu.columns[~ocu.columns.isin(ocu_response_values.columns)]]
    ↳# get features
```

```
[127]: # Splitting data - train.csv, test.csv
ocu_test = ocu[ocu.price_doc.isnull()].drop(columns=ocu_response_values.columns)
inv_test = inv[inv.price_doc.isnull()].drop(columns=inv_response_values.columns)
ocu_train = ocu[~ocu.price_doc.isnull()]
inv_train = inv[~inv.price_doc.isnull()]
```

```
[128]: # Get response variables values
ocu_price = ocu_train.price_doc
ocu_logprice = ocu_train.log_price
ocu_pricesq = ocu_train.price_sq

inv_price = inv_train.price_doc
inv_logprice = inv_train.log_price
inv_pricesq = inv_train.price_sq
```

```
[129]: final_ocu_train = ocu_train.drop(columns = ['price_doc', 'log_price', 'price_sq'])
final_inv_train = inv_train.drop(columns = ['price_doc', 'log_price', 'price_sq'])
```

```
[130]: from_float_to_int =
    ↳['material', 'state', 'build_year', 'floor', 'max_floor', 'num_room', 'hospital_beds_raion', 'floors']
    ↳+ list(groups_dfs['buildings'].columns) + list(full_data.
    ↳filter(like='missing').columns)
```

```
[131]: final_inv_train[from_float_to_int] = final_inv_train[from_float_to_int].
    ↳astype(int)
final_ocu_train[from_float_to_int] = final_ocu_train[from_float_to_int].
    ↳astype(int)
```

## 8 Model Building and Hyperparameter Tuning

There are several methods for hyperparameter tuning so in order to choose which method to use we read about the ones we know. While researching for reading materials we encountered bayesian optimization method for tuning. (Putatunda, S. et al) Showed that hyperopt (a library for bayesian optimization in python) gave the best results in terms of time complexity and minimizing the loss function compared to random search and grid search thus we decided to use it in our project as well.



```
[238]: def trainTestSplit(x,y):
        x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3,
        ↪random_state=22)
        x_val, x_test, y_val, y_test = train_test_split(x_test, y_test, test_size=0.
        ↪5, random_state=22)
        return x_train, x_val, x_test, y_train, y_val, y_test

[239]: ocu_x_train, ocu_x_val, ocu_x_test, ocu_y_train, ocu_y_val, ocu_y_test =
        ↪trainTestSplit(final_ocu_train, ocu_logprice)
        inv_x_train, inv_x_val, inv_x_test, inv_y_train, inv_y_val, inv_y_test =
        ↪trainTestSplit(final_inv_train, inv_logprice)

[240]: def plotImportance(n_features, importance, features, model,Product):
        sorted_imp = importance.argsort()[::-1]
        df_imp = pd.DataFrame(dict(features = features[sorted_imp], imp =
        ↪importance[sorted_imp]))
        fig = px.bar(df_imp[:n_features], y='features', x='imp')
        fig.update_layout(title = f'Feature Importance - {model} {Product}')
        fig.show()
        print(features[importance.argmax()])
```

## 8.1 XGBoost

We created an objective function for each type of apartments.

```
[241]: def xgb_inv_tuning_allfeatures(parameters):
        model = XGBRegressor(**parameters)
        evaluation = [(inv_x_train,inv_y_train),(inv_x_val, inv_y_val)]
        model.fit(inv_x_train,inv_y_train,eval_set=evaluation,verbose=False)
        preds = model.predict(inv_x_val)
        rmse = mean_squared_error(inv_y_val, preds, squared=False)
        print("Score:",rmse)
        return {'loss':rmse,'status':STATUS_OK,'model':model}

[242]: def xgb_ocu_tuning_allfeatures(parameters):
        model = XGBRegressor(**parameters)
        evaluation = [(ocu_x_train,ocu_y_train),(ocu_x_val, ocu_y_val)]
        model.fit(ocu_x_train, ocu_y_train,eval_set=evaluation,verbose=False)
        preds = model.predict(ocu_x_val)
        rmse = mean_squared_error(ocu_y_val, preds, squared=False)
        print("Score:",rmse)
        return {'loss':rmse,'status':STATUS_OK,'model':model}
```

we defined xgboost parameter space as (Kapoor & Perrone, 2021) did

```
[243]: xgb_parameters = {'max_depth':hp.randint("max_depth", 2,8),
                        'eta':hp.loguniform('eta',np.log((1/10)**3), np.log(1)),
                        'colsample_bytree': hp.uniform('colsample_bytree',0.3,1),
```

```

        'subsample': hp.uniform('subsample',0.5,1),
        'reg_lambda': hp.loguniform('reg_lambda',np.log((1/10)**6),np.
↪log(20)),
        'reg_alpha': hp.loguniform('reg_alpha',np.log((1/10)**6),np.
↪log(20)),
        'gamma': hp.loguniform('gamma', np.log((1/10)**6),np.log(64)),
        'n_estimators': hp.randint('n_estimators',100,1024),
        'eval_metric':'rmse',
        'objective':'reg:squarederror'
    }

```

```
[244]: xgb_trials_ocu = Trials()
```

```
[245]: xgb_trials_inv = Trials()
```

```

[246]: xgb_ocu_best_parameters = fmin(
        fn=xgb_ocu_tuning_allfeatures,
        space = xgb_parameters,
        algo=tpe.suggest,
        max_evals = 50,
        trials=xgb_trials_ocu
    )

print(xgb_ocu_best_parameters)

```

Score:

0.1058841780318957

Score:

0.11760002354235286

Score:

0.11243087575672941

Score:

0.24060875364947543

Score:

0.24511412104881367

Score:

0.14878641873855952

Score:

0.10953223382234713

Score:

0.1835038071362695

Score:

0.12745476681769424

Score:

0.29803440732223524

Score:

0.23810987577761583

Score:

0.15848260810087186  
Score:  
0.18837042037773216  
Score:  
0.2813977395594886  
Score:  
0.19923688360831746  
Score:  
0.11654816048142522  
Score:  
0.12395860184795071  
Score:  
0.37200025166924333  
Score:  
0.12089973907258454  
Score:  
0.21334093942880827  
Score:  
0.1112088978961913  
Score:  
0.12312344756530295  
Score:  
0.11236894022987384  
Score:  
0.11335735138611501  
Score:  
0.19391468129846934  
Score:  
0.13596687735197857  
Score:  
0.20883406337506294  
Score:  
0.11510706838921458  
Score:  
0.1077578131701131  
Score:  
0.10978600129542317  
Score:  
0.11415945882246847  
Score:  
0.11503978225490993  
Score:  
0.11668929415623612  
Score:  
0.10830222249519537  
Score:  
0.10895546970934925  
Score:

```

0.127058205756004
Score:
0.11613781442191491
Score:
0.2303979501671179
Score:
0.12039804630252057
Score:
0.12257767236801981
Score:
0.11837344983528544
Score:
0.11335939746853689
Score:
0.17798761424432913
Score:
0.13528101929824526
Score:
0.1518207944514198
Score:
0.11645348245217892
Score:
0.11375553251210391
Score:
0.13601846248070645
Score:
0.12151958031341414
Score:
0.13146835281146416
100%|      | 50/50 [01:41<00:00, 2.03s/trial, best loss: 0.1058841780318957]
{'colsample_bytree': 0.9372904286503072, 'eta': 0.1623697585579469, 'gamma':
0.0008233655894597945, 'max_depth': 6, 'n_estimators': 134, 'reg_alpha':
0.21573418575839792, 'reg_lambda': 0.0005802138188644974, 'subsample':
0.5570964152767883}

```

```

[247]: best_model_ocu = xgb_trials_ocu.results[np.argmin([r['loss'] for r in
↳xgb_trials_ocu.results if 'loss' in r])]['model'] # gets the best model

```

```

[248]: top_models = 5 # get the top 5 models... we will use average on the predictions.
↳.

```

```

[249]: top_xgb_ocu = sorted(xgb_trials_ocu.results, key= lambda x: x['loss'] if 'loss'
↳in x else 9999)[:top_models] # gets the top 5 models

```

```

[250]: xgb_inv_best_parameters = fmin(
    fn=xgb_inv_tuning_allfeatures,
    space = xgb_parameters,
    algo=tpe.suggest,

```

```
max_evals = 50,  
    trials=xgb_trials_inv  
)  
  
print(xgb_inv_best_parameters)
```

Score:  
0.2644836297439095  
Score:  
0.237984461666739  
Score:  
0.210273106002578  
Score:  
0.1913569946088281  
Score:  
0.20951332148997517  
Score:  
0.28066070348554417  
Score:  
0.22007190133446491  
Score:  
0.19764589327468923  
Score:  
0.1913246020769075  
Score:  
0.24152545107071222  
Score:  
0.20332354595408839  
Score:  
0.19377021365863262  
Score:  
0.19011617733955652  
Score:  
0.19446846784182495  
Score:  
0.1891460588945189  
Score:  
0.20454553558655522  
Score:  
0.3287878977157227  
Score:  
0.19548133299062423  
Score:  
0.21539542026975106  
Score:  
0.3376580378296252  
Score:  
0.19929156283175697

Score:  
0.190075905765267  
Score:  
0.18874587798440692  
Score:  
0.1881638907524976  
Score:  
0.188457930114363  
Score:  
0.2755012072676936  
Score:  
0.19089195515664698  
Score:  
0.27687410051956995  
Score:  
0.18979205129823307  
Score:  
0.20307321240006596  
Score:  
0.1922816831163264  
Score:  
0.28715084202275964  
64%| | 32/50 [01:58<01:30, 5.02s/trial, best loss:  
0.1881638907524976] [CV 2/5] END gamma=0.4, learning\_rate=0.01, max\_depth=3,  
n\_estimators=500;; score=-5127318584929.354 total time= 1.4s  
[CV 4/5] END gamma=0.4, learning\_rate=0.01, max\_depth=8, n\_estimators=600;;  
score=-4317272443258.196 total time= 2.9s  
[CV 2/5] END gamma=0.0, learning\_rate=0.01, max\_depth=7, n\_estimators=700;;  
score=-4192620222329.378 total time= 2.9s  
[CV 4/5] END gamma=0.4, learning\_rate=0.01, max\_depth=3, n\_estimators=500;;  
score=-5232050735579.723 total time= 1.5s  
[CV 5/5] END gamma=0.4, learning\_rate=0.01, max\_depth=8, n\_estimators=600;;  
score=-4580326896855.904 total time= 3.0s  
[CV 3/5] END gamma=0.0, learning\_rate=0.01, max\_depth=7, n\_estimators=700;;  
score=-4508570745843.917 total time= 2.8s  
[CV 1/5] END gamma=0.4, learning\_rate=0.01, max\_depth=3, n\_estimators=500;;  
score=-5164109389688.033 total time= 1.5s  
[CV 3/5] END gamma=0.4, learning\_rate=0.01, max\_depth=6, n\_estimators=900;;  
score=-4504080967896.182 total time= 3.1s  
[CV 5/5] END gamma=0.0, learning\_rate=0.01, max\_depth=7, n\_estimators=700;;  
score=-4595175520303.857 total time= 2.8s  
[CV 1/5] END gamma=0.4, learning\_rate=0.01, max\_depth=8, n\_estimators=600;;  
score=-4321104468889.524 total time= 3.0s  
[CV 1/5] END gamma=0.0, learning\_rate=0.01, max\_depth=7, n\_estimators=700;;  
score=-4331478831631.282 total time= 3.0s  
[CV 2/5] END gamma=0.0, learning\_rate=0.01, max\_depth=3, n\_estimators=500;;  
score=-5127318584929.354 total time= 1.4s  
[CV 3/5] END gamma=0.4, learning\_rate=0.01, max\_depth=3, n\_estimators=500;;

```

score=-5360790269415.222 total time= 1.5s
[CV 1/5] END gamma=0.4, learning_rate=0.01, max_depth=6, n_estimators=900;,
score=-4303690138607.251 total time= 3.2s
[CV 4/5] END gamma=0.0, learning_rate=0.01, max_depth=7, n_estimators=700;,
score=-4312570188166.866 total time= 2.8s
[CV 5/5] END gamma=0.4, learning_rate=0.01, max_depth=3, n_estimators=500;,
score=-5426181289294.336 total time= 1.5s
[CV 2/5] END gamma=0.4, learning_rate=0.01, max_depth=6, n_estimators=900;,
score=-4191622197420.055 total time= 3.2s
[CV 1/5] END gamma=0.0, learning_rate=0.01, max_depth=3, n_estimators=500;,
score=-5164109389688.033 total time= 1.5s
[CV 3/5] END gamma=0.0, learning_rate=0.01, max_depth=3, n_estimators=500;,
score=-5360790269415.222 total time= 1.3s
[CV 2/5] END gamma=0.4, learning_rate=0.01, max_depth=8, n_estimators=600;,
score=-4176254383141.411 total time= 3.0s
[CV 5/5] END gamma=0.4, learning_rate=0.01, max_depth=6, n_estimators=900;,
score=-4613396917786.215 total time= 3.2s
[CV 4/5] END gamma=0.0, learning_rate=0.01, max_depth=3, n_estimators=500;,
score=-5232050735579.723 total time= 1.3s
[CV 3/5] END gamma=0.4, learning_rate=0.01, max_depth=8, n_estimators=600;,
score=-4488427717458.747 total time= 2.9s
[CV 4/5] END gamma=0.4, learning_rate=0.01, max_depth=6, n_estimators=900;,
score=-4321196746451.036 total time= 3.2s
[CV 5/5] END gamma=0.0, learning_rate=0.01, max_depth=3, n_estimators=500;,
score=-5426181289294.336 total time= 1.3s
Score:
0.19445228859836627
Score:
0.19794270711307135
Score:
0.19028979628278397
Score:
0.22037142354119835
Score:
0.19148575567196613
Score:
0.19284961295365233
Score:
0.22837141806007372
Score:
0.19147764565405398
Score:
0.22101366067907519
Score:
0.2704537532194982
Score:
0.22802909721355305
Score:

```

```

0.21628700098937664
Score:
0.1918366088728885
Score:
0.1893375517304326
Score:
0.21708204548126495
Score:
0.1906910116265093
Score:
0.2410008190273595
Score:
0.18866423954997255
100%|      | 50/50 [03:01<00:00, 3.62s/trial, best loss: 0.1881638907524976]
{'colsample_bytree': 0.4017398986910317, 'eta': 0.014062305138921366, 'gamma':
1.5298470665441234e-05, 'max_depth': 7, 'n_estimators': 698, 'reg_alpha':
0.0013533958007616535, 'reg_lambda': 9.091311069282547e-06, 'subsample':
0.9522562407615086}

```

```

[251]: best_model_inv = xgb_trials_inv.results[np.argmin([r['loss'] for r in
↳xgb_trials_inv.results if 'loss' in r ])]['model']


[252]: top_xgb_inv = sorted(xgb_trials_inv.results, key= lambda x: x['loss'] if 'loss'
↳in x else 9999)[:top_models]

[253]: plotImportance(n_features=20, importance=best_model_inv.feature_importances_,
↳features=inv_x_train.columns, model="XGB", Product="INV")

```




## Feature Importance - XGB INV



cafe\_count\_5000\_price\_high\_binary

```
[254]: plotImportance(n_features=20, importance=best_model_ocu.feature_importances_,
    ↪ features= ocu_x_train.columns, model="XGB", Product="OCU")
```

## Feature Importance - XGB OCU



sub\_area

## 8.2 Random Forest

```
[255]: rf_parameters = {
    'n_estimators': hp.choice('n_estimators', np.arange(50,1050,50)),
    'max_features': hp.choice('max_features', ['log2','sqrt']),
    'criterion': hp.choice('criterion', ['friedman_mse', 'squared_error']),
    'max_depth': scope.int(hp.uniform('max_depth', 5,20)),
    'max_samples': hp.uniform('max_samples', 0,0.8),
    'min_samples_leaf': scope.int(hp.uniform('min_samples_leaf',1,5)),
    'min_samples_split': scope.int(hp.uniform('min_samples_split',2,6))
}
```

```
[256]: def rf_inv_tuning_allfeatures(parameters):
    model = RandomForestRegressor(**parameters)
    model.fit(inv_x_train,inv_y_train)
    preds = model.predict(inv_x_val)
    rmse = mean_squared_error(inv_y_val, preds, squared=False)
    print("Score:",rmse)
    return {'loss':rmse,'status':STATUS_OK,'model':model}
```

```
[257]: def rf_ocu_tuning_allfeatures(parameters):
        model = RandomForestRegressor(**parameters)
        model.fit(ocu_x_train, ocu_y_train)
        preds = model.predict(ocu_x_val)
        rmse = mean_squared_error(ocu_y_val, preds, squared=False)
        print("Score:",rmse)
        return {'loss':rmse, 'status':STATUS_OK, 'model':model}
```

```
[258]: rf_trials_inv = Trials()
```

```
[259]: rf_inv_best_parameters = fmin(
        fn=rf_inv_tuning_allfeatures,
        space = rf_parameters,
        algo=tpe.suggest,
        max_evals = 20,
        trials=rf_trials_inv
    )
    print(rf_inv_best_parameters)
```

```
Score:
0.22299355804043655
Score:
0.25935196368644237
Score:
0.21612811558531436
Score:
0.2089327336511747
Score:
0.26499524273401126
Score:
0.20532434921530165
Score:
0.2633114956032697
Score:
0.2252921786079209
Score:
0.20292189754500922
Score:
0.2380140243122244
Score:
0.2821369101963713
Score:
0.20042785505870644
Score:
0.21054567611959404
Score:
0.21399065674874168
Score:
```

```

0.2311431083680959
Score:
0.25176518915744345
Score:
0.21538544204616264
Score:
0.26868475006811027
Score:
0.24079665126419023
Score:
0.21764865318343715
100%|      | 20/20 [02:40<00:00, 8.04s/trial, best loss: 0.20042785505870644]
{'criterion': 0, 'max_depth': 15.874326197235035, 'max_features': 1,
'max_samples': 0.658117314106804, 'min_samples_leaf': 1.8732209888405222,
'min_samples_split': 2.676734090766692, 'n_estimators': 12}

```

```
[260]: rf_trials_ocu = Trials()
```

```

[261]: rf_ocu_best_parameters = fmin(
        fn=rf_ocu_tuning_allfeatures,
        space = rf_parameters,
        algo=tpe.suggest,
        max_evals = 20,
        trials=rf_trials_ocu
    )
print(rf_ocu_best_parameters)

```

```

Score:
0.18007082690660559
Score:
0.18707279774761876
Score:
0.19483020273773885
Score:
0.1329407121782871
Score:
0.1498122550728186
Score:
0.15415476376610596
Score:
0.17535495731279
Score:
0.1565775068465129
Score:
0.15084188954129624
Score:
0.13779421928847332
Score:

```

```

0.22146974159414087
Score:
0.16374522672300335
Score:
0.16102651536475218
Score:
0.12958525507623309
Score:
0.17348097021807224
Score:
0.21786415805731146
Score:
0.13635373388878674
Score:
0.14317843043125195
Score:
0.16021760062243592
Score:
0.20190852630409437
100%|      | 20/20 [00:40<00:00, 2.02s/trial, best loss: 0.12958525507623309]
{'criterion': 0, 'max_depth': 15.768681850617753, 'max_features': 0,
'max_samples': 0.7134946183782201, 'min_samples_leaf': 1.7281239019674,
'min_samples_split': 5.03356610318302, 'n_estimators': 14}

```

```

[262]: best_rf_ocu = rf_trials_ocu.results[np.argmin([r['loss'] for r in rf_trials_ocu.
↳results if 'loss' in r]))['model']
best_rf_inv = rf_trials_inv.results[np.argmin([r['loss'] for r in rf_trials_inv.
↳results if 'loss' in r]))['model']

top5_rf_ocu = sorted(rf_trials_ocu.results, key= lambda x: x['loss'] if 'loss' in
↳x else 9999)[:top_models]
top5_rf_inv = sorted(rf_trials_inv.results, key= lambda x: x['loss'] if 'loss' in
↳x else 9999)[:top_models]


```

```

[263]: plotImportance(n_features=20, importance=best_rf_ocu.feature_importances_,
↳features=ocu_x_train.columns, model="RF", Product="OCU")

```


### Feature Importance - RF OCU



full\_sq

```
[264]: plotImportance(n_features=20, importance=best_rf_inv.feature_importances_,
    ↪ features=inv_x_train.columns, model="RF", Product="INV")
```

## Feature Importance - RF INV



area\_diff

## 8.3 Test Predictions

```
[265]: inv_xgb_testpreds = best_model_inv.predict(inv_x_test)
inv_rf_testpreds = best_rf_inv.predict(inv_x_test)

ocu_xgb_testpreds = best_model_ocu.predict(ocu_x_test)
ocu_rf_testpreds = best_rf_ocu.predict(ocu_x_test)

test_inv_preds = pd.DataFrame(dict(year = inv_x_test.year, quarter = inv_x_test.
    ↳ quarter, xgb = np.expm1(inv_xgb_testpreds), rf=np.
    ↳ expm1(inv_rf_testpreds), true = np.expm1(inv_y_test)))
#test_inv_preds = pd.DataFrame(dict(year = inv_x_test.year, quarter = inv_x_test.
    ↳ quarter, xgb = inv_xgb_testpreds*inv_x_test['full_sq'], rf=
    ↳ inv_rf_testpreds*inv_x_test['full_sq'], true =
    ↳ inv_y_test*inv_x_test['full_sq']))
test_ocu_preds = pd.DataFrame(dict(year = ocu_x_test.year, quarter = ocu_x_test.
    ↳ quarter, xgb = np.expm1(ocu_xgb_testpreds), rf=np.expm1(ocu_rf_testpreds),
    ↳ true = np.expm1(ocu_y_test)))
```


```
[266]: q = test_inv_preds.groupby(['year', 'quarter'])
q2 = test_ocu_preds.groupby(['year', 'quarter'])
```

```
[267]: melted_q = q.mean().reset_index().melt(id_vars='year', var_name='model',
        value_name='prediction')
melted_q['year'] = melted_q['year'].apply(lambda year: str(year))
sns.set_style("whitegrid")
plt.figure(figsize=(10, 6))

# Plot the lines for different models
sns.lineplot(data=melted_q[melted_q['model'] != 'quarter'], x='year',
             y='prediction', hue='model')

plt.xlabel('Year')
plt.ylabel('Price')
plt.title('Predictions VS True prices (Investment)')

plt.legend(title='Model')
plt.show()
```



```
[268]: melted_q2 = q2.mean().reset_index().melt(id_vars='year', var_name='model',
        value_name='prediction')
melted_q2['year'] = melted_q2['year'].apply(lambda year: str(year))
```



```


sns.set_style("whitegrid")
plt.figure(figsize=(10, 6))

# Plot the lines for different models
sns.lineplot(data=melted_q2[melted_q2['model'] != 'quarter'], x='year', y=
    y='prediction', hue='model')

plt.xlabel('Year')
plt.ylabel('Price')
plt.title('Predictions VS True prices (Owner Ocuppier)')

plt.legend(title='Model')
plt.show()

```



It seems that for both xgboost and random forest models it's easier to predict the prices for the Owner Ocuppier product type.

XGboost - Give better results in both cases.

XGBoost - Builds trees sequentially, optimize a specific objective function.

Random Forest - builds trees independently with no interaction between them.

The sequential nature of xgboost allows it to learn complex patterns as the current tree learns from previous trees mistakes (reduces bias but prone to overfitting) while random forest may capture different pattern due to the randomness during the training phase.

As mentioned above, random forest builds trees independently where each of the trees learn a certain part of the data and at the end the predictions are summarized using average. This method reduces the variance of the model and allows a better generalization to unseen data.

So we believe that the difference between the two models predictions is probably because xgboost can learn complex patterns while a bagging ensemble model (random forest) is good when the model has low bias already.

Though XGBOOST predictions seem to be a little better, we still wanted to try and ensemble them.

## 9 XGBoost without Hyper Tuning

Though hyper parameter tuning algorithm can help get better results we decided to try choosing the right parameters ourselves. We did some research online, and decided to create 5 parameters sets so we ended up having 5 xgboost models for each product type.

```
[133]: import xgboost as xgb
def no_tune_train(params, x_train,y_train):
    dtrain = xgb.DMatrix(x_train, y_train)

    boosts = xgb.cv(params, dtrain, num_boost_round = 10000, verbose_eval = 100, early_stopping_rounds = 50)

    xgb_model = xgb.train(params, dtrain,num_boost_round=len(boosts))
    return xgb_model
```

### 9.1 Log Price

```
[270]: ocu_xgb_no_tune_params1 = {
    'subsample':0.7,
    'max_depth': 4,
    'colsample_bytree':0.7,
    'lambda': 200,
    'alpha':8,
    'eta': 0.05
}
ocu_xgb_no_tune_params2 = {
    'subsample':1,
    'max_depth':4,
    'colsample_bytree':0.7,
    'lambda': 200,
    'alpha':8,
    'eta':0.05
}
ocu_xgb_no_tune_params3 = {
    'subsample':1,
```

```

    'max_depth':4,
    'colsample_bytree':1,
    'lambda': 200,
    'alpha':8,
    'eta':0.05,
    'nthread':8
}

ocu_xgb_no_tune_params4 = {
    'subsample':0.9,
    'max_depth':4,
    'colsample_bytree':0.7,
    'lambda': 200,
    'alpha':8,
    'eta':0.05
}

ocu_xgb_no_tune_params5 = {
    'subsample':0.8,
    'max_depth':7,
    'colsample_bytree':0.7,
    'lambda': 200,
    'alpha':8,
    'eta':0.05
}

ocu_notue_xgb1 = no_tune_train(params=ocu_xgb_no_tune_params1,
    ↪x_train=final_ocu_train,y_train=ocu_train.log_price)
ocu_notue_xgb2 = no_tune_train(params=ocu_xgb_no_tune_params2,
    ↪x_train=final_ocu_train,y_train=ocu_train.log_price)
ocu_notue_xgb3 = no_tune_train(params=ocu_xgb_no_tune_params3,
    ↪x_train=final_ocu_train,y_train=ocu_train.log_price)
ocu_notue_xgb4 = no_tune_train(params=ocu_xgb_no_tune_params4,
    ↪x_train=final_ocu_train,y_train=ocu_train.log_price)
ocu_notue_xgb5 = no_tune_train(params=ocu_xgb_no_tune_params5,
    ↪x_train=final_ocu_train,y_train=ocu_train.log_price)

```

[0]	train-rmse:0.43817+0.00228	test-rmse:0.43827+0.00482
[100]	train-rmse:0.15363+0.00092	test-rmse:0.16070+0.00237
[200]	train-rmse:0.13024+0.00113	test-rmse:0.14034+0.00156
[300]	train-rmse:0.12142+0.00133	test-rmse:0.13335+0.00159
[400]	train-rmse:0.11649+0.00152	test-rmse:0.12965+0.00158
[500]	train-rmse:0.11323+0.00163	test-rmse:0.12726+0.00158
[600]	train-rmse:0.11071+0.00168	test-rmse:0.12545+0.00170
[700]	train-rmse:0.10888+0.00170	test-rmse:0.12425+0.00172
[800]	train-rmse:0.10744+0.00173	test-rmse:0.12330+0.00187
[900]	train-rmse:0.10623+0.00173	test-rmse:0.12252+0.00191
[1000]	train-rmse:0.10515+0.00172	test-rmse:0.12179+0.00200

[1100]	train-rmse:0.10424+0.00173	test-rmse:0.12121+0.00206
[1200]	train-rmse:0.10349+0.00176	test-rmse:0.12071+0.00210
[1300]	train-rmse:0.10280+0.00174	test-rmse:0.12031+0.00212
[1400]	train-rmse:0.10219+0.00175	test-rmse:0.11994+0.00218
[1500]	train-rmse:0.10166+0.00172	test-rmse:0.11967+0.00220
[1600]	train-rmse:0.10120+0.00169	test-rmse:0.11941+0.00225
[1700]	train-rmse:0.10077+0.00164	test-rmse:0.11914+0.00233
[1800]	train-rmse:0.10035+0.00162	test-rmse:0.11892+0.00237
[1900]	train-rmse:0.10001+0.00158	test-rmse:0.11873+0.00243
[2000]	train-rmse:0.09966+0.00157	test-rmse:0.11854+0.00246
[2100]	train-rmse:0.09936+0.00157	test-rmse:0.11838+0.00248
[2200]	train-rmse:0.09905+0.00156	test-rmse:0.11819+0.00250
[2300]	train-rmse:0.09877+0.00155	test-rmse:0.11804+0.00254
[2400]	train-rmse:0.09849+0.00154	test-rmse:0.11787+0.00258
[2500]	train-rmse:0.09821+0.00152	test-rmse:0.11770+0.00259
[2600]	train-rmse:0.09797+0.00153	test-rmse:0.11760+0.00260
[2700]	train-rmse:0.09775+0.00153	test-rmse:0.11748+0.00261
[2800]	train-rmse:0.09752+0.00151	test-rmse:0.11736+0.00265
[2900]	train-rmse:0.09728+0.00151	test-rmse:0.11723+0.00268
[3000]	train-rmse:0.09706+0.00151	test-rmse:0.11709+0.00271
[3100]	train-rmse:0.09685+0.00149	test-rmse:0.11698+0.00272
[3200]	train-rmse:0.09667+0.00150	test-rmse:0.11690+0.00274
[3300]	train-rmse:0.09645+0.00150	test-rmse:0.11679+0.00275
[3400]	train-rmse:0.09629+0.00148	test-rmse:0.11671+0.00276
[3500]	train-rmse:0.09610+0.00148	test-rmse:0.11664+0.00278
[3600]	train-rmse:0.09593+0.00149	test-rmse:0.11657+0.00276
[3700]	train-rmse:0.09577+0.00147	test-rmse:0.11648+0.00279
[3800]	train-rmse:0.09560+0.00146	test-rmse:0.11639+0.00281
[3900]	train-rmse:0.09545+0.00147	test-rmse:0.11632+0.00282
[4000]	train-rmse:0.09530+0.00146	test-rmse:0.11624+0.00284
[4100]	train-rmse:0.09516+0.00146	test-rmse:0.11617+0.00282
[4200]	train-rmse:0.09503+0.00144	test-rmse:0.11610+0.00285
[4300]	train-rmse:0.09489+0.00145	test-rmse:0.11603+0.00284
[4400]	train-rmse:0.09477+0.00144	test-rmse:0.11596+0.00288
[4500]	train-rmse:0.09466+0.00144	test-rmse:0.11590+0.00287
[4600]	train-rmse:0.09452+0.00143	test-rmse:0.11584+0.00288
[4700]	train-rmse:0.09441+0.00142	test-rmse:0.11578+0.00288
[4800]	train-rmse:0.09430+0.00140	test-rmse:0.11573+0.00291
[4900]	train-rmse:0.09419+0.00139	test-rmse:0.11569+0.00292
[5000]	train-rmse:0.09408+0.00140	test-rmse:0.11564+0.00293
[5100]	train-rmse:0.09397+0.00140	test-rmse:0.11560+0.00293
[5200]	train-rmse:0.09387+0.00139	test-rmse:0.11556+0.00294
[5300]	train-rmse:0.09376+0.00139	test-rmse:0.11551+0.00295
[5400]	train-rmse:0.09366+0.00139	test-rmse:0.11546+0.00296
[5500]	train-rmse:0.09356+0.00140	test-rmse:0.11542+0.00296
[5600]	train-rmse:0.09346+0.00140	test-rmse:0.11538+0.00296
[5700]	train-rmse:0.09337+0.00141	test-rmse:0.11532+0.00297
[5800]	train-rmse:0.09328+0.00141	test-rmse:0.11530+0.00297

[5900]	train-rmse:0.09317+0.00140	test-rmse:0.11524+0.00299
[6000]	train-rmse:0.09309+0.00141	test-rmse:0.11521+0.00299
[6100]	train-rmse:0.09299+0.00140	test-rmse:0.11514+0.00300
[6200]	train-rmse:0.09290+0.00141	test-rmse:0.11513+0.00300
[6300]	train-rmse:0.09281+0.00142	test-rmse:0.11510+0.00300
[6400]	train-rmse:0.09271+0.00140	test-rmse:0.11505+0.00304
[6500]	train-rmse:0.09263+0.00141	test-rmse:0.11500+0.00305
[6600]	train-rmse:0.09256+0.00141	test-rmse:0.11496+0.00305
[6700]	train-rmse:0.09248+0.00140	test-rmse:0.11492+0.00306
[6751]	train-rmse:0.09244+0.00141	test-rmse:0.11492+0.00306
[0]	train-rmse:0.43740+0.00235	test-rmse:0.43751+0.00472
[100]	train-rmse:0.14626+0.00112	test-rmse:0.15512+0.00156
[200]	train-rmse:0.12483+0.00098	test-rmse:0.13708+0.00176
[300]	train-rmse:0.11613+0.00133	test-rmse:0.13024+0.00191
[400]	train-rmse:0.11141+0.00134	test-rmse:0.12663+0.00221
[500]	train-rmse:0.10841+0.00141	test-rmse:0.12445+0.00229
[595]	train-rmse:0.10766+0.00124	test-rmse:0.12389+0.00243
[0]	train-rmse:0.43685+0.00219	test-rmse:0.43702+0.00485
[100]	train-rmse:0.14571+0.00117	test-rmse:0.15488+0.00162
[200]	train-rmse:0.12427+0.00096	test-rmse:0.13703+0.00194
[300]	train-rmse:0.11555+0.00144	test-rmse:0.13039+0.00201
[400]	train-rmse:0.11079+0.00154	test-rmse:0.12684+0.00213
[500]	train-rmse:0.10772+0.00160	test-rmse:0.12467+0.00220
[581]	train-rmse:0.10743+0.00137	test-rmse:0.12449+0.00240
[0]	train-rmse:0.43758+0.00227	test-rmse:0.43771+0.00473
[100]	train-rmse:0.14824+0.00094	test-rmse:0.15651+0.00217
[200]	train-rmse:0.12622+0.00099	test-rmse:0.13785+0.00190
[300]	train-rmse:0.11766+0.00122	test-rmse:0.13118+0.00212
[400]	train-rmse:0.11271+0.00139	test-rmse:0.12746+0.00218
[500]	train-rmse:0.10951+0.00148	test-rmse:0.12521+0.00222
[600]	train-rmse:0.10721+0.00153	test-rmse:0.12363+0.00230
[700]	train-rmse:0.10556+0.00158	test-rmse:0.12257+0.00236
[800]	train-rmse:0.10441+0.00161	test-rmse:0.12184+0.00239
[900]	train-rmse:0.10356+0.00168	test-rmse:0.12133+0.00243
[1000]	train-rmse:0.10280+0.00168	test-rmse:0.12083+0.00249
[1100]	train-rmse:0.10225+0.00170	test-rmse:0.12050+0.00253
[1200]	train-rmse:0.10181+0.00170	test-rmse:0.12023+0.00260
[1300]	train-rmse:0.10139+0.00170	test-rmse:0.11995+0.00261
[1400]	train-rmse:0.10104+0.00168	test-rmse:0.11975+0.00263
[1500]	train-rmse:0.10072+0.00171	test-rmse:0.11956+0.00265
[1600]	train-rmse:0.10044+0.00169	test-rmse:0.11938+0.00268
[1700]	train-rmse:0.10016+0.00169	test-rmse:0.11924+0.00267
[1800]	train-rmse:0.09994+0.00169	test-rmse:0.11911+0.00270
[1900]	train-rmse:0.09975+0.00166	test-rmse:0.11900+0.00273
[2000]	train-rmse:0.09958+0.00164	test-rmse:0.11890+0.00273
[2100]	train-rmse:0.09940+0.00167	test-rmse:0.11880+0.00275
[2200]	train-rmse:0.09926+0.00164	test-rmse:0.11872+0.00276
[2300]	train-rmse:0.09911+0.00164	test-rmse:0.11862+0.00278

[2400]	train-rmse:0.09896+0.00164	test-rmse:0.11853+0.00278
[2500]	train-rmse:0.09881+0.00162	test-rmse:0.11844+0.00279
[2600]	train-rmse:0.09868+0.00163	test-rmse:0.11837+0.00279
[2700]	train-rmse:0.09855+0.00163	test-rmse:0.11828+0.00280
[2800]	train-rmse:0.09845+0.00161	test-rmse:0.11822+0.00282
[2900]	train-rmse:0.09834+0.00161	test-rmse:0.11816+0.00282
[3000]	train-rmse:0.09823+0.00160	test-rmse:0.11810+0.00284
[3100]	train-rmse:0.09813+0.00160	test-rmse:0.11803+0.00285
[3200]	train-rmse:0.09803+0.00160	test-rmse:0.11798+0.00286
[3300]	train-rmse:0.09792+0.00160	test-rmse:0.11791+0.00286
[3400]	train-rmse:0.09781+0.00159	test-rmse:0.11785+0.00285
[3500]	train-rmse:0.09771+0.00160	test-rmse:0.11779+0.00287
[3600]	train-rmse:0.09762+0.00159	test-rmse:0.11775+0.00287
[3700]	train-rmse:0.09754+0.00157	test-rmse:0.11770+0.00289
[3800]	train-rmse:0.09746+0.00157	test-rmse:0.11765+0.00290
[3900]	train-rmse:0.09740+0.00158	test-rmse:0.11762+0.00291
[4000]	train-rmse:0.09732+0.00158	test-rmse:0.11756+0.00291
[4100]	train-rmse:0.09721+0.00157	test-rmse:0.11749+0.00292
[4200]	train-rmse:0.09713+0.00160	test-rmse:0.11746+0.00292
[4300]	train-rmse:0.09707+0.00159	test-rmse:0.11742+0.00293
[4400]	train-rmse:0.09700+0.00159	test-rmse:0.11738+0.00294
[4500]	train-rmse:0.09695+0.00160	test-rmse:0.11735+0.00294
[4600]	train-rmse:0.09691+0.00159	test-rmse:0.11733+0.00294
[4700]	train-rmse:0.09684+0.00157	test-rmse:0.11730+0.00295
[4800]	train-rmse:0.09676+0.00156	test-rmse:0.11726+0.00297
[4900]	train-rmse:0.09669+0.00157	test-rmse:0.11722+0.00297
[5000]	train-rmse:0.09663+0.00157	test-rmse:0.11720+0.00299
[5100]	train-rmse:0.09657+0.00157	test-rmse:0.11717+0.00299
[5200]	train-rmse:0.09653+0.00157	test-rmse:0.11714+0.00300
[5300]	train-rmse:0.09648+0.00158	test-rmse:0.11712+0.00300
[5400]	train-rmse:0.09641+0.00157	test-rmse:0.11708+0.00300
[5500]	train-rmse:0.09637+0.00157	test-rmse:0.11706+0.00300
[5600]	train-rmse:0.09631+0.00157	test-rmse:0.11703+0.00300
[5700]	train-rmse:0.09626+0.00155	test-rmse:0.11699+0.00302
[5800]	train-rmse:0.09621+0.00154	test-rmse:0.11697+0.00303
[5900]	train-rmse:0.09616+0.00155	test-rmse:0.11694+0.00303
[6000]	train-rmse:0.09611+0.00155	test-rmse:0.11692+0.00304
[6100]	train-rmse:0.09607+0.00156	test-rmse:0.11690+0.00304
[6200]	train-rmse:0.09601+0.00156	test-rmse:0.11686+0.00304
[6300]	train-rmse:0.09595+0.00156	test-rmse:0.11684+0.00304
[6400]	train-rmse:0.09590+0.00156	test-rmse:0.11681+0.00304
[6500]	train-rmse:0.09586+0.00156	test-rmse:0.11680+0.00305
[6600]	train-rmse:0.09582+0.00155	test-rmse:0.11677+0.00306
[6700]	train-rmse:0.09578+0.00154	test-rmse:0.11675+0.00307
[6800]	train-rmse:0.09573+0.00155	test-rmse:0.11674+0.00307
[6900]	train-rmse:0.09570+0.00155	test-rmse:0.11671+0.00307
[7000]	train-rmse:0.09566+0.00155	test-rmse:0.11670+0.00306
[7100]	train-rmse:0.09563+0.00154	test-rmse:0.11669+0.00306

[7200]	train-rmse:0.09559+0.00153	test-rmse:0.11667+0.00307
[7300]	train-rmse:0.09556+0.00153	test-rmse:0.11665+0.00308
[7400]	train-rmse:0.09552+0.00152	test-rmse:0.11661+0.00308
[7500]	train-rmse:0.09547+0.00151	test-rmse:0.11660+0.00309
[7600]	train-rmse:0.09545+0.00151	test-rmse:0.11658+0.00309
[7700]	train-rmse:0.09541+0.00151	test-rmse:0.11656+0.00310
[7800]	train-rmse:0.09537+0.00151	test-rmse:0.11654+0.00310
[7900]	train-rmse:0.09535+0.00150	test-rmse:0.11653+0.00310
[8000]	train-rmse:0.09532+0.00150	test-rmse:0.11651+0.00310
[8100]	train-rmse:0.09528+0.00150	test-rmse:0.11649+0.00310
[8200]	train-rmse:0.09525+0.00150	test-rmse:0.11647+0.00311
[8300]	train-rmse:0.09522+0.00149	test-rmse:0.11646+0.00312
[8400]	train-rmse:0.09519+0.00150	test-rmse:0.11644+0.00312
[8500]	train-rmse:0.09517+0.00149	test-rmse:0.11643+0.00312
[8600]	train-rmse:0.09514+0.00149	test-rmse:0.11642+0.00312
[8700]	train-rmse:0.09510+0.00149	test-rmse:0.11639+0.00313
[8800]	train-rmse:0.09507+0.00149	test-rmse:0.11638+0.00313
[8900]	train-rmse:0.09503+0.00148	test-rmse:0.11637+0.00313
[9000]	train-rmse:0.09500+0.00148	test-rmse:0.11635+0.00314
[9100]	train-rmse:0.09498+0.00148	test-rmse:0.11634+0.00314
[9200]	train-rmse:0.09496+0.00148	test-rmse:0.11633+0.00314
[9300]	train-rmse:0.09493+0.00148	test-rmse:0.11632+0.00314
[9400]	train-rmse:0.09490+0.00147	test-rmse:0.11630+0.00314
[9500]	train-rmse:0.09486+0.00147	test-rmse:0.11628+0.00314
[9600]	train-rmse:0.09484+0.00147	test-rmse:0.11627+0.00315
[9700]	train-rmse:0.09481+0.00147	test-rmse:0.11625+0.00315
[9800]	train-rmse:0.09478+0.00146	test-rmse:0.11623+0.00316
[9900]	train-rmse:0.09476+0.00146	test-rmse:0.11622+0.00316
[9999]	train-rmse:0.09473+0.00145	test-rmse:0.11621+0.00316
[0]	train-rmse:0.43751+0.00223	test-rmse:0.43769+0.00477
[100]	train-rmse:0.13752+0.00116	test-rmse:0.15028+0.00192
[200]	train-rmse:0.11485+0.00122	test-rmse:0.13210+0.00147
[300]	train-rmse:0.10650+0.00156	test-rmse:0.12609+0.00176
[400]	train-rmse:0.10206+0.00168	test-rmse:0.12315+0.00179
[500]	train-rmse:0.09932+0.00172	test-rmse:0.12147+0.00192
[600]	train-rmse:0.09748+0.00172	test-rmse:0.12039+0.00202
[700]	train-rmse:0.09615+0.00184	test-rmse:0.11967+0.00203
[800]	train-rmse:0.09518+0.00180	test-rmse:0.11912+0.00212
[900]	train-rmse:0.09449+0.00183	test-rmse:0.11876+0.00214
[1000]	train-rmse:0.09379+0.00177	test-rmse:0.11836+0.00219
[1100]	train-rmse:0.09327+0.00178	test-rmse:0.11807+0.00219
[1200]	train-rmse:0.09282+0.00181	test-rmse:0.11783+0.00223
[1300]	train-rmse:0.09237+0.00180	test-rmse:0.11762+0.00224
[1400]	train-rmse:0.09199+0.00180	test-rmse:0.11742+0.00227
[1500]	train-rmse:0.09168+0.00182	test-rmse:0.11728+0.00225
[1600]	train-rmse:0.09135+0.00185	test-rmse:0.11712+0.00227
[1700]	train-rmse:0.09105+0.00181	test-rmse:0.11696+0.00234
[1800]	train-rmse:0.09076+0.00180	test-rmse:0.11681+0.00236

[1900]	train-rmse:0.09054+0.00179	test-rmse:0.11673+0.00237
[2000]	train-rmse:0.09030+0.00180	test-rmse:0.11661+0.00237
[2100]	train-rmse:0.09008+0.00178	test-rmse:0.11650+0.00240
[2200]	train-rmse:0.08989+0.00176	test-rmse:0.11641+0.00241
[2300]	train-rmse:0.08973+0.00177	test-rmse:0.11634+0.00243
[2400]	train-rmse:0.08955+0.00176	test-rmse:0.11627+0.00245
[2500]	train-rmse:0.08940+0.00174	test-rmse:0.11618+0.00246
[2600]	train-rmse:0.08925+0.00173	test-rmse:0.11611+0.00246
[2700]	train-rmse:0.08910+0.00174	test-rmse:0.11603+0.00249
[2800]	train-rmse:0.08893+0.00173	test-rmse:0.11597+0.00248
[2900]	train-rmse:0.08880+0.00175	test-rmse:0.11589+0.00251
[3000]	train-rmse:0.08867+0.00173	test-rmse:0.11582+0.00251
[3100]	train-rmse:0.08853+0.00176	test-rmse:0.11575+0.00251
[3200]	train-rmse:0.08840+0.00175	test-rmse:0.11570+0.00250
[3300]	train-rmse:0.08827+0.00178	test-rmse:0.11564+0.00251
[3400]	train-rmse:0.08817+0.00178	test-rmse:0.11560+0.00251
[3500]	train-rmse:0.08805+0.00178	test-rmse:0.11555+0.00251
[3600]	train-rmse:0.08792+0.00177	test-rmse:0.11550+0.00249
[3700]	train-rmse:0.08779+0.00175	test-rmse:0.11545+0.00251
[3800]	train-rmse:0.08767+0.00173	test-rmse:0.11539+0.00252
[3900]	train-rmse:0.08756+0.00171	test-rmse:0.11535+0.00253
[4000]	train-rmse:0.08745+0.00172	test-rmse:0.11531+0.00252
[4100]	train-rmse:0.08732+0.00173	test-rmse:0.11525+0.00252
[4200]	train-rmse:0.08724+0.00173	test-rmse:0.11521+0.00254
[4300]	train-rmse:0.08716+0.00172	test-rmse:0.11516+0.00253
[4400]	train-rmse:0.08707+0.00171	test-rmse:0.11513+0.00255
[4500]	train-rmse:0.08698+0.00170	test-rmse:0.11509+0.00254
[4600]	train-rmse:0.08690+0.00169	test-rmse:0.11507+0.00254
[4700]	train-rmse:0.08681+0.00167	test-rmse:0.11502+0.00256
[4800]	train-rmse:0.08674+0.00166	test-rmse:0.11498+0.00256
[4900]	train-rmse:0.08668+0.00165	test-rmse:0.11495+0.00257
[5000]	train-rmse:0.08659+0.00165	test-rmse:0.11492+0.00258
[5100]	train-rmse:0.08652+0.00165	test-rmse:0.11489+0.00256
[5200]	train-rmse:0.08646+0.00166	test-rmse:0.11487+0.00257
[5300]	train-rmse:0.08638+0.00165	test-rmse:0.11483+0.00258
[5400]	train-rmse:0.08631+0.00165	test-rmse:0.11480+0.00258
[5500]	train-rmse:0.08625+0.00164	test-rmse:0.11476+0.00258
[5600]	train-rmse:0.08618+0.00163	test-rmse:0.11473+0.00258
[5700]	train-rmse:0.08609+0.00162	test-rmse:0.11468+0.00260
[5800]	train-rmse:0.08604+0.00161	test-rmse:0.11466+0.00260
[5900]	train-rmse:0.08596+0.00160	test-rmse:0.11462+0.00261
[6000]	train-rmse:0.08588+0.00160	test-rmse:0.11460+0.00260
[6100]	train-rmse:0.08580+0.00160	test-rmse:0.11456+0.00259
[6200]	train-rmse:0.08573+0.00160	test-rmse:0.11452+0.00259
[6300]	train-rmse:0.08566+0.00162	test-rmse:0.11450+0.00260
[6400]	train-rmse:0.08559+0.00162	test-rmse:0.11447+0.00260
[6500]	train-rmse:0.08554+0.00161	test-rmse:0.11444+0.00260
[6600]	train-rmse:0.08548+0.00161	test-rmse:0.11443+0.00260



[6700]	train-rmse:0.08541+0.00161	test-rmse:0.11440+0.00259
[6800]	train-rmse:0.08535+0.00161	test-rmse:0.11437+0.00260
[6900]	train-rmse:0.08529+0.00163	test-rmse:0.11435+0.00259
[7000]	train-rmse:0.08524+0.00162	test-rmse:0.11433+0.00260
[7100]	train-rmse:0.08518+0.00162	test-rmse:0.11430+0.00260
[7200]	train-rmse:0.08513+0.00161	test-rmse:0.11427+0.00260
[7300]	train-rmse:0.08509+0.00160	test-rmse:0.11425+0.00261
[7400]	train-rmse:0.08504+0.00159	test-rmse:0.11423+0.00261
[7500]	train-rmse:0.08498+0.00158	test-rmse:0.11421+0.00261
[7600]	train-rmse:0.08493+0.00158	test-rmse:0.11419+0.00260
[7700]	train-rmse:0.08488+0.00158	test-rmse:0.11416+0.00261
[7800]	train-rmse:0.08483+0.00157	test-rmse:0.11414+0.00261
[7900]	train-rmse:0.08477+0.00157	test-rmse:0.11411+0.00262
[8000]	train-rmse:0.08474+0.00156	test-rmse:0.11410+0.00262
[8100]	train-rmse:0.08469+0.00154	test-rmse:0.11408+0.00262
[8200]	train-rmse:0.08465+0.00153	test-rmse:0.11406+0.00263
[8300]	train-rmse:0.08461+0.00153	test-rmse:0.11404+0.00264
[8400]	train-rmse:0.08457+0.00153	test-rmse:0.11403+0.00264
[8500]	train-rmse:0.08453+0.00153	test-rmse:0.11401+0.00265
[8600]	train-rmse:0.08449+0.00153	test-rmse:0.11400+0.00265
[8700]	train-rmse:0.08445+0.00153	test-rmse:0.11397+0.00265
[8800]	train-rmse:0.08441+0.00152	test-rmse:0.11395+0.00266
[8900]	train-rmse:0.08436+0.00152	test-rmse:0.11393+0.00266
[8944]	train-rmse:0.08435+0.00153	test-rmse:0.11393+0.00266

```
[271]: inv_xgb_no_tune_params1 = {
    'subsample':0.7,
    'max_depth': 4,
    'colsample_bytree':0.6,
    'lambda': 100,
    'alpha':8,
    'eta': 0.05
}

inv_xgb_no_tune_params2 = {
    'subsample':1,
    'max_depth':4,
    'colsample_bytree':0.5,
    'lambda': 100,
    'alpha':8,
    'eta':0.05
}

inv_xgb_no_tune_params3 = {
    'subsample':1,
    'max_depth':4,
    'colsample_bytree':1,
    'lambda': 100,
```

```

    'alpha':8,
    'eta':0.05
}

inv_xgb_no_tune_params4 = {
    'subsample':0.9,
    'max_depth':4,
    'colsample_bytree':0.8,
    'lambda': 100,
    'alpha':8,
    'eta':0.05
}

inv_xgb_no_tune_params5 = {
    'subsample':0.7,
    'max_depth':6,
    'colsample_bytree':1,
    'lambda': 100,
    'alpha':8,
    'eta':0.05
}

inv_notue_xgb1 = no_tune_train(params=inv_xgb_no_tune_params1,
    ↪x_train=final_inv_train,y_train=inv_train.log_price)
inv_notue_xgb2 = no_tune_train(params=inv_xgb_no_tune_params2,
    ↪x_train=final_inv_train,y_train=inv_train.log_price)
inv_notue_xgb3 = no_tune_train(params=inv_xgb_no_tune_params3,
    ↪x_train=final_inv_train,y_train=inv_train.log_price)
inv_notue_xgb4 = no_tune_train(params=inv_xgb_no_tune_params4,
    ↪x_train=final_inv_train,y_train=inv_train.log_price)
inv_notue_xgb5 = no_tune_train(params=inv_xgb_no_tune_params5,
    ↪x_train=final_inv_train,y_train=inv_train.log_price)

```

[0]	train-rmse:0.37664+0.00145	test-rmse:0.37672+0.00315
[100]	train-rmse:0.19010+0.00260	test-rmse:0.19851+0.00594
[200]	train-rmse:0.17815+0.00263	test-rmse:0.19003+0.00638
[300]	train-rmse:0.17254+0.00256	test-rmse:0.18758+0.00650
[400]	train-rmse:0.16838+0.00252	test-rmse:0.18633+0.00642
[500]	train-rmse:0.16497+0.00249	test-rmse:0.18569+0.00646
[600]	train-rmse:0.16198+0.00244	test-rmse:0.18511+0.00644
[700]	train-rmse:0.15934+0.00238	test-rmse:0.18488+0.00639
[800]	train-rmse:0.15695+0.00237	test-rmse:0.18468+0.00637
[900]	train-rmse:0.15475+0.00242	test-rmse:0.18445+0.00633
[1000]	train-rmse:0.15270+0.00241	test-rmse:0.18428+0.00632
[1051]	train-rmse:0.15171+0.00240	test-rmse:0.18428+0.00632
[0]	train-rmse:0.37625+0.00147	test-rmse:0.37639+0.00313

[100]	train-rmse:0.18728+0.00264	test-rmse:0.19715+0.00600
[200]	train-rmse:0.17591+0.00280	test-rmse:0.18924+0.00620
[300]	train-rmse:0.17054+0.00273	test-rmse:0.18692+0.00619
[400]	train-rmse:0.16648+0.00275	test-rmse:0.18581+0.00614
[500]	train-rmse:0.16299+0.00268	test-rmse:0.18523+0.00619
[600]	train-rmse:0.16008+0.00270	test-rmse:0.18486+0.00614
[700]	train-rmse:0.15745+0.00270	test-rmse:0.18464+0.00618
[800]	train-rmse:0.15563+0.00307	test-rmse:0.18448+0.00615
[900]	train-rmse:0.15493+0.00404	test-rmse:0.18446+0.00612
[946]	train-rmse:0.15486+0.00415	test-rmse:0.18446+0.00612
[0]	train-rmse:0.37614+0.00152	test-rmse:0.37628+0.00306
[100]	train-rmse:0.18636+0.00275	test-rmse:0.19686+0.00598
[200]	train-rmse:0.17476+0.00262	test-rmse:0.18916+0.00633
[300]	train-rmse:0.16912+0.00253	test-rmse:0.18695+0.00649
[400]	train-rmse:0.16486+0.00250	test-rmse:0.18596+0.00649
[500]	train-rmse:0.16124+0.00259	test-rmse:0.18541+0.00641
[600]	train-rmse:0.15812+0.00265	test-rmse:0.18515+0.00645
[700]	train-rmse:0.15529+0.00265	test-rmse:0.18496+0.00645
[800]	train-rmse:0.15311+0.00240	test-rmse:0.18487+0.00647
[866]	train-rmse:0.15296+0.00237	test-rmse:0.18485+0.00648
[0]	train-rmse:0.37632+0.00151	test-rmse:0.37646+0.00312
[100]	train-rmse:0.18730+0.00275	test-rmse:0.19750+0.00596
[200]	train-rmse:0.17543+0.00293	test-rmse:0.18933+0.00628
[300]	train-rmse:0.16971+0.00281	test-rmse:0.18697+0.00639
[400]	train-rmse:0.16531+0.00272	test-rmse:0.18580+0.00633
[500]	train-rmse:0.16168+0.00269	test-rmse:0.18525+0.00633
[600]	train-rmse:0.15841+0.00274	test-rmse:0.18489+0.00634
[700]	train-rmse:0.15556+0.00263	test-rmse:0.18466+0.00633
[800]	train-rmse:0.15302+0.00263	test-rmse:0.18456+0.00631
[900]	train-rmse:0.15066+0.00255	test-rmse:0.18443+0.00628
[1000]	train-rmse:0.14847+0.00250	test-rmse:0.18435+0.00630
[1100]	train-rmse:0.14640+0.00247	test-rmse:0.18432+0.00631
[1200]	train-rmse:0.14449+0.00246	test-rmse:0.18422+0.00632
[1261]	train-rmse:0.14339+0.00244	test-rmse:0.18423+0.00633
[0]	train-rmse:0.37640+0.00143	test-rmse:0.37663+0.00311
[100]	train-rmse:0.17997+0.00286	test-rmse:0.19515+0.00638
[200]	train-rmse:0.16605+0.00277	test-rmse:0.18813+0.00661
[300]	train-rmse:0.15812+0.00263	test-rmse:0.18626+0.00649
[400]	train-rmse:0.15205+0.00244	test-rmse:0.18541+0.00645
[500]	train-rmse:0.14694+0.00247	test-rmse:0.18506+0.00642
[600]	train-rmse:0.14255+0.00233	test-rmse:0.18490+0.00651
[700]	train-rmse:0.13864+0.00221	test-rmse:0.18473+0.00647
[800]	train-rmse:0.13521+0.00214	test-rmse:0.18468+0.00652
[900]	train-rmse:0.13213+0.00212	test-rmse:0.18457+0.00655
[1000]	train-rmse:0.12938+0.00203	test-rmse:0.18453+0.00660
[1034]	train-rmse:0.12850+0.00200	test-rmse:0.18458+0.00660

## 9.2 Price SQ

```
[303]: ocu_xgb_no_tune_params1 = {
        'subsample':0.7,
        'max_depth': 9,
        'colsample_bytree':0.7,
        'lambda': 200,
        'eta': 0.1,
        'nthread':8
    }

    ocu_xgb_no_tune_params2 = {
        'subsample':1,
        'max_depth':9,
        'colsample_bytree':0.7,
        'lambda': 200,
        'eta':0.1,
        'nthread':8
    }

    ocu_xgb_no_tune_params3 = {
        'subsample':1,
        'max_depth':9,
        'colsample_bytree':1,
        'lambda': 200,
        'eta':0.1,
        'nthread':8
    }

    ocu_xgb_no_tune_params4 = {
        'subsample':0.9,
        'max_depth':7,
        'colsample_bytree':0.7,
        'lambda': 200,
        'eta':0.1,
        'nthread':8
    }

    ocu_xgb_no_tune_params5 = {
        'subsample':0.8,
        'max_depth':5,
        'colsample_bytree':0.7,
        'lambda': 200,
        'eta':0.1,
        'nthread':8
    }
```

```

ocu_xgb_no_tune_params5 = {
    'subsample':0.8,
    'max_depth':5,
    'colsample_bytree':0.7,
    'lambda': 200,
    'eta':0.1,
    'nthread':8
}

ocu_pricesq_xgb1 = no_tune_train(params=ocu_xgb_no_tune_params1,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_sq)
ocu_pricesq_xgb2 = no_tune_train(params=ocu_xgb_no_tune_params2,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_sq)
ocu_pricesq_xgb3 = no_tune_train(params=ocu_xgb_no_tune_params3,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_sq)
#ocu_pricesq_xgb4 = no_tune_train(params=ocu_xgb_no_tune_params4,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_sq)
#ocu_pricesq_xgb5 = no_tune_train(params=ocu_xgb_no_tune_params5,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_sq)

```

```

[0]      train-rmse:38484.04351+82.81214 test-rmse:38511.40089+121.35604
[100]    train-rmse:13424.62413+254.73118 test-rmse:16532.09987+654.36717
[200]    train-rmse:10627.66035+241.24923 test-rmse:15392.24064+817.18197
[300]    train-rmse:9016.62751+248.22539 test-rmse:14983.03044+855.50944
[400]    train-rmse:7917.29103+222.79735 test-rmse:14828.63538+832.76186
[500]    train-rmse:7065.66106+185.11194 test-rmse:14743.64086+827.03857
[600]    train-rmse:6350.22851+155.12067 test-rmse:14715.67248+815.60607
[700]    train-rmse:5761.05577+140.79376 test-rmse:14696.74167+809.29647
[800]    train-rmse:5252.27029+125.26369 test-rmse:14686.02123+797.07021
[900]    train-rmse:4816.80461+112.31727 test-rmse:14675.51928+789.44450
[1000]   train-rmse:4437.63322+100.90364 test-rmse:14672.16563+781.59906
[1029]   train-rmse:4338.73054+104.52579 test-rmse:14672.98742+781.00732
[0]      train-rmse:38325.34227+82.76809 test-rmse:38361.97021+125.98146
[100]    train-rmse:12264.10733+282.85695 test-rmse:15941.27286+640.49545
[200]    train-rmse:9571.51963+262.57019 test-rmse:14960.51730+765.37213
[300]    train-rmse:7977.62319+78.87495 test-rmse:14679.75808+794.12570
[400]    train-rmse:6926.04498+103.97775 test-rmse:14598.29623+798.35386
[500]    train-rmse:6111.25220+106.54485 test-rmse:14566.52595+799.86332
[600]    train-rmse:5467.41191+118.65159 test-rmse:14546.56774+787.89674
[663]    train-rmse:5131.79120+141.73843 test-rmse:14549.51370+789.06929
[0]      train-rmse:38285.06338+75.46018 test-rmse:38323.41696+129.52488
[100]    train-rmse:12269.69857+199.75133 test-rmse:16083.18546+670.70088
[200]    train-rmse:9638.59525+247.09952 test-rmse:15068.71348+793.10939
[300]    train-rmse:7991.22889+57.25974 test-rmse:14767.46078+800.76465
[400]    train-rmse:6832.73762+30.87685 test-rmse:14696.33936+766.84479
[500]    train-rmse:5967.88760+118.60616 test-rmse:14690.46550+742.77208

```

[530] train-rmse:5750.46650+136.85066 test-rmse:14693.80237+744.51364

```
[309]: inv_xgb_no_tune_params1 = {
    'subsample':0.7,
    'max_depth': 6,
    'colsample_bytree':0.6,
    'lambda': 8,
    'eta': 0.02,
    'nthread':8
}

inv_xgb_no_tune_params2 = {
    'subsample':0.9,
    'max_depth':6,
    'colsample_bytree':0.5,
    'lambda': 8,
    'eta':0.02,
    'nthread':8
}

inv_xgb_no_tune_params3 = {
    'subsample':1,
    'max_depth':6,
    'colsample_bytree':1,
    'lambda': 8,
    'eta':0.02,
    'nthread':8
}

inv_xgb_no_tune_params4 = {
    'subsample':0.9,
    'max_depth':6,
    'colsample_bytree':0.8,
    'lambda': 8,
    'eta':0.02,
    'nthread':8
}

inv_xgb_no_tune_params5 = {
    'subsample':0.7,
    'max_depth':6,
    'colsample_bytree':1,
    'lambda': 8,
    'eta':0.02,
    'nthread':8
}
```

```

inv_pricesq_xgb1 = no_tune_train(params=inv_xgb_no_tune_params1,
    ↪x_train=final_inv_train,y_train=inv_train.price_sq)
inv_pricesq_xgb2 = no_tune_train(params=inv_xgb_no_tune_params2,
    ↪x_train=final_inv_train,y_train=inv_train.price_sq)
inv_pricesq_xgb3 = no_tune_train(params=inv_xgb_no_tune_params3,
    ↪x_train=final_inv_train,y_train=inv_train.price_sq)
inv_pricesq_xgb4 = no_tune_train(params=inv_xgb_no_tune_params4,
    ↪x_train=final_inv_train,y_train=inv_train.price_sq)
inv_pricesq_xgb5 = no_tune_train(params=inv_xgb_no_tune_params5,
    ↪x_train=final_inv_train,y_train=inv_train.price_sq)

```

[0]	train-rmse:40112.21781+175.26007	test-rmse:40151.52418+345.31237
[100]	train-rmse:25713.19561+231.26029	test-rmse:28243.93444+570.67381
[200]	train-rmse:22325.83311+231.46293	test-rmse:26547.64555+654.96826
[300]	train-rmse:20572.27552+226.64285	test-rmse:26035.94167+685.93400
[400]	train-rmse:19274.49162+218.20516	test-rmse:25815.40726+690.83347
[500]	train-rmse:18143.59007+220.51647	test-rmse:25690.77861+706.13558
[600]	train-rmse:17098.61699+194.43773	test-rmse:25602.38414+700.71122
[700]	train-rmse:16191.43047+183.59547	test-rmse:25568.85820+708.72245
[800]	train-rmse:15317.22658+194.48421	test-rmse:25544.64502+701.07093
[900]	train-rmse:14523.65541+184.46306	test-rmse:25524.58302+697.47492
[1000]	train-rmse:13764.89153+192.50943	test-rmse:25513.95836+701.97912
[1096]	train-rmse:13066.95545+183.99071	test-rmse:25512.94185+708.13624
[0]	train-rmse:40104.87528+181.34007	test-rmse:40144.53793+344.42132
[100]	train-rmse:25573.80319+211.94962	test-rmse:28225.29001+546.58794
[200]	train-rmse:22125.25758+225.48297	test-rmse:26526.03713+616.55481
[300]	train-rmse:20356.30941+192.95062	test-rmse:25993.02056+670.48326
[400]	train-rmse:19055.85392+209.78975	test-rmse:25773.97448+676.99983
[500]	train-rmse:17975.95951+202.84791	test-rmse:25653.16361+673.04108
[600]	train-rmse:17017.54720+193.38738	test-rmse:25581.96372+676.04512
[700]	train-rmse:16109.07425+213.92915	test-rmse:25539.27983+671.51846
[800]	train-rmse:15304.90073+221.86642	test-rmse:25515.50266+675.53331
[900]	train-rmse:14528.64375+211.31627	test-rmse:25497.50124+665.35837
[1000]	train-rmse:13840.42151+210.31587	test-rmse:25483.04309+669.26762
[1100]	train-rmse:13187.12322+202.90460	test-rmse:25478.29638+662.99636
[1123]	train-rmse:13045.78137+193.52022	test-rmse:25479.14973+664.86461
[0]	train-rmse:40091.48104+176.62077	test-rmse:40142.81723+349.31856
[100]	train-rmse:25230.82412+237.48703	test-rmse:28210.36184+476.21528
[200]	train-rmse:21806.41111+268.46419	test-rmse:26626.35423+537.39232
[300]	train-rmse:20043.92496+256.15611	test-rmse:26142.33494+544.08696
[400]	train-rmse:18827.14995+239.72498	test-rmse:25938.31806+556.64564
[500]	train-rmse:17851.84147+248.96391	test-rmse:25838.92680+561.28374
[600]	train-rmse:17012.78311+245.99652	test-rmse:25791.86855+567.33443
[700]	train-rmse:16210.58251+252.21551	test-rmse:25765.34613+564.67262
[800]	train-rmse:15479.54069+267.64126	test-rmse:25736.91844+566.50680
[900]	train-rmse:14802.79544+253.15374	test-rmse:25725.68236+568.74879
[1000]	train-rmse:14213.70850+275.44110	test-rmse:25716.35282+574.22306

[1023]	train-rmse:14081.39975+284.61799	test-rmse:25716.26264+573.64021
[0]	train-rmse:40100.30738+179.80512	test-rmse:40142.23219+343.76530
[100]	train-rmse:25338.60307+227.13645	test-rmse:28167.47178+584.69803
[200]	train-rmse:21843.27759+223.35781	test-rmse:26537.58469+652.09021
[300]	train-rmse:20045.65537+211.10064	test-rmse:26045.88003+661.40152
[400]	train-rmse:18729.16594+213.14176	test-rmse:25848.98584+647.93477
[500]	train-rmse:17598.80925+193.07936	test-rmse:25729.37333+647.64883
[600]	train-rmse:16618.91879+193.68059	test-rmse:25670.73440+641.47182
[700]	train-rmse:15705.15004+176.72161	test-rmse:25639.43614+642.70056
[800]	train-rmse:14901.33047+180.09017	test-rmse:25618.40493+645.15552
[900]	train-rmse:14104.36939+178.64508	test-rmse:25598.74759+636.15278
[1000]	train-rmse:13385.27053+168.18778	test-rmse:25583.91390+630.86253
[1100]	train-rmse:12697.63255+185.21431	test-rmse:25580.00242+629.08821
[1128]	train-rmse:12525.32874+178.89176	test-rmse:25580.82581+631.77730
[0]	train-rmse:40103.01254+180.71323	test-rmse:40138.94820+345.24360
[100]	train-rmse:25462.45520+200.58399	test-rmse:28212.33727+540.62766
[200]	train-rmse:22017.62097+217.78137	test-rmse:26550.50226+639.96330
[300]	train-rmse:20262.26150+202.91240	test-rmse:26064.55233+668.81168
[400]	train-rmse:18879.71051+176.92455	test-rmse:25832.43824+701.68727
[500]	train-rmse:17725.28978+163.96314	test-rmse:25716.77768+693.56615
[600]	train-rmse:16666.77675+134.60030	test-rmse:25657.19211+697.88980
[700]	train-rmse:15684.15170+138.78549	test-rmse:25623.50753+696.17407
[800]	train-rmse:14819.52773+125.28874	test-rmse:25601.33754+690.66189
[900]	train-rmse:13977.16086+118.24544	test-rmse:25592.65241+694.78896
[938]	train-rmse:13670.83779+125.35484	test-rmse:25595.15545+696.90130

### 9.3 Price Doc

```
[274]: ocu_price_xgb1 = no_tune_train(params=ocu_xgb_no_tune_params1,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_doc)
ocu_price_xgb2 = no_tune_train(params=ocu_xgb_no_tune_params2,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_doc)
ocu_price_xgb3 = no_tune_train(params=ocu_xgb_no_tune_params3,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_doc)
ocu_price_xgb4 = no_tune_train(params=ocu_xgb_no_tune_params4,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_doc)
ocu_price_xgb5 = no_tune_train(params=ocu_xgb_no_tune_params5,
    ↪x_train=final_ocu_train,y_train=ocu_train.price_doc)
```

[0]	train-rmse:3997561.35034+105907.26534	test-
	rmse:3997081.92588+227108.23764	
[100]	train-rmse:1394180.32318+86822.30984	test-
	rmse:1673751.15972+266801.50298	
[200]	train-rmse:1051343.93733+71406.55071	test-
	rmse:1497363.61191+245541.14749	
[300]	train-rmse:869121.88629+62188.59237	test-
	rmse:1436081.40288+231210.73280	
[400]	train-rmse:746920.60417+54914.22926	test-



rmse:1407370.71659+222621.85112	
[500] train-rmse:657828.42981+46388.24218	test-
rmse:1392457.41027+217308.31632	
[600] train-rmse:587418.87263+38347.58312	test-
rmse:1386128.96860+213790.19186	
[700] train-rmse:529879.66017+32439.42607	test-
rmse:1380192.71594+212187.48130	
[800] train-rmse:483781.76497+29393.89820	test-
rmse:1375431.63458+209616.37949	
[900] train-rmse:443424.80433+25430.11737	test-
rmse:1372541.69579+207543.18844	
[1000] train-rmse:407289.65558+22673.81195	test-
rmse:1370148.80911+205875.23850	
[1100] train-rmse:375589.71539+19566.97559	test-
rmse:1368918.56551+204935.82996	
[1200] train-rmse:349257.13216+15859.37876	test-
rmse:1366730.11108+203297.71671	
[1300] train-rmse:325308.27184+13025.27191	test-
rmse:1365600.23417+202121.56421	
[1400] train-rmse:302809.34006+11362.30371	test-
rmse:1364136.23722+200970.25694	
[1500] train-rmse:282005.92583+9487.32808	test-
rmse:1363208.78678+199406.67977	
[1600] train-rmse:263749.76846+7819.41576	test-
rmse:1362316.06255+198625.90066	
[1610] train-rmse:262018.96954+7742.22473	test-
rmse:1362382.12922+198567.95347	
[0] train-rmse:3978031.79744+106207.97446	test-
rmse:3977753.44184+226479.15722	
[100] train-rmse:1248745.58872+92606.70357	test-
rmse:1646493.33832+239105.42166	
[200] train-rmse:904591.02401+73968.62705	test-
rmse:1505622.38333+211284.00806	
[300] train-rmse:727679.84730+66292.46162	test-
rmse:1462294.76952+202792.85029	
[400] train-rmse:616156.73383+61664.51860	test-
rmse:1447805.11636+197219.92882	
[500] train-rmse:535950.80763+56300.16763	test-
rmse:1440216.01910+196990.70706	
[600] train-rmse:471992.74527+53704.25628	test-
rmse:1439148.71753+194216.39226	
[626] train-rmse:458183.64695+52270.28194	test-
rmse:1439232.00450+193360.81779	
[0] train-rmse:3975060.63967+106008.51476	test-
rmse:3974454.77499+228557.03379	
[100] train-rmse:1247946.14588+86255.35580	test-
rmse:1665299.74475+222225.38388	
[200] train-rmse:901868.40513+75082.27595	test-

rmse:1513483.83563+197845.33536	
[300] train-rmse:725082.21264+68646.16521	test-
rmse:1466005.43412+197298.04604	
[400] train-rmse:617047.47961+64738.15837	test-
rmse:1450227.88181+191227.39511	
[500] train-rmse:538284.40295+59413.04001	test-
rmse:1440286.98212+192480.34926	
[600] train-rmse:476692.25396+57885.64957	test-
rmse:1436645.82945+191843.26492	
[700] train-rmse:427307.67592+54636.27095	test-
rmse:1435350.04289+191429.16404	
[737] train-rmse:411263.14789+52587.23798	test-
rmse:1435956.65036+191096.84349	
[0] train-rmse:3984757.12494+106809.38134	test-
rmse:3985838.44425+226177.01859	
[100] train-rmse:1331189.64699+83768.44691	test-
rmse:1647793.48111+254199.11299	
[200] train-rmse:997806.40222+66781.86007	test-
rmse:1488428.56847+229620.07790	
[300] train-rmse:820876.59977+57749.60709	test-
rmse:1434257.57585+216740.94759	
[400] train-rmse:709729.59530+53927.52529	test-
rmse:1412130.72759+206956.76367	
[500] train-rmse:627523.84348+49761.25172	test-
rmse:1402497.62110+203519.27867	
[600] train-rmse:563851.25879+45790.70835	test-
rmse:1395961.07597+202246.39288	
[700] train-rmse:509540.10028+40941.08015	test-
rmse:1392063.94834+201005.92191	
[800] train-rmse:464995.89564+37105.48977	test-
rmse:1390047.33828+197339.24772	
[900] train-rmse:426708.10981+34861.87501	test-
rmse:1386681.07121+195064.15872	
[1000] train-rmse:393128.34315+32147.49630	test-
rmse:1385392.41189+192932.16615	
[1100] train-rmse:363571.15845+29416.86011	test-
rmse:1383479.45949+191619.70405	
[1200] train-rmse:337179.09100+27003.40852	test-
rmse:1383271.10611+190477.82924	
[1268] train-rmse:320922.72691+25821.72209	test-
rmse:1382731.02821+190024.99982	
[0] train-rmse:3992403.96377+106915.93132	test-
rmse:3992740.33130+226019.18639	
[100] train-rmse:1428979.74952+90263.31373	test-
rmse:1667233.66481+243986.47294	
[200] train-rmse:1108232.68972+73569.48101	test-
rmse:1492304.20052+217852.33233	
[300] train-rmse:942841.63809+63181.87466	test-

```

rmse:1431992.08613+206694.54979
[400]   train-rmse:838206.22716+57201.44562   test-
rmse:1400696.74763+202063.78003
[500]   train-rmse:758287.91466+51587.89689   test-
rmse:1386948.26794+199504.61293
[600]   train-rmse:696563.20358+46119.32101   test-
rmse:1379162.98271+198542.67495
[700]   train-rmse:643826.18608+41276.44584   test-
rmse:1371636.04107+197337.53571
[800]   train-rmse:601242.40969+38817.89750   test-
rmse:1365479.49925+199287.39000
[900]   train-rmse:564054.97774+36152.25326   test-
rmse:1360999.65899+197804.41170
[1000]  train-rmse:530983.91608+32408.56501   test-
rmse:1357040.21418+197366.07974
[1100]  train-rmse:500330.09741+29148.05657   test-
rmse:1354163.73257+197442.53700
[1200]  train-rmse:473446.11305+26681.67241   test-
rmse:1351172.95591+196829.87273
[1300]  train-rmse:449622.32073+24139.51927   test-
rmse:1348660.11430+195789.33249
[1400]  train-rmse:426752.22542+22186.84452   test-
rmse:1346112.74600+195905.78086
[1500]  train-rmse:406594.07579+20640.65933   test-
rmse:1344299.24611+195096.84050
[1600]  train-rmse:387985.98385+18832.22095   test-
rmse:1342151.83547+194244.79270
[1700]  train-rmse:370899.51843+17513.54765   test-
rmse:1340552.93239+194364.65011
[1782]  train-rmse:357566.62267+16522.56775   test-
rmse:1340465.62127+193425.69979

```

```

[275]: inv_price_xgb1 = no_tune_train(params=inv_xgb_no_tune_params1,
    ↪x_train=final_inv_train,y_train=inv_train.price_doc)
inv_price_xgb2 = no_tune_train(params=inv_xgb_no_tune_params2,
    ↪x_train=final_inv_train,y_train=inv_train.price_doc)
inv_price_xgb3 = no_tune_train(params=inv_xgb_no_tune_params3,
    ↪x_train=final_inv_train,y_train=inv_train.price_doc)
inv_price_xgb4 = no_tune_train(params=inv_xgb_no_tune_params4,
    ↪x_train=final_inv_train,y_train=inv_train.price_doc)
inv_price_xgb5 = no_tune_train(params=inv_xgb_no_tune_params5,
    ↪x_train=final_inv_train,y_train=inv_train.price_doc)

```

```

[0]      train-rmse:4043486.19608+36468.84202   test-
rmse:4044329.11215+75340.12220
[100]    train-rmse:2249147.07222+23808.76128   test-
rmse:2381206.05804+76191.10618
[200]    train-rmse:1845647.60587+25251.86052   test-

```

rmse:2105554.70170+75403.99085	
[300] train-rmse:1665873.98608+25838.22730	test-
rmse:2021312.20021+72613.03314	
[400] train-rmse:1551986.08521+26007.90318	test-
rmse:1985194.00300+72208.09522	
[500] train-rmse:1463856.55551+25284.63476	test-
rmse:1964009.74085+72673.81490	
[600] train-rmse:1390000.65129+22823.39715	test-
rmse:1950926.57437+72698.40867	
[700] train-rmse:1326289.33063+20527.95404	test-
rmse:1941020.80665+74043.51078	
[800] train-rmse:1268927.08281+18281.34543	test-
rmse:1933900.38222+75395.58635	
[900] train-rmse:1217075.57841+17523.84799	test-
rmse:1928562.42327+75680.89146	
[1000] train-rmse:1168848.58893+16602.98446	test-
rmse:1924649.44471+75254.87038	
[1100] train-rmse:1125195.02139+17243.66110	test-
rmse:1922216.16729+75060.90572	
[1200] train-rmse:1083337.20123+16731.74384	test-
rmse:1919814.65123+75143.44725	
[1300] train-rmse:1044385.35173+15069.48272	test-
rmse:1918121.53693+74582.63198	
[1400] train-rmse:1008416.55665+14662.74867	test-
rmse:1917356.21028+74053.97044	
[1500] train-rmse:973388.73458+14571.86667	test-
rmse:1915639.89278+73864.41274	
[1600] train-rmse:940774.30977+14457.74543	test-
rmse:1915332.06271+73891.79529	
[1700] train-rmse:909161.28372+13439.87048	test-
rmse:1914032.87467+73727.51454	
[1758] train-rmse:891251.96502+13158.31824	test-
rmse:1914355.49274+73534.23542	
[0] train-rmse:4041471.41179+36545.53622	test-
rmse:4041871.39023+75163.62067	
[100] train-rmse:2180124.46646+23672.12454	test-
rmse:2343486.30162+75334.89217	
[200] train-rmse:1770879.89749+26434.86679	test-
rmse:2088696.62701+70151.31055	
[300] train-rmse:1594265.41564+26510.19340	test-
rmse:2021775.02678+70177.46108	
[400] train-rmse:1484458.46467+24560.78909	test-
rmse:1994530.70672+67368.51788	
[500] train-rmse:1402576.48665+25066.42280	test-
rmse:1980329.67462+65518.93756	
[600] train-rmse:1332983.15808+23742.16278	test-
rmse:1967989.74936+65259.13694	
[700] train-rmse:1277335.09354+21306.34515	test-

rmse:1960454.57306+67227.13531	
[800] train-rmse:1228226.04951+23214.15579	test-
rmse:1954142.82927+68029.48527	
[900] train-rmse:1186654.93469+23537.91964	test-
rmse:1949768.98062+69676.23347	
[1000] train-rmse:1149611.49689+23524.92588	test-
rmse:1946091.92848+71399.21782	
[1100] train-rmse:1115209.25906+26115.29020	test-
rmse:1943029.39941+71585.39492	
[1200] train-rmse:1080604.25419+27548.32816	test-
rmse:1940840.26185+72499.12017	
[1300] train-rmse:1048709.88782+28171.65430	test-
rmse:1938668.02241+72928.39297	
[1400] train-rmse:1021457.85193+31470.77282	test-
rmse:1937199.00448+73202.71581	
[1500] train-rmse:993718.34810+34352.60862	test-
rmse:1936021.80372+73648.59426	
[1600] train-rmse:965323.75885+36255.84834	test-
rmse:1934426.70583+73757.97814	
[1700] train-rmse:940535.13217+36819.36237	test-
rmse:1933278.27178+74201.38760	
[1800] train-rmse:917688.00554+37384.27542	test-
rmse:1932573.03721+74345.45829	
[1900] train-rmse:893951.71080+36788.34627	test-
rmse:1932148.92171+74343.92994	
[2000] train-rmse:870540.14355+35485.32042	test-
rmse:1931916.94930+74327.35828	
[2005] train-rmse:869668.32329+35648.72730	test-
rmse:1931878.97248+74262.25697	
[0] train-rmse:4041226.53873+36792.39114	test-
rmse:4042016.89226+74635.88386	
[100] train-rmse:2147387.76784+24249.80386	test-
rmse:2328734.07918+77243.79879	
[200] train-rmse:1733676.90007+27554.46697	test-
rmse:2082331.03388+73905.70690	
[300] train-rmse:1560548.05884+28406.66973	test-
rmse:2024539.83973+70593.55976	
[400] train-rmse:1452841.41044+25564.32407	test-
rmse:2001491.24006+66866.38264	
[500] train-rmse:1376052.44596+21962.81536	test-
rmse:1989048.14714+64405.62733	
[600] train-rmse:1313167.39070+19351.46242	test-
rmse:1981140.79211+63804.30732	
[700] train-rmse:1256474.96026+19755.16100	test-
rmse:1974529.85863+64606.73918	
[800] train-rmse:1209634.13501+17179.12316	test-
rmse:1969359.11122+66762.24200	
[900] train-rmse:1167652.46646+17089.40011	test-

rmse:1966219.20757+67356.59756	
[1000] train-rmse:1124627.47855+16742.56655	test-
rmse:1963054.65878+68078.35713	
[1100] train-rmse:1085194.11627+19453.75139	test-
rmse:1960317.06427+67625.94408	
[1200] train-rmse:1046791.20419+23744.55959	test-
rmse:1958135.99043+67150.89346	
[1300] train-rmse:1012028.63218+25252.97531	test-
rmse:1956874.38948+65752.36160	
[1400] train-rmse:982789.27433+23248.05093	test-
rmse:1955373.82594+64979.56943	
[1500] train-rmse:956043.48882+25576.63589	test-
rmse:1953972.27009+64989.37629	
[1600] train-rmse:931886.12399+27086.56076	test-
rmse:1952778.04328+64898.10389	
[1700] train-rmse:907490.03355+27734.03570	test-
rmse:1951976.37321+65107.49833	
[1770] train-rmse:890952.28457+27016.52955	test-
rmse:1952453.24044+65214.26672	
[0] train-rmse:4041984.42475+36868.35073	test-
rmse:4042764.12688+75021.88277	
[100] train-rmse:2175149.44773+24481.00680	test-
rmse:2339045.26282+76054.41297	
[200] train-rmse:1764710.81164+26031.60538	test-
rmse:2083766.72698+74729.70154	
[300] train-rmse:1588183.88334+24810.60626	test-
rmse:2015922.29775+74631.16914	
[400] train-rmse:1475264.71278+20932.07205	test-
rmse:1988053.93699+74137.51349	
[500] train-rmse:1388662.34406+19175.83817	test-
rmse:1970995.86367+73731.79423	
[600] train-rmse:1315728.30461+17071.22357	test-
rmse:1959296.49250+74657.90874	
[700] train-rmse:1253868.28067+15735.59066	test-
rmse:1951499.29773+75217.57252	
[800] train-rmse:1199866.25213+14479.32097	test-
rmse:1944826.47649+75851.14457	
[900] train-rmse:1149549.09536+14738.72244	test-
rmse:1940257.42801+76298.52679	
[1000] train-rmse:1101754.40605+15346.50042	test-
rmse:1936436.03673+76624.09108	
[1100] train-rmse:1058772.42677+15347.14280	test-
rmse:1933766.21159+76683.62407	
[1200] train-rmse:1017500.60026+16556.81313	test-
rmse:1931016.23058+76884.53109	
[1300] train-rmse:979042.80579+16947.82591	test-
rmse:1929333.04791+76346.96563	
[1343] train-rmse:962730.65365+16515.36384	test-

```

rmse:1929400.83698+76142.40635
[0]      train-rmse:4043522.05654+36766.05716      test-
rmse:4044544.21878+75053.53737
[100]    train-rmse:2224504.60118+24339.71964      test-
rmse:2368560.70577+76465.27177
[200]    train-rmse:1816567.09249+26049.02661      test-
rmse:2094882.35156+75024.60426
[300]    train-rmse:1636782.75636+27360.25282      test-
rmse:2014820.61809+72343.32953
[400]    train-rmse:1521394.28778+26136.25626      test-
rmse:1982381.55533+70592.63641
[500]    train-rmse:1431080.85251+24504.66126      test-
rmse:1962757.48940+71704.99363
[600]    train-rmse:1357123.29865+22810.52549      test-
rmse:1949418.15580+72945.04377
[700]    train-rmse:1292604.42354+19765.62509      test-
rmse:1941151.67791+74068.21655
[800]    train-rmse:1235670.29448+18664.83472      test-
rmse:1935334.65777+75699.22464
[900]    train-rmse:1183329.53238+18687.00806      test-
rmse:1930773.31688+76831.76962
[1000]   train-rmse:1135287.95086+18013.08607      test-
rmse:1927799.27970+76296.92360
[1100]   train-rmse:1090039.31820+17119.13112      test-
rmse:1924822.16660+76986.08356
[1200]   train-rmse:1046774.60390+17441.20928      test-
rmse:1922503.80137+76730.09554
[1300]   train-rmse:1007038.42853+16750.53477      test-
rmse:1921082.15892+76400.12381
[1400]   train-rmse:969586.98041+16741.59316      test-
rmse:1919998.46766+76617.30364
[1500]   train-rmse:934117.79187+16267.62226      test-
rmse:1919145.43514+77261.02644
[1600]   train-rmse:900575.23520+15485.12903      test-
rmse:1918614.26699+77186.80492
[1700]   train-rmse:868028.32141+15662.44445      test-
rmse:1918059.08445+77011.90647
[1800]   train-rmse:836743.47058+15422.11641      test-
rmse:1917894.21762+76857.20362
[1900]   train-rmse:807782.98975+14863.04531      test-
rmse:1916901.22863+77562.87533
[1957]   train-rmse:791963.81313+14666.68188      test-
rmse:1916835.61480+78059.77616

```

## 9.4 No Tune Top Features Models

```
[276]: def getTopFeatures(model, n_features=100):
        scores = model.get_score()
        scores_items = scores.items()
        sorted_imp = sorted(list(scores_items), key=lambda x:x[1], reverse=True)
        return [f for f, v in sorted_imp][:n_features]

[277]: ocu_topf_xgb1 = no_tune_train(params=ocu_xgb_no_tune_params1,
    ↪x_train=final_ocu_train[getTopFeatures(ocu_pricesq_xgb1)],y_train=ocu_train.
    ↪price_sq)
ocu_topf_xgb2 = no_tune_train(params=ocu_xgb_no_tune_params2,
    ↪x_train=final_ocu_train[getTopFeatures(ocu_pricesq_xgb2)],y_train=ocu_train.
    ↪price_sq)
ocu_topf_xgb3 = no_tune_train(params=ocu_xgb_no_tune_params3,
    ↪x_train=final_ocu_train[getTopFeatures(ocu_pricesq_xgb3)],y_train=ocu_train.
    ↪price_sq)
#ocu_topf_xgb4 = no_tune_train(params=ocu_xgb_no_tune_params4,
    ↪x_train=final_ocu_train[getTopFeatures(ocu_pricesq_xgb4)],y_train=ocu_train.
    ↪price_sq)
#ocu_topf_xgb5 = no_tune_train(params=ocu_xgb_no_tune_params5,
    ↪x_train=final_ocu_train[getTopFeatures(ocu_pricesq_xgb5)],y_train=ocu_train.
    ↪price_sq)

[0]      train-rmse:38480.25836+81.13618 test-rmse:38509.61063+127.08556
[100]    train-rmse:13450.93157+236.11228 test-rmse:16551.25581+712.61958
[200]    train-rmse:10660.64047+224.24871 test-rmse:15417.84622+891.01862
[300]    train-rmse:9034.54293+216.08831 test-rmse:15002.09528+952.65054
[400]    train-rmse:7923.52501+190.03904 test-rmse:14835.55598+937.81474
[500]    train-rmse:7076.68285+153.40698 test-rmse:14772.91208+929.39184
[600]    train-rmse:6357.57283+124.30423 test-rmse:14746.82396+925.44247
[700]    train-rmse:5760.42348+111.06358 test-rmse:14718.41485+923.41680
[800]    train-rmse:5253.35127+91.99815 test-rmse:14708.35794+914.98522
[900]    train-rmse:4811.45094+81.62272 test-rmse:14703.31453+912.89991
[984]    train-rmse:4491.67784+81.37224 test-rmse:14701.78423+903.22484
[0]      train-rmse:38313.47596+72.78976 test-rmse:38342.36822+144.02281
[100]    train-rmse:12316.29503+250.01243 test-rmse:16048.47925+660.59775
[200]    train-rmse:9512.16455+238.50445 test-rmse:15043.77187+772.75257
[300]    train-rmse:7944.55545+116.26590 test-rmse:14748.35372+820.66088
[400]    train-rmse:6874.55676+58.06755 test-rmse:14656.63145+809.46041
[500]    train-rmse:6085.15487+39.49892 test-rmse:14609.27512+818.50555
[600]    train-rmse:5457.29092+53.22647 test-rmse:14595.59835+807.53000
[631]    train-rmse:5301.29442+64.90939 test-rmse:14592.37970+808.34383
[0]      train-rmse:38288.36350+76.36158 test-rmse:38323.12761+128.15257
[100]    train-rmse:12372.90869+255.93006 test-rmse:16099.24765+655.62374
[200]    train-rmse:9689.96468+311.59243 test-rmse:15051.00492+771.45411
[300]    train-rmse:8030.22770+119.78028 test-rmse:14729.33478+755.18571
[398]    train-rmse:6939.96994+83.70394 test-rmse:14710.85477+741.51071
```



```
[278]: inv_topf_xgb1 = no_tune_train(params=inv_xgb_no_tune_params1,
    ↪x_train=final_inv_train[getTopFeatures(inv_pricesq_xgb1)],y_train=inv_train.
    ↪price_sq)
inv_topf_xgb2 = no_tune_train(params=inv_xgb_no_tune_params2,
    ↪x_train=final_inv_train[getTopFeatures(inv_pricesq_xgb2)],y_train=inv_train.
    ↪price_sq)
inv_topf_xgb3 = no_tune_train(params=inv_xgb_no_tune_params3,
    ↪x_train=final_inv_train[getTopFeatures(inv_pricesq_xgb3)],y_train=inv_train.
    ↪price_sq)
inv_topf_xgb4 = no_tune_train(params=inv_xgb_no_tune_params4,
    ↪x_train=final_inv_train[getTopFeatures(inv_pricesq_xgb4)],y_train=inv_train.
    ↪price_sq)
inv_topf_xgb5 = no_tune_train(params=inv_xgb_no_tune_params5,
    ↪x_train=final_inv_train[getTopFeatures(inv_pricesq_xgb5)],y_train=inv_train.
    ↪price_sq)
```

[0]	train-rmse:40188.62492+177.40797	test-rmse:40212.88273+347.07107
[100]	train-rmse:28392.54466+216.86992	test-rmse:29643.56770+518.69098
[200]	train-rmse:25253.95135+208.64492	test-rmse:27501.59138+578.58448
[300]	train-rmse:23733.77698+201.55973	test-rmse:26750.27888+614.49033
[400]	train-rmse:22685.10013+211.08082	test-rmse:26393.06148+610.95818
[500]	train-rmse:21798.86978+206.76134	test-rmse:26186.44478+605.41645
[600]	train-rmse:21018.03054+197.75076	test-rmse:26051.76119+599.71006
[700]	train-rmse:20292.82960+190.10698	test-rmse:25961.49427+596.47046
[800]	train-rmse:19624.90386+197.89920	test-rmse:25898.35222+591.47515
[900]	train-rmse:18997.72127+206.65107	test-rmse:25856.91614+594.80101
[1000]	train-rmse:18404.91357+205.97959	test-rmse:25827.42656+596.52622
[1100]	train-rmse:17838.09382+209.15900	test-rmse:25803.19312+608.44313
[1200]	train-rmse:17289.39440+197.81026	test-rmse:25781.71462+618.38652
[1300]	train-rmse:16777.89312+195.15464	test-rmse:25762.74809+619.28543
[1400]	train-rmse:16286.83209+195.65997	test-rmse:25748.90355+616.85970
[1500]	train-rmse:15807.54567+192.79624	test-rmse:25743.09850+615.38244
[1587]	train-rmse:15404.31244+178.15734	test-rmse:25741.66095+620.75626
[0]	train-rmse:40187.93095+174.97001	test-rmse:40214.55040+346.12069
[100]	train-rmse:27971.82087+220.13231	test-rmse:29411.08706+517.91060
[200]	train-rmse:24817.64186+205.57655	test-rmse:27371.62960+561.31726
[300]	train-rmse:23361.85378+156.89010	test-rmse:26668.24135+607.56686
[400]	train-rmse:22382.49977+150.43176	test-rmse:26341.30107+622.81527
[500]	train-rmse:21643.32341+157.54743	test-rmse:26145.00262+608.69736
[600]	train-rmse:21050.87644+151.45903	test-rmse:26030.61498+611.70209
[700]	train-rmse:20478.34004+149.68113	test-rmse:25937.11594+600.42143
[800]	train-rmse:19962.05830+165.21195	test-rmse:25873.23178+602.74874
[900]	train-rmse:19464.35173+172.27927	test-rmse:25839.53285+602.31145
[1000]	train-rmse:18978.42563+223.98332	test-rmse:25801.21838+604.52684
[1100]	train-rmse:18546.94747+219.40645	test-rmse:25780.87241+603.65471
[1200]	train-rmse:18133.13932+219.05227	test-rmse:25762.65042+607.65149
[1300]	train-rmse:17744.62011+225.07041	test-rmse:25750.76321+603.99949

[1400]	train-rmse:17365.01708+231.02701	test-rmse:25743.91109+601.96341
[1500]	train-rmse:16994.48046+226.72679	test-rmse:25735.96554+605.35267
[1600]	train-rmse:16630.78477+227.45289	test-rmse:25732.68458+606.77851
[1700]	train-rmse:16287.15180+219.07090	test-rmse:25726.22624+605.10600
[1800]	train-rmse:15954.08343+247.47148	test-rmse:25720.46310+602.05059
[1900]	train-rmse:15620.54710+245.79497	test-rmse:25717.95941+601.33518
[1960]	train-rmse:15392.44470+239.53008	test-rmse:25719.75958+603.33606
[0]	train-rmse:40170.65943+175.97934	test-rmse:40199.03484+346.54597
[100]	train-rmse:27756.86964+184.54660	test-rmse:29330.72921+562.90134
[200]	train-rmse:24570.25212+202.90736	test-rmse:27350.36134+564.50490
[300]	train-rmse:23077.73266+202.69012	test-rmse:26688.37384+586.75233
[400]	train-rmse:22120.46195+204.63214	test-rmse:26397.20994+594.67928
[500]	train-rmse:21347.00612+167.70847	test-rmse:26218.96759+586.35482
[600]	train-rmse:20718.28297+169.54591	test-rmse:26117.49068+584.63736
[700]	train-rmse:20151.09661+207.12199	test-rmse:26056.17381+602.39327
[800]	train-rmse:19616.32382+253.15304	test-rmse:26000.49501+600.71442
[900]	train-rmse:19106.41110+304.33337	test-rmse:25977.33688+600.20992
[1000]	train-rmse:18601.27208+329.12612	test-rmse:25950.42745+596.18063
[1100]	train-rmse:18131.82159+351.74902	test-rmse:25938.31369+608.93124
[1200]	train-rmse:17718.93856+372.16970	test-rmse:25923.79128+615.06662
[1300]	train-rmse:17254.15139+363.94660	test-rmse:25911.54838+617.77232
[1400]	train-rmse:16820.38421+361.56976	test-rmse:25908.86748+623.39152
[1500]	train-rmse:16381.06445+372.70024	test-rmse:25902.67643+629.53789
[1596]	train-rmse:16019.69431+367.42542	test-rmse:25900.33207+634.90301
[0]	train-rmse:40175.63966+178.41277	test-rmse:40205.92439+344.81681
[100]	train-rmse:27935.51560+200.47362	test-rmse:29421.82850+521.07143
[200]	train-rmse:24723.62504+193.45018	test-rmse:27343.74825+551.34975
[300]	train-rmse:23190.22962+173.88321	test-rmse:26640.06179+584.36887
[400]	train-rmse:22148.36304+179.04470	test-rmse:26311.72839+580.19395
[500]	train-rmse:21272.83296+216.44463	test-rmse:26122.11669+562.89769
[600]	train-rmse:20519.79434+240.15753	test-rmse:25998.86927+558.64119
[700]	train-rmse:19832.91416+269.24801	test-rmse:25912.42447+556.49551
[800]	train-rmse:19200.64319+254.82880	test-rmse:25852.46725+559.45630
[900]	train-rmse:18610.67244+268.17517	test-rmse:25815.32478+567.56869
[1000]	train-rmse:18026.14861+244.01982	test-rmse:25785.06333+573.92962
[1100]	train-rmse:17459.88193+240.07534	test-rmse:25761.43337+573.40287
[1200]	train-rmse:16925.49182+217.20829	test-rmse:25744.31568+578.00609
[1300]	train-rmse:16428.48231+209.37279	test-rmse:25730.98595+582.98151
[1400]	train-rmse:15925.97466+192.67954	test-rmse:25723.72613+582.24616
[1500]	train-rmse:15452.08746+162.71861	test-rmse:25719.40868+582.72074
[1599]	train-rmse:14994.87538+162.73872	test-rmse:25715.55738+580.02934
[0]	train-rmse:40186.19248+176.82888	test-rmse:40217.10260+346.45892
[100]	train-rmse:28190.51441+202.27347	test-rmse:29548.13205+558.61784
[200]	train-rmse:25021.05225+212.05065	test-rmse:27437.54462+577.76091
[300]	train-rmse:23494.40959+215.08911	test-rmse:26706.16056+595.33185
[400]	train-rmse:22421.10088+202.95816	test-rmse:26374.68863+592.99432
[500]	train-rmse:21519.00599+212.02507	test-rmse:26170.10374+590.25967
[600]	train-rmse:20709.72781+218.14082	test-rmse:26045.07051+598.49376

[700]	train-rmse:19966.02672+202.07606	test-rmse:25960.20179+597.80586
[800]	train-rmse:19254.93749+208.85360	test-rmse:25897.12068+590.44165
[900]	train-rmse:18606.31104+202.90210	test-rmse:25860.18920+589.13329
[1000]	train-rmse:18000.08812+193.48419	test-rmse:25827.76936+595.37427
[1100]	train-rmse:17413.36043+190.54787	test-rmse:25805.45991+605.92677
[1200]	train-rmse:16842.71905+169.44983	test-rmse:25792.03178+609.55155
[1300]	train-rmse:16313.13698+166.42451	test-rmse:25778.57865+614.42751
[1400]	train-rmse:15805.52165+156.78438	test-rmse:25775.89303+611.91107
[1410]	train-rmse:15753.86174+158.32541	test-rmse:25775.32894+611.60849

## 10 Predictions for Kaggle

```
[304]: mn = 84/90
      mn_inv = 77/90
```

### 10.1 No Tune Price SQ

```
[305]: inv_pricesq_preds = [inv_pricesq_xgb1.predict(xgb.
      ↪DMatrix(inv_test))*inv_test['full_sq'],
inv_pricesq_xgb2.predict(xgb.DMatrix(inv_test))*inv_test['full_sq'],
inv_pricesq_xgb3.predict(xgb.DMatrix(inv_test))*inv_test['full_sq'],
inv_pricesq_xgb4.predict(xgb.DMatrix(inv_test))*inv_test['full_sq'],
inv_pricesq_xgb5.predict(xgb.DMatrix(inv_test))*inv_test['full_sq']
]

ocu_pricesq_preds = [ocu_pricesq_xgb1.predict(xgb.
      ↪DMatrix(ocu_test))*ocu_test['full_sq'],
ocu_pricesq_xgb2.predict(xgb.DMatrix(ocu_test))*ocu_test['full_sq'],
ocu_pricesq_xgb3.predict(xgb.DMatrix(ocu_test))*ocu_test['full_sq'],
#ocu_pricesq_xgb4.predict(xgb.DMatrix(ocu_test))*ocu_test['full_sq'],
#ocu_pricesq_xgb5.predict(xgb.DMatrix(ocu_test))*ocu_test['full_sq']
]
```

```
[306]: # Average/ Median the predictions
ocu_pricesq_preds = np.median(np.array(ocu_pricesq_preds),axis=0)
inv_pricesq_preds = np.median(np.array(inv_pricesq_preds),axis=0)
```

```
[307]: # Create the data frame of the predictions.
inv_pricesq_notune = pd.DataFrame(dict(id=new_inv[new_inv['price_doc'].
      ↪isnull()]['id'], price_doc = inv_pricesq_preds*mn_inv))
ocu_pricesq_notune = pd.DataFrame(dict(id=new_ocu[new_ocu['price_doc'].
      ↪isnull()]['id'], price_doc = ocu_pricesq_preds*mn))
xgb_pricesq_finals = pd.concat([ocu_pricesq_notune,inv_pricesq_notune])
```

## 10.2 No Tune Log Price

```
[283]: # Get the predictions for each model
inv_notunepreds = [np.expm1(inv_notue_xgb1.predict(xgb.DMatrix(inv_test))),
                   np.expm1(inv_notue_xgb2.predict(xgb.DMatrix(inv_test))),
                   np.expm1(inv_notue_xgb3.predict(xgb.DMatrix(inv_test))),
                   np.expm1(inv_notue_xgb4.predict(xgb.DMatrix(inv_test))),
                   np.expm1(inv_notue_xgb5.predict(xgb.DMatrix(inv_test)))]

ocu_notunepreds = [np.expm1(ocu_notue_xgb1.predict(xgb.DMatrix(ocu_test))),
                   np.expm1(ocu_notue_xgb2.predict(xgb.DMatrix(ocu_test))),
                   np.expm1(ocu_notue_xgb3.predict(xgb.DMatrix(ocu_test))),
                   np.expm1(ocu_notue_xgb4.predict(xgb.DMatrix(ocu_test))),
                   np.expm1(ocu_notue_xgb5.predict(xgb.DMatrix(ocu_test)))]
```

```
[284]: # Average the predictions
ocu_notune_preds = np.mean(np.array(ocu_notunepreds),axis=0)
inv_notune_preds = np.mean(np.array(inv_notunepreds),axis=0)
```

```
[285]: # Create the data frame of the predictions.
inv_xgb_notune = pd.DataFrame(dict(id=new_inv[new_inv['price_doc'].
    ↪isnull()]['id'], price_doc = inv_notune_preds*mn_inv))
ocu_xgb_notune = pd.DataFrame(dict(id=new_ocu[new_ocu['price_doc'].
    ↪isnull()]['id'], price_doc = ocu_notune_preds*mn))
xgb_notune_finals = pd.concat([ocu_xgb_notune,inv_xgb_notune])
```

## 10.3 No Tune Price Doc

```
[286]: inv_price_preds = [inv_price_xgb1.predict(xgb.DMatrix(inv_test)),
                           inv_price_xgb2.predict(xgb.DMatrix(inv_test)),
                           inv_price_xgb3.predict(xgb.DMatrix(inv_test)),
                           inv_price_xgb4.predict(xgb.DMatrix(inv_test)),
                           inv_price_xgb5.predict(xgb.DMatrix(inv_test))]

ocu_price_preds = [ocu_price_xgb1.predict(xgb.DMatrix(ocu_test)),
                   ocu_price_xgb2.predict(xgb.DMatrix(ocu_test)),
                   ocu_price_xgb3.predict(xgb.DMatrix(ocu_test)),
                   ocu_price_xgb4.predict(xgb.DMatrix(ocu_test)),
                   ocu_price_xgb5.predict(xgb.DMatrix(ocu_test))]
```

```
[287]: # Average the predictions
ocu_price_preds = np.mean(np.array(ocu_price_preds),axis=0)
inv_price_preds = np.mean(np.array(inv_price_preds),axis=0)
```

```
[288]: # Create the data frame of the predictions.
```

```

inv_xgb_price = pd.DataFrame(dict(id=new_inv[new_inv['price_doc']].
    ↳isnull())['id'], price_doc = inv_price_preds*mn_inv))
ocu_xgb_price = pd.DataFrame(dict(id=new_ocu[new_ocu['price_doc']].
    ↳isnull())['id'], price_doc = ocu_price_preds*mn))
xgb_price_finals = pd.concat([ocu_xgb_price,inv_xgb_price])

```

## 10.4 No Tune Top Features Price SQ

```

[316]: inv_topf_preds = [inv_topf_xgb1.predict(xgb.
    ↳DMatrix(inv_test[getTopFeatures(inv_pricesq_xgb1)])))*inv_test['full_sq'],
inv_topf_xgb2.predict(xgb.
    ↳DMatrix(inv_test[getTopFeatures(inv_pricesq_xgb2)])))*inv_test['full_sq'],
inv_topf_xgb3.predict(xgb.
    ↳DMatrix(inv_test[getTopFeatures(inv_pricesq_xgb3)])))*inv_test['full_sq'],
inv_topf_xgb4.predict(xgb.
    ↳DMatrix(inv_test[getTopFeatures(inv_pricesq_xgb4)])))*inv_test['full_sq'],
inv_topf_xgb5.predict(xgb.
    ↳DMatrix(inv_test[getTopFeatures(inv_pricesq_xgb5)])))*inv_test['full_sq']]

ocu_topf_preds = [ocu_topf_xgb1.predict(xgb.
    ↳DMatrix(ocu_test[getTopFeatures(ocu_pricesq_xgb1)])))*ocu_test['full_sq'],
ocu_topf_xgb2.predict(xgb.
    ↳DMatrix(ocu_test[getTopFeatures(ocu_pricesq_xgb2)])))*ocu_test['full_sq'],
ocu_topf_xgb3.predict(xgb.
    ↳DMatrix(ocu_test[getTopFeatures(ocu_pricesq_xgb3)])))*ocu_test['full_sq'],
#ocu_topf_xgb4.predict(xgb.
    ↳DMatrix(ocu_test[getTopFeatures(ocu_pricesq_xgb4)])))*ocu_test['full_sq'],
#ocu_topf_xgb5.predict(xgb.
    ↳DMatrix(ocu_test[getTopFeatures(ocu_pricesq_xgb5)])))*ocu_test['full_sq']
]

```

```

[317]: inv_topf_preds = np.mean(np.array(inv_topf_preds),axis=0)
ocu_topf_preds = np.mean(np.array(ocu_topf_preds),axis=0)

```

```

[318]: inv_xgb_topf= pd.DataFrame(dict(id=new_inv[new_inv['price_doc']].
    ↳isnull())['id'], price_doc = inv_topf_preds*mn_inv))
ocu_xgb_topf = pd.DataFrame(dict(id=new_ocu[new_ocu['price_doc']].
    ↳isnull())['id'], price_doc = ocu_topf_preds*mn))
xgb_topf_finals = pd.concat([ocu_xgb_topf,inv_xgb_topf])
xgb_topf_finals['price_doc'] = xgb_topf_finals['price_doc']

```

## 10.5 Top 5 models - bayesian optimization

```
[292]: top_inv_preds = [np.exp(model['model'].predict(inv_test)) - 1 for model in
    ↪top_xgb_inv]
top_ocu_preds = [np.exp(model['model'].predict(ocu_test)) - 1 for model in
    ↪top_xgb_ocu]

top_rf_inv_preds = [np.expm1(model['model'].predict(inv_test)) for model in
    ↪top5_rf_inv]
top_rf_ocu_preds = [np.expm1(model['model'].predict(ocu_test)) for model in
    ↪top5_rf_ocu]
```

```
[293]: top_avg_inv_price = np.array(top_inv_preds).mean(axis=0)
top_avg_ocu_price = np.array(top_ocu_preds).mean(axis=0)

top_avg_inv_rf = np.array(top_rf_inv_preds).mean(axis=0)
top_avg_ocu_rf = np.array(top_rf_ocu_preds).mean(axis=0)
```

```
[294]: inv_xgb_top = pd.DataFrame(dict(id = new_inv[new_inv['price_doc'].
    ↪isnull()]['id'], price_doc = top_avg_inv_price*mn_inv))
ocu_xgb_top = pd.DataFrame(dict(id = new_ocu[new_ocu['price_doc'].
    ↪isnull()]['id'], price_doc = top_avg_ocu_price*mn))
xgb_top_finals = pd.concat([ocu_xgb_top, inv_xgb_top])

inv_rf_top = pd.DataFrame(dict(id = new_inv[new_inv['price_doc'].
    ↪isnull()]['id'], price_doc = top_avg_inv_rf*mn_inv))
ocu_rf_top = pd.DataFrame(dict(id = new_ocu[new_ocu['price_doc'].
    ↪isnull()]['id'], price_doc = top_avg_ocu_rf*mn))
rf_top_finals = pd.concat([ocu_rf_top, inv_rf_top])
```

## 10.6 Top 1 Model - bayesian optimization

```
[295]: inv_xgb_preds = np.exp(best_model_inv.predict(inv_test)) - 1
#inv_xgb_preds = np.exp(best_model_inv.predict(inv_test))*inv_test['full_sq']
ocu_xgb_preds = np.exp(best_model_ocu.predict(ocu_test)) - 1

inv_xgb = pd.DataFrame(dict(id = new_inv[new_inv['price_doc'].isnull()]['id'],
    ↪price_doc = inv_xgb_preds*mn_inv))
#inv_xgb = pd.DataFrame(dict(id=new_inv[new_inv['price_doc'].isnull()]['id'],
    ↪reset_index(drop=True), price_doc=inv_xgb_preds.reset_index(drop=True)))
ocu_xgb = pd.DataFrame(dict(id = new_ocu[new_ocu['price_doc'].isnull()]['id'],
    ↪price_doc = ocu_xgb_preds*mn))
xgb_finals = pd.concat([ocu_xgb, inv_xgb])
```

```

rf_inv_preds = np.exp(best_rf_inv.predict(inv_test)) - 1
# rf_inv_preds = best_rf_inv.predict(inv_test)*inv_test['full_sq']
rf_ocu_preds = np.exp(best_rf_ocu.predict(ocu_test)) - 1
inv_rf = pd.DataFrame(dict(id = new_inv[new_inv['price_doc'].isnull()]['id'],
    ↪ price_doc = rf_inv_preds*mn_inv))
# inv_rf = pd.DataFrame(dict(id=inv_done[inv_done['price_doc'].isnull()]['id'],
    ↪ reset_index(drop=True), price_doc=rf_inv_preds.reset_index(drop=True)))
ocu_rf = pd.DataFrame(dict(id = new_ocu[new_ocu['price_doc'].isnull()]['id'],
    ↪ price_doc = rf_ocu_preds*mn))
rf_finals = pd.concat([ocu_rf, inv_rf])

```

## 10.7 Ensemble

Top 1 Model - Investment (Bayesian Optimization) , Pricesq Investment Model

For OwnerOccupier - pricesq model only.

```

[313]: ensemble_inv = pd.DataFrame(dict(id=inv_xgb['id'],price_doc =
    ↪ inv_xgb['price_doc']*0.12 + inv_pricesq_notune['price_doc']*0.88))
ensemble_inv = pd.concat([ocu_pricesq_notune,ensemble_inv])

```

## 11 Kaggle Submission Files

```

[308]: xgb_pricesq_finals.to_csv("./price_sq_xgb.csv",index=False)

```

```

[314]: # Ensemble
ensemble_inv.to_csv("./ensemble_inv_normalocu.csv",index=False)

```

```

[315]: # Random Forest Predictions
rf_finals.to_csv("./rf_submission.csv", index=False)
# XGboost Predictions
xgb_finals.to_csv("./xgb_submission.csv", index=False)

# # XGboost (without algorithm for tuning)
xgb_notune_finals.to_csv("./xgb_notune.csv", index=False)

# # XGboost (without algorithm for tuning) TOP FEATURES
xgb_topf_finals.to_csv("./xgb_topf_notune.csv", index=False)

# # XGBOOST (without algorithm for tuning) PRICE SQ
xgb_pricesq_finals.to_csv("./price_sq_xgb.csv",index=False)

# # XGboost (without algorithm for tuning) PRICE DOC

```

```

xgb_price_finals.to_csv("./price_doc_xgb.csv", index=False)

# Top 5 xgboost models
xgb_top_finals.to_csv("./xgb_top5_submission.csv", index=False)

# Top 5 random forest models
rf_top_finals.to_csv("./rf_top5_submission.csv", index=False)

# # Weighted Average
pd.DataFrame(dict(id=rf_finals['id'], price_doc=(xgb_finals['price_doc']*0.
↪9999999 + rf_finals['price_doc']*0.01))).to_csv("./top1_rf_xgb_ensemble.
↪csv", index=False)
pd.DataFrame(dict(id=xgb_top_finals['id'], ↵
↪price_doc=(xgb_top_finals['price_doc']*0.9999999 + ↵
↪rf_top_finals['price_doc']*0.01))).to_csv("./top5_rf_xgb_ensemble.
↪csv", index=False)
pd.DataFrame(dict(id=xgb_notune_finals['id'], ↵
↪price_doc=xgb_notune_finals['price_doc']*0.9 + xgb_finals['price_doc']*0.1)).
↪to_csv("./xgb_stack_ensemble.csv", index=False)
pd.DataFrame(dict(id=xgb_notune_finals['id'], ↵
↪price_doc=xgb_notune_finals['price_doc']*0.9 + xgb_top_finals['price_doc']*0.
↪1))).to_csv("./xgb_stack_top5_ensemble.csv", index=False)
pd.DataFrame(dict(id=xgb_pricesq_finals['id'], ↵
↪price_doc=xgb_pricesq_finals['price_doc']*0.9+ ↵
↪xgb_notune_finals['price_doc']*(1-0.9))).to_csv("./
↪xgb_pricesq_stack_ensemble.csv", index=False)

```

## 12 Bibliography

Kapoor, S., & Perrone, V. (2021). A simple and fast baseline for tuning large XGBoost models. arXiv preprint arXiv:2111.06924.

Putatunda, S., & Rama, K. (2018). A Comparative Analysis of Hyperopt as Against Other Approaches for Hyper-Parameter Optimization of XGBoost. International Conference on Signal Processing and Machine Learning.